covid | the 10 largest locations without Florida pooled | evaluation: Jun 2025 to Sep 2026 | loss in counts per 100k residents

Bars, top to bottom: the 10 best on this page (quantile loss after peak, all designs and horizons)

- Leaky MultiQT, adaptive rate
- Spread scale: most recent starts, learned + sharing
- Spread scale: most recent starts, equal
- Weekly starts: most recent, learned + sharing
- Power scale: most recent starts, learned + sharing
- Power scale: most recent starts, equal
- Delphi REF-op (covid prod)
- Leaky MultiQT (log1p tracker), adaptive rate
- Fixed MultiQT, adaptive rate
- Weekly starts: most recent, slow + fast, learned + sharing

Colours as in the trade-off plots of R/04.

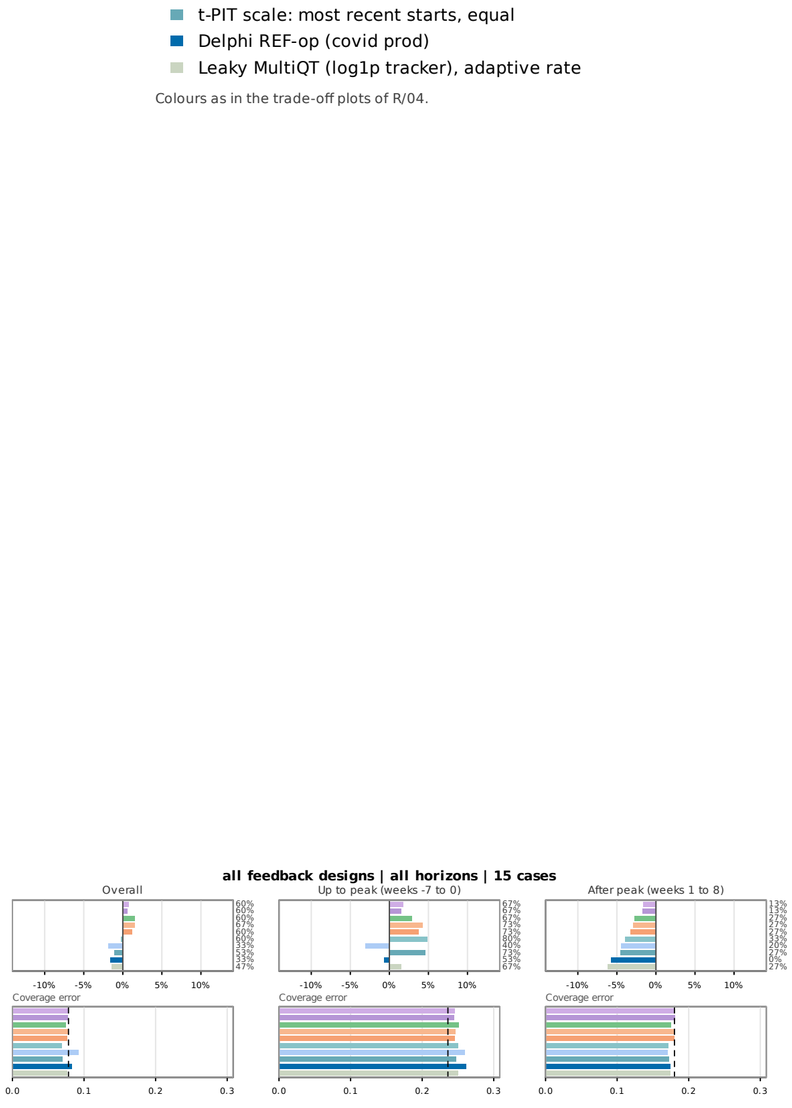
Loss bar: 100 * (1 - pooled loss / pooled loss of the base), pooled over the node's cases (location x horizon x feedback design), each case's mean loss weighted by its number of weeks; to the right: positive = smaller loss than the base.
Numbers right of each loss panel: share of the node's cases in which the method's loss is below the base's (nodes with at least 5 cases). Up to peak / after peak: the weeks around every local peak of the latest counts.
Below each loss panel: coverage error, the mean over the quantile levels of |coverage - alpha| over the weeks of the window, averaged over the cases of the node weighted by their weeks; dashed: the base forecasts.

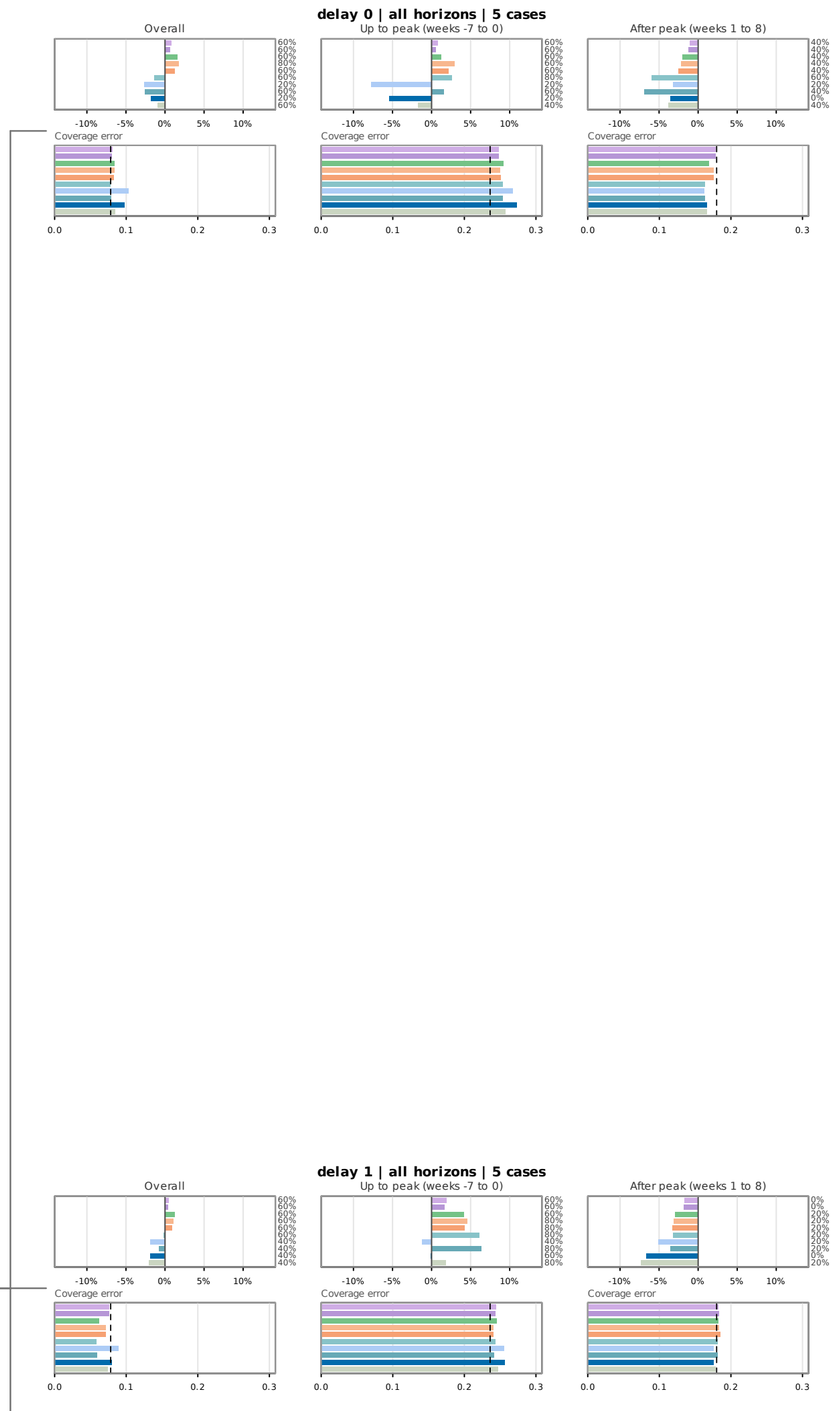
covid | United States (US) | evaluation: Jun 2025 to Sep 2026 | loss in counts per 100k residents

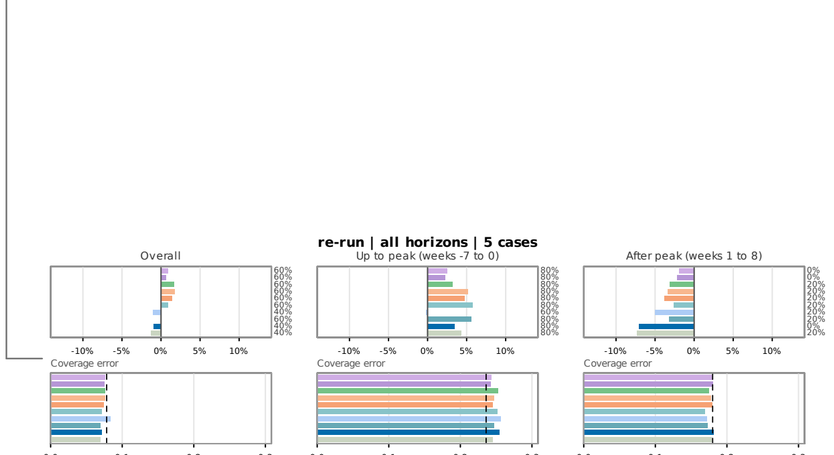
Bars, top to bottom: the 10 best on this page (quantile loss after peak, all designs and horizons)

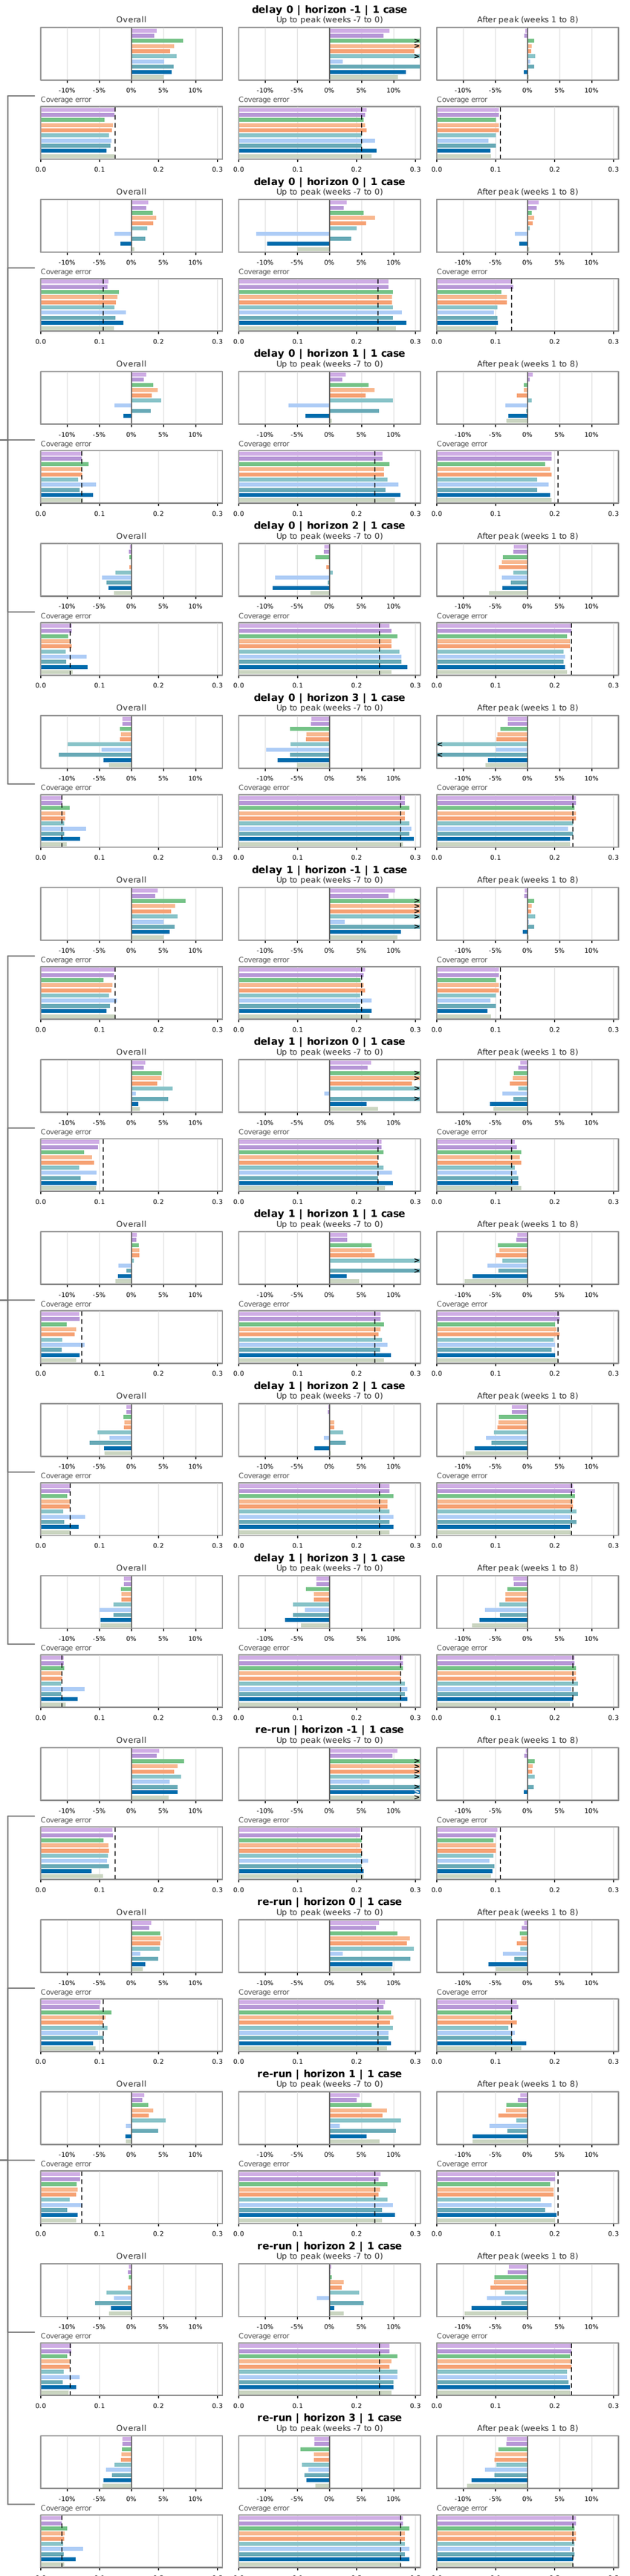
- Spread scale: most recent starts, learned + sharing
- Spread scale: most recent starts, equal
- Weekly starts: most recent, learned + sharing
- Power scale: most recent starts, learned + sharing
- Power scale: most recent starts, equal
- t-PIT scale: most recent starts, learned + sharing
- Leaky MultiQT, adaptive rate
- t-PIT scale: most recent starts, equal
- Delphi REF-op (covid prod)
- Leaky MultiQT (log1p tracker), adaptive rate

Colours as in the trade-off plots of R/04.









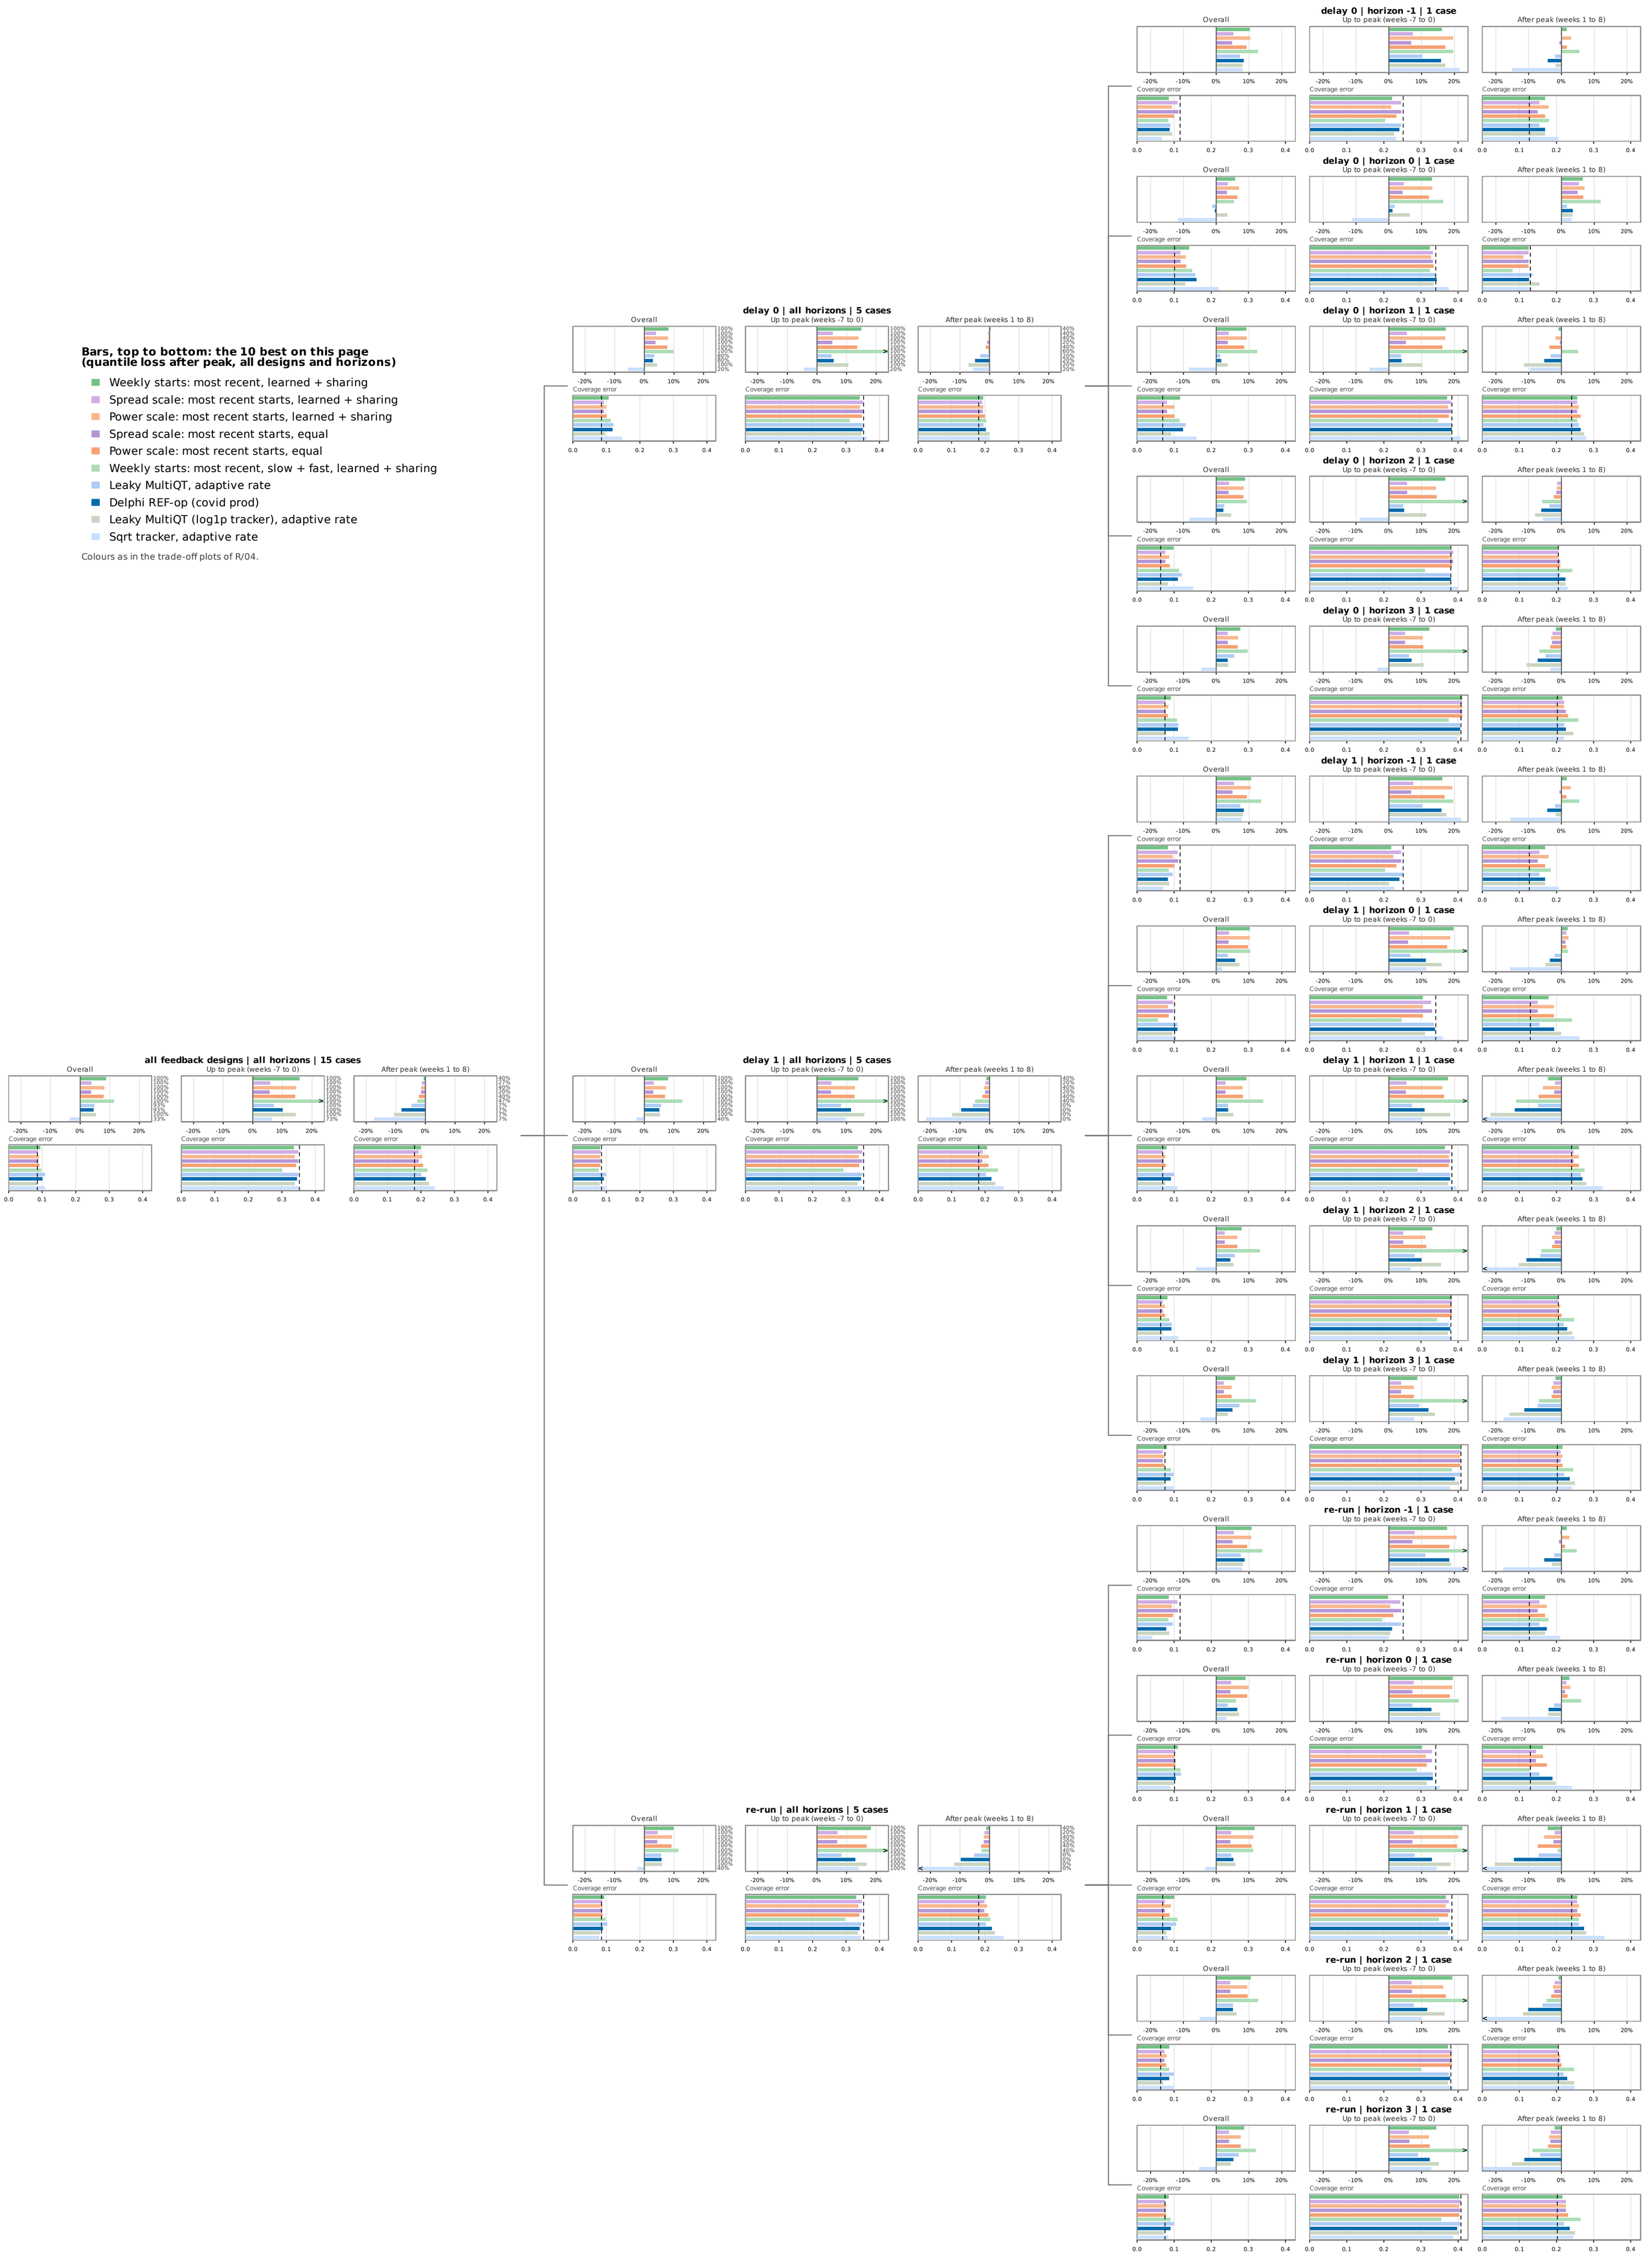
Loss bar: 100 * (1 - pooled loss / pooled loss of the base), pooled over the node's cases (location x horizon x feedback design), each case's mean loss weighted by its number of weeks; to the right: positive = smaller loss than the base.
Numbers right of each loss panel: share of the node's cases in which the method's loss is below the base's (nodes with at least 5 cases). Up to peak / after peak: the weeks around every local peak of the latest counts.
Below each loss panel: coverage error, the mean over the quantile levels of |coverage - alpha| over the weeks of the window, averaged over the cases of the node weighted by their weeks; dashed: the base forecasts.

covid | California (06) | evaluation: Jun 2025 to Sep 2026 | loss in counts per 100k residents

Bars, top to bottom: the 10 best on this page (quantile loss after peak, all designs and horizons)

- Weekly starts: most recent, learned + sharing
- Spread scale: most recent starts, learned + sharing
- Power scale: most recent starts, learned + sharing
- Spread scale: most recent starts, equal
- Power scale: most recent starts, equal
- Weekly starts: most recent, slow + fast, learned + sharing
- Leaky MultiQT, adaptive rate
- Delphi REF-op (covid prod)
- Leaky MultiQT (log1p tracker), adaptive rate
- Sqrt tracker, adaptive rate

Colours as in the trade-off plots of R/04.



Loss bar: 100 * (1 - pooled loss / pooled loss of the base), pooled over the node's cases (location x horizon x feedback design), each case's mean loss weighted by its number of weeks; to the right: positive = smaller loss than the base.
Numbers right of each loss panel: share of the node's cases in which the method's loss is below the base's (nodes with at least 5 cases). Up to peak / after peak: the weeks around every local peak of the latest counts.
Below each loss panel: coverage error, the mean over the quantile levels of |coverage - alpha| over the weeks of the window, averaged over the cases of the node weighted by their weeks; dashed: the base forecasts.

covid | Texas (48) | evaluation: Jun 2025 to Sep 2026 | loss in counts per 100k residents

Bars, top to bottom: the 10 best on this page (quantile loss after peak, all designs and horizons)

- Weekly starts: most recent, learned + sharing
- Spread scale: most recent starts, learned + sharing
- Spread scale: most recent starts, equal
- Power scale: most recent starts, learned + sharing
- Leaky MultiQT, adaptive rate
- Power scale: most recent starts, equal
- Weekly starts: most recent, slow + fast, learned + sharing
- Delphi REF-op (covid prod)
- Leaky MultiQT (log1p tracker), adaptive rate
- Fixed MultiQT, adaptive rate

Colours as in the trade-off plots of R/04.



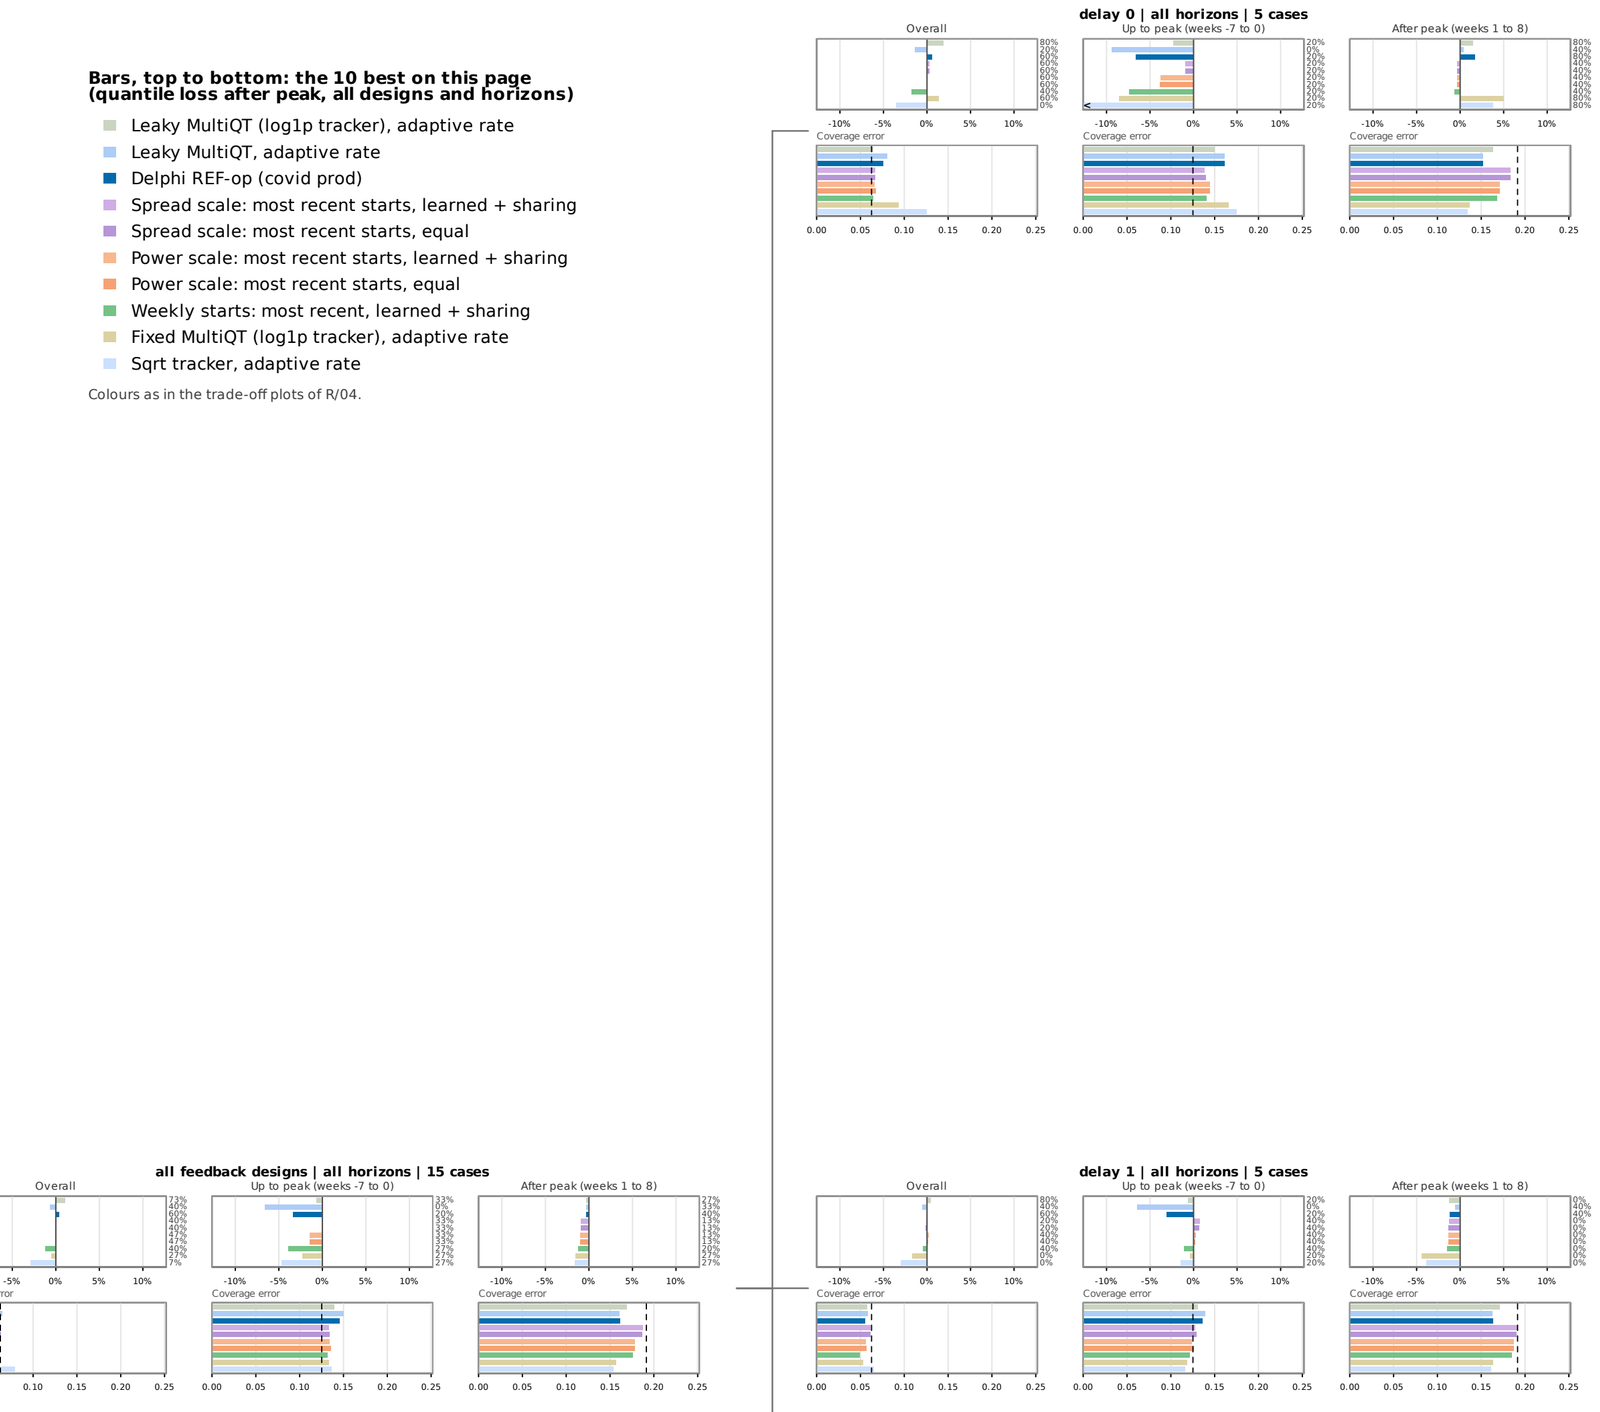
Loss bar: 100 * (1 - pooled loss / pooled loss of the base), pooled over the node's cases (location x horizon x feedback design), each case's mean loss weighted by its number of weeks; to the right: positive = smaller loss than the base.
Numbers right of each loss panel: share of the node's cases in which the method's loss is below the base's (nodes with at least 5 cases). Up to peak / after peak: the weeks around every local peak of the latest counts.
Below each loss panel: coverage error, the mean over the quantile levels of |coverage - alpha| over the weeks of the window, averaged over the cases of the node weighted by their weeks; dashed: the base forecasts.

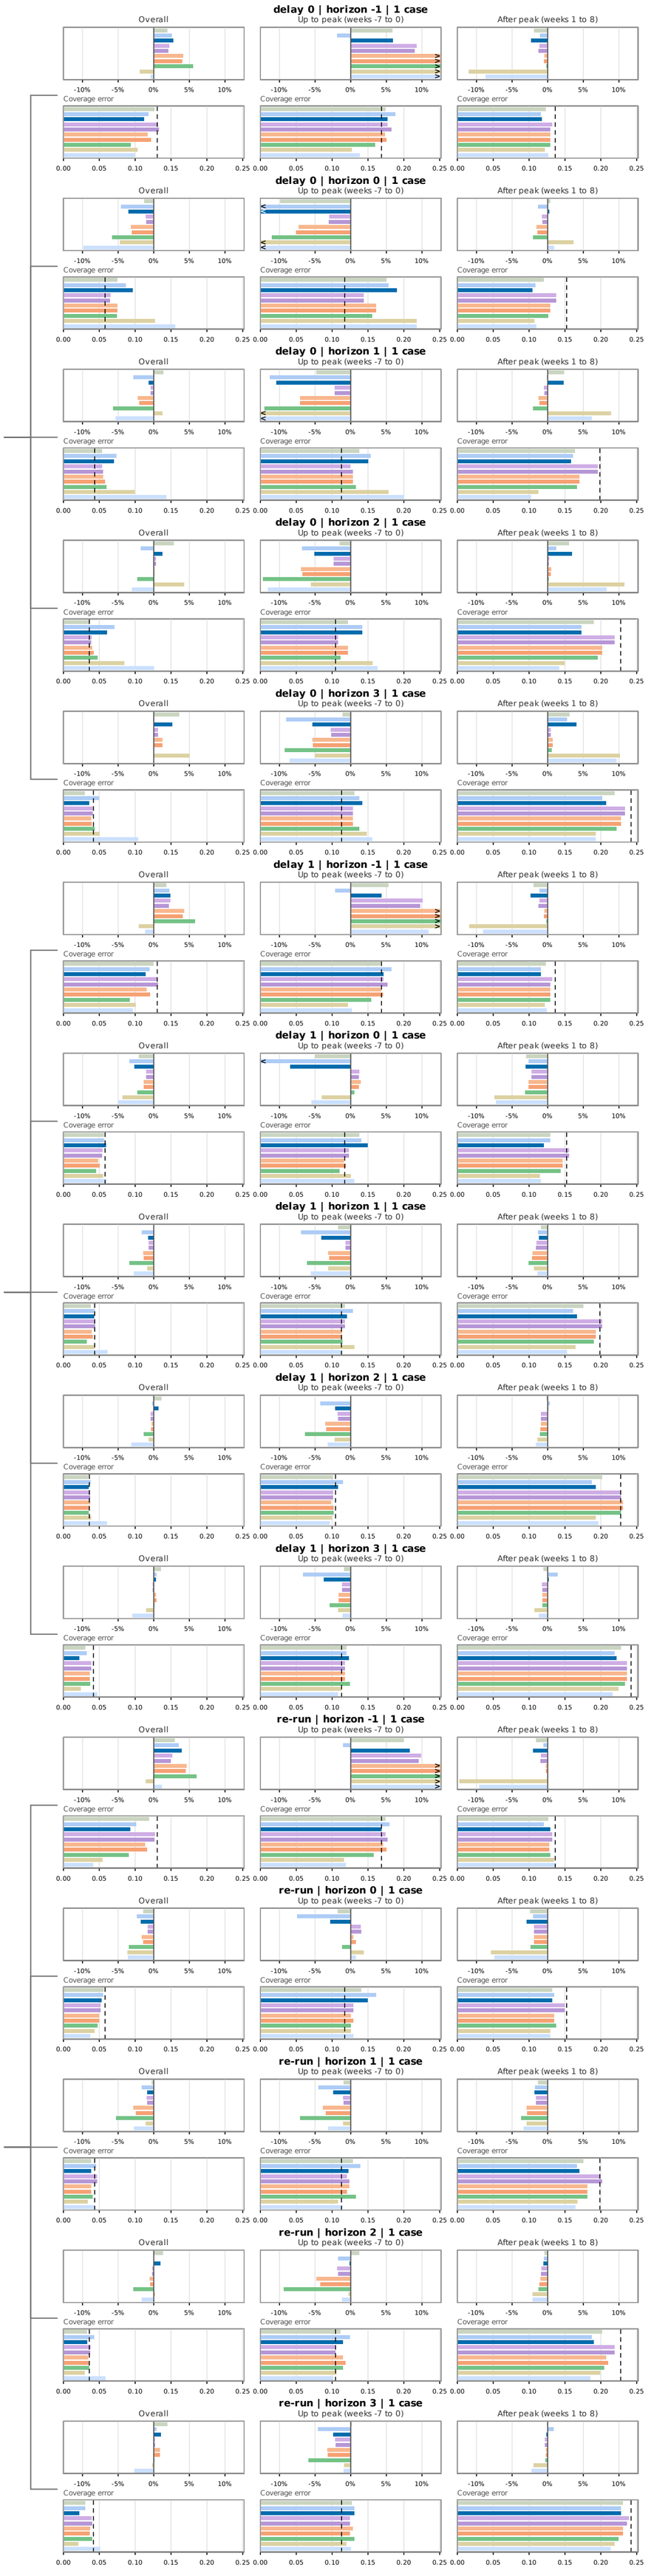
covid | New York (36) | evaluation: Jun 2025 to Sep 2026 | loss in counts per 100k residents

Bars, top to bottom: the 10 best on this page (quantile loss after peak, all designs and horizons)

- Leaky MultiQT (log1p tracker), adaptive rate
- Leaky MultiQT, adaptive rate
- Delphi REF-op (covid prod)
- Spread scale: most recent starts, learned + sharing
- Spread scale: most recent starts, equal
- Power scale: most recent starts, learned + sharing
- Power scale: most recent starts, equal
- Weekly starts: most recent, learned + sharing
- Fixed MultiQT (log1p tracker), adaptive rate
- Sqrt tracker, adaptive rate

Colours as in the trade-off plots of R/04.





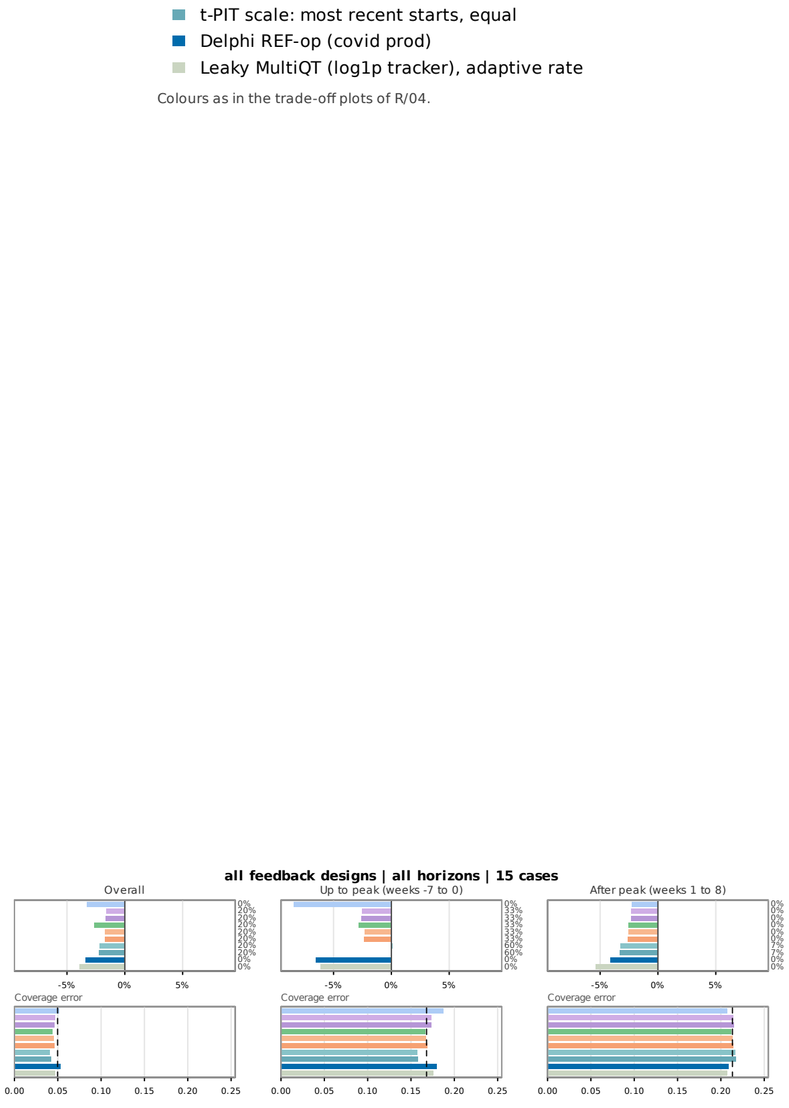
Loss bar: 100 * (1 - pooled loss / pooled loss of the base), pooled over the node's cases (location x horizon x feedback design), each case's mean loss weighted by its number of weeks; to the right: positive = smaller loss than the base.
Numbers right of each loss panel: share of the node's cases in which the method's loss is below the base's (nodes with at least 5 cases). Up to peak / after peak: the weeks around every local peak of the latest counts.
Below each loss panel: coverage error, the mean over the quantile levels of |coverage - alpha| over the weeks of the window, averaged over the cases of the node weighted by their weeks; dashed: the base forecasts.

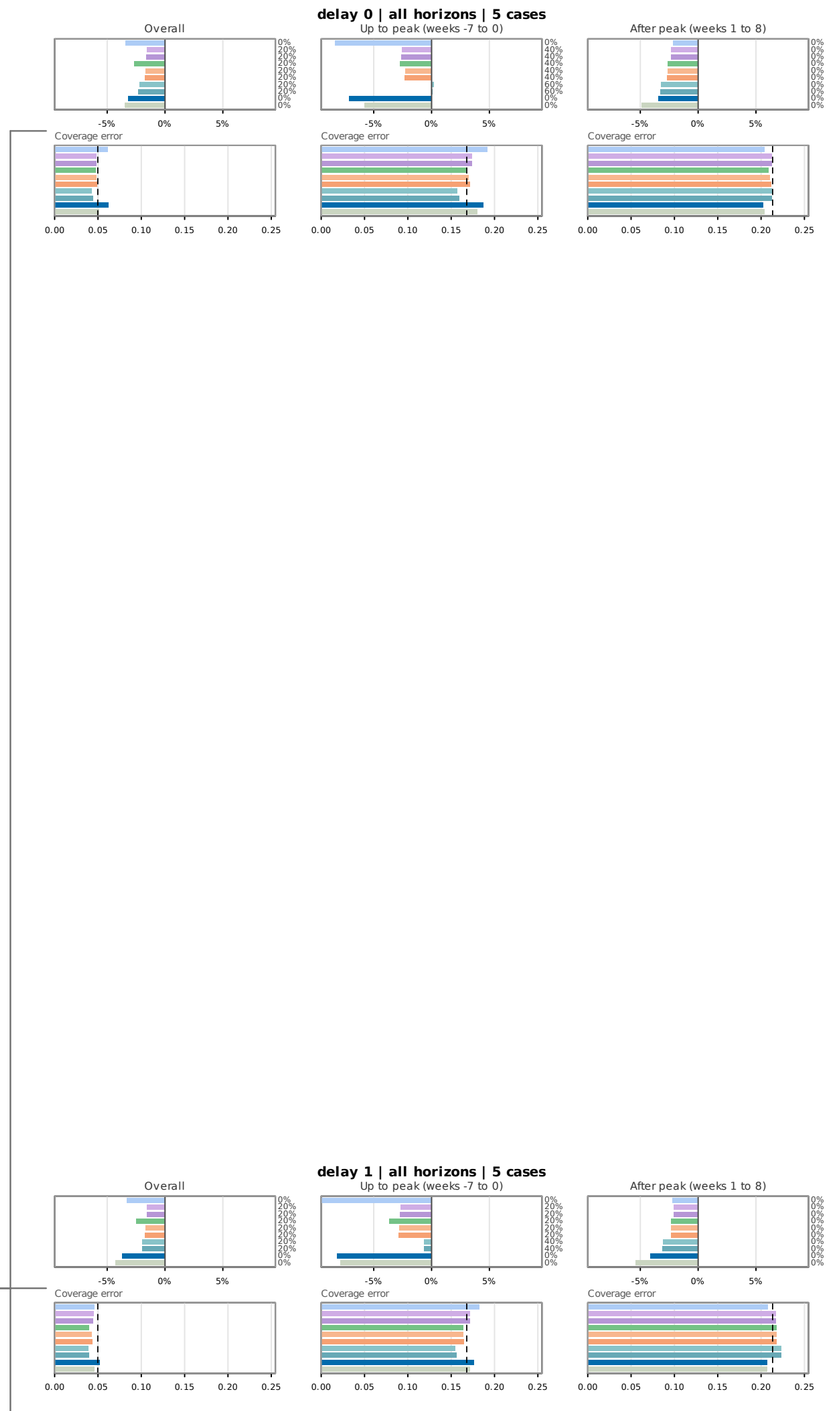
covid | Pennsylvania (42) | evaluation: Jun 2025 to Sep 2026 | loss in counts per 100k residents

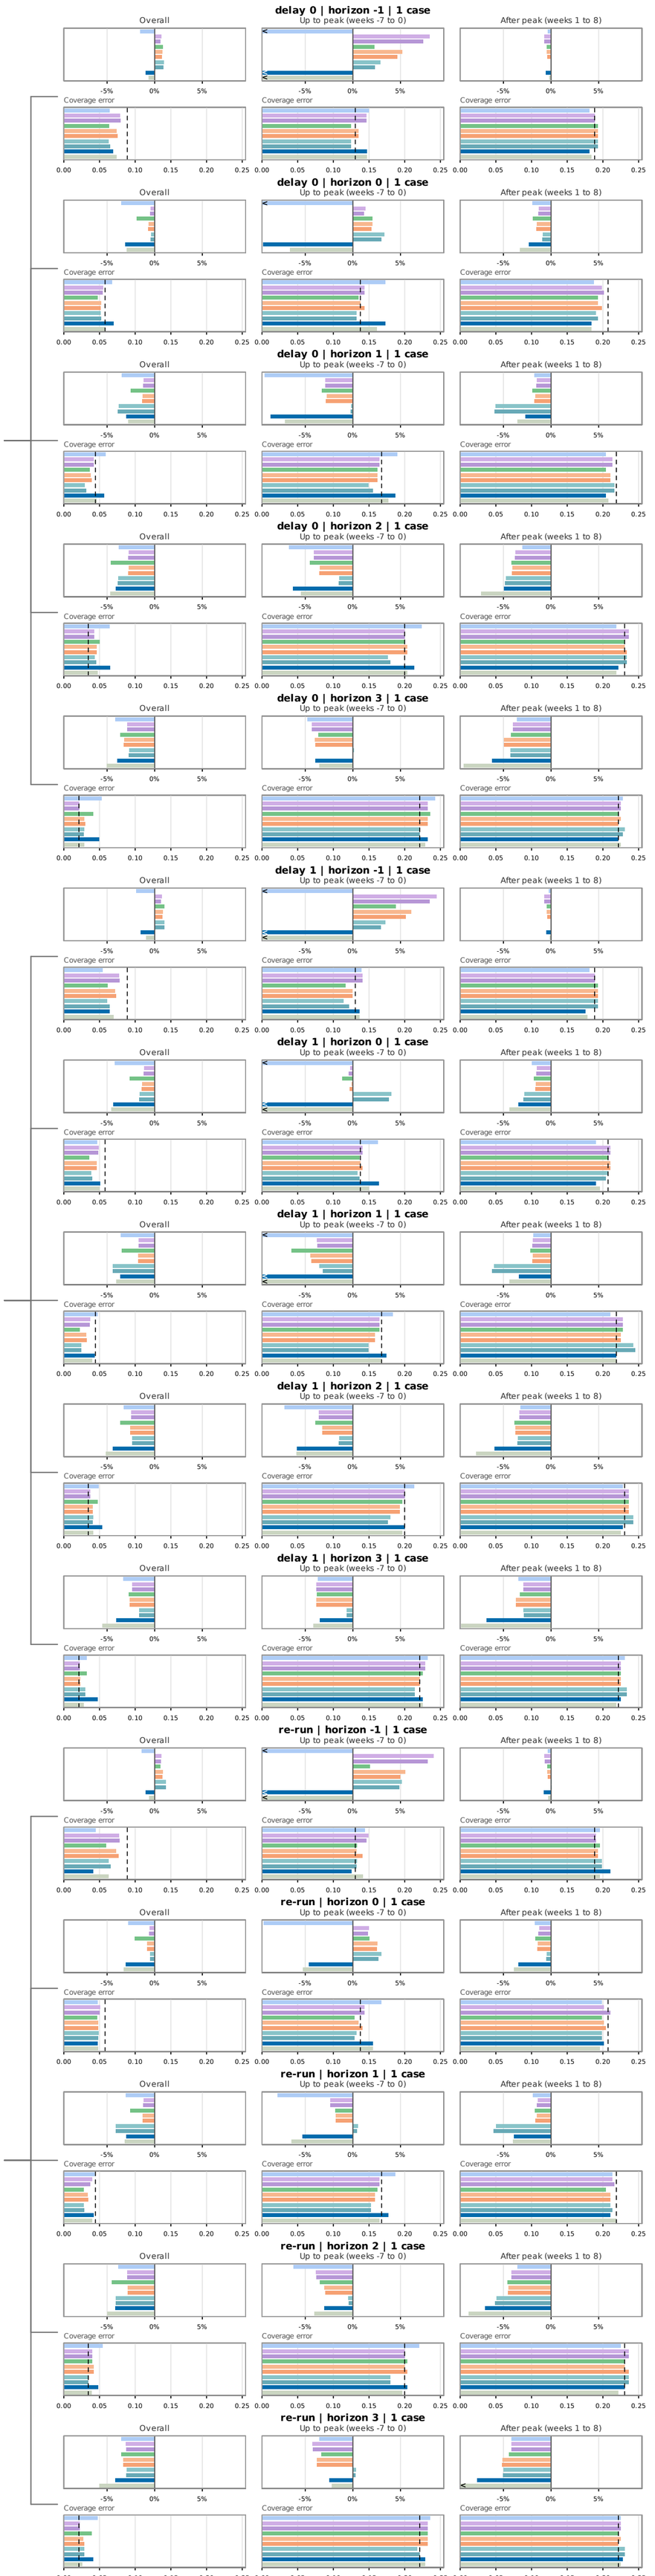
Bars, top to bottom: the 10 best on this page (quantile loss after peak, all designs and horizons)

- Leaky MultiQT, adaptive rate
- Spread scale: most recent starts, learned + sharing
- Spread scale: most recent starts, equal
- Weekly starts: most recent, learned + sharing
- Power scale: most recent starts, learned + sharing
- Power scale: most recent starts, equal
- t-PIT scale: most recent starts, learned + sharing
- t-PIT scale: most recent starts, equal
- Delphi REF-op (covid prod)
- Leaky MultiQT (log1p tracker), adaptive rate

Colours as in the trade-off plots of R/04.







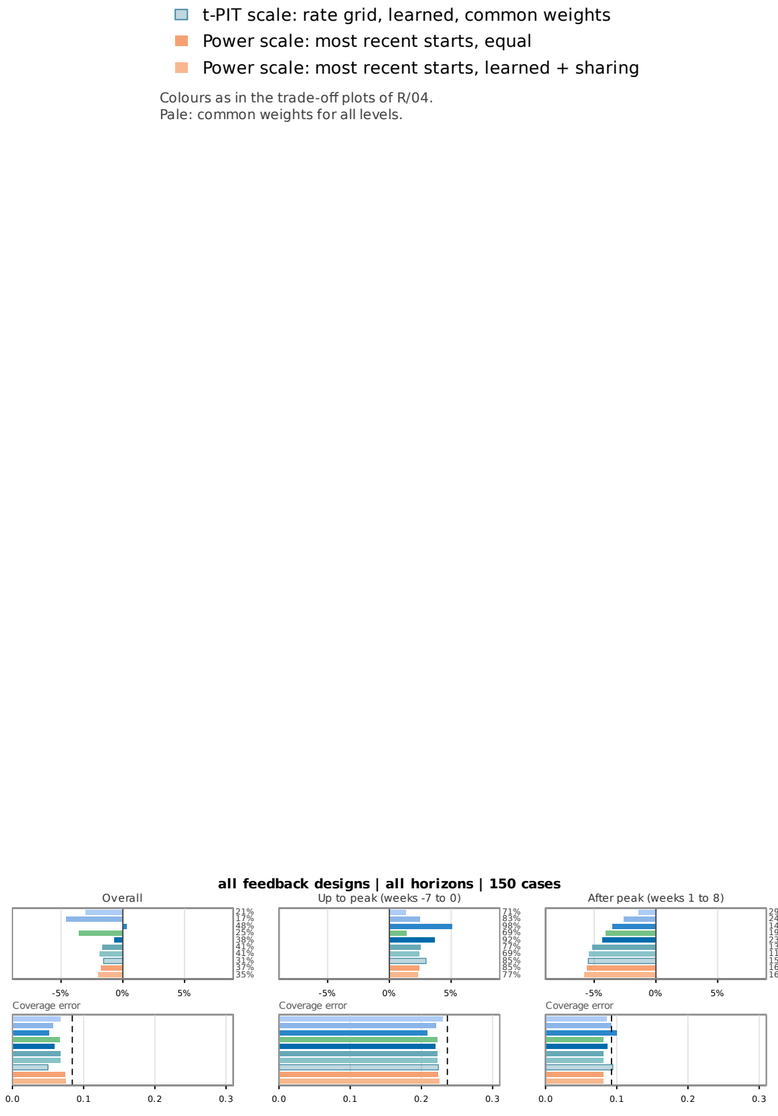
Loss bar: 100 * (1 - pooled loss / pooled loss of the base), pooled over the node's cases (location x horizon x feedback design), each case's mean loss weighted by its number of weeks; to the right: positive = smaller loss than the base.
Numbers right of each loss panel: share of the node's cases in which the method's loss is below the base's (nodes with at least 5 cases). Up to peak / after peak: the weeks around every local peak of the latest counts.
Below each loss panel: coverage error, the mean over the quantile levels of |coverage - alpha| over the weeks of the window, averaged over the cases of the node weighted by their weeks; dashed: the base forecasts.

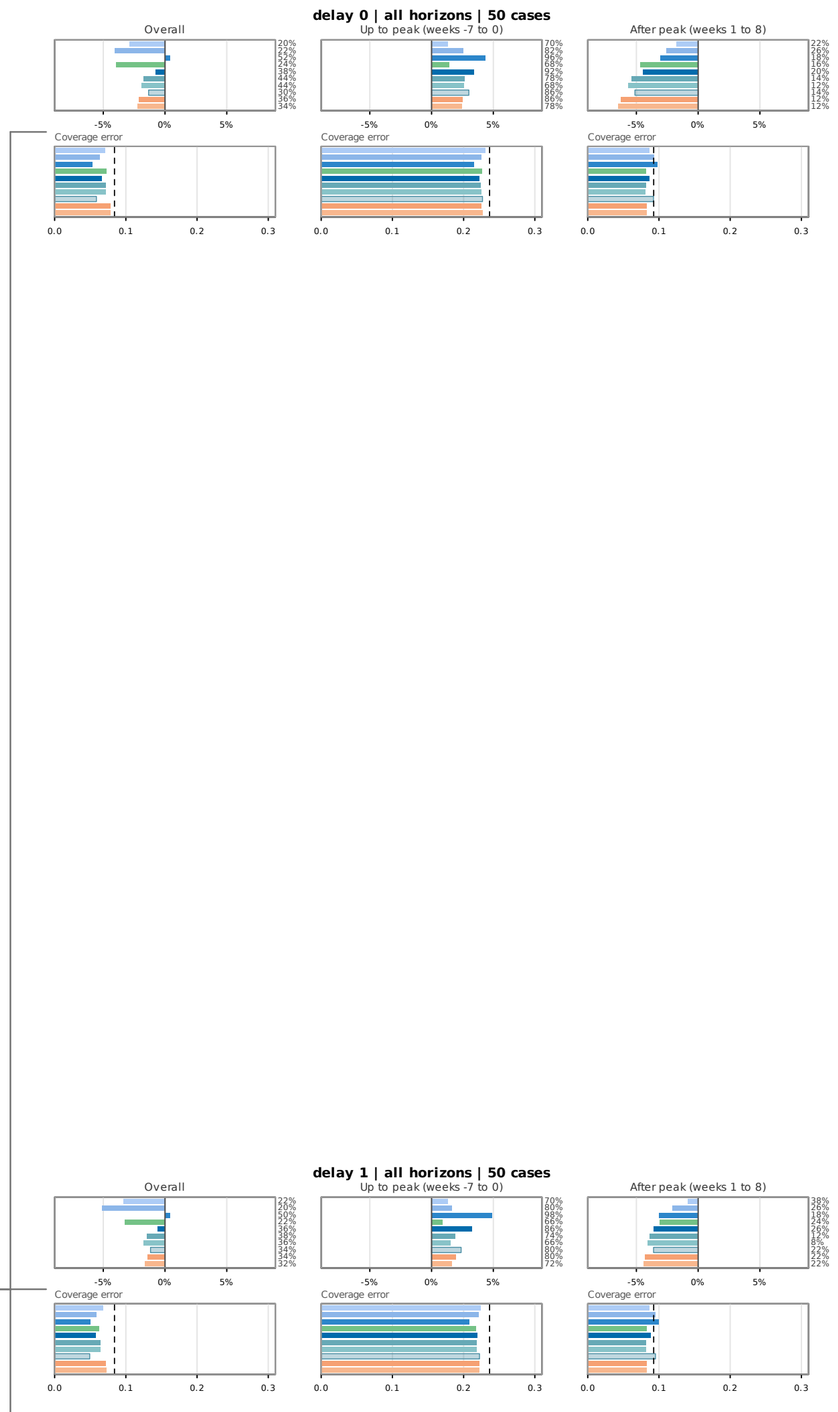
flu | the 10 largest locations without Florida pooled | evaluation: Jun 2025 to Sep 2026 | loss in counts per 100k residents

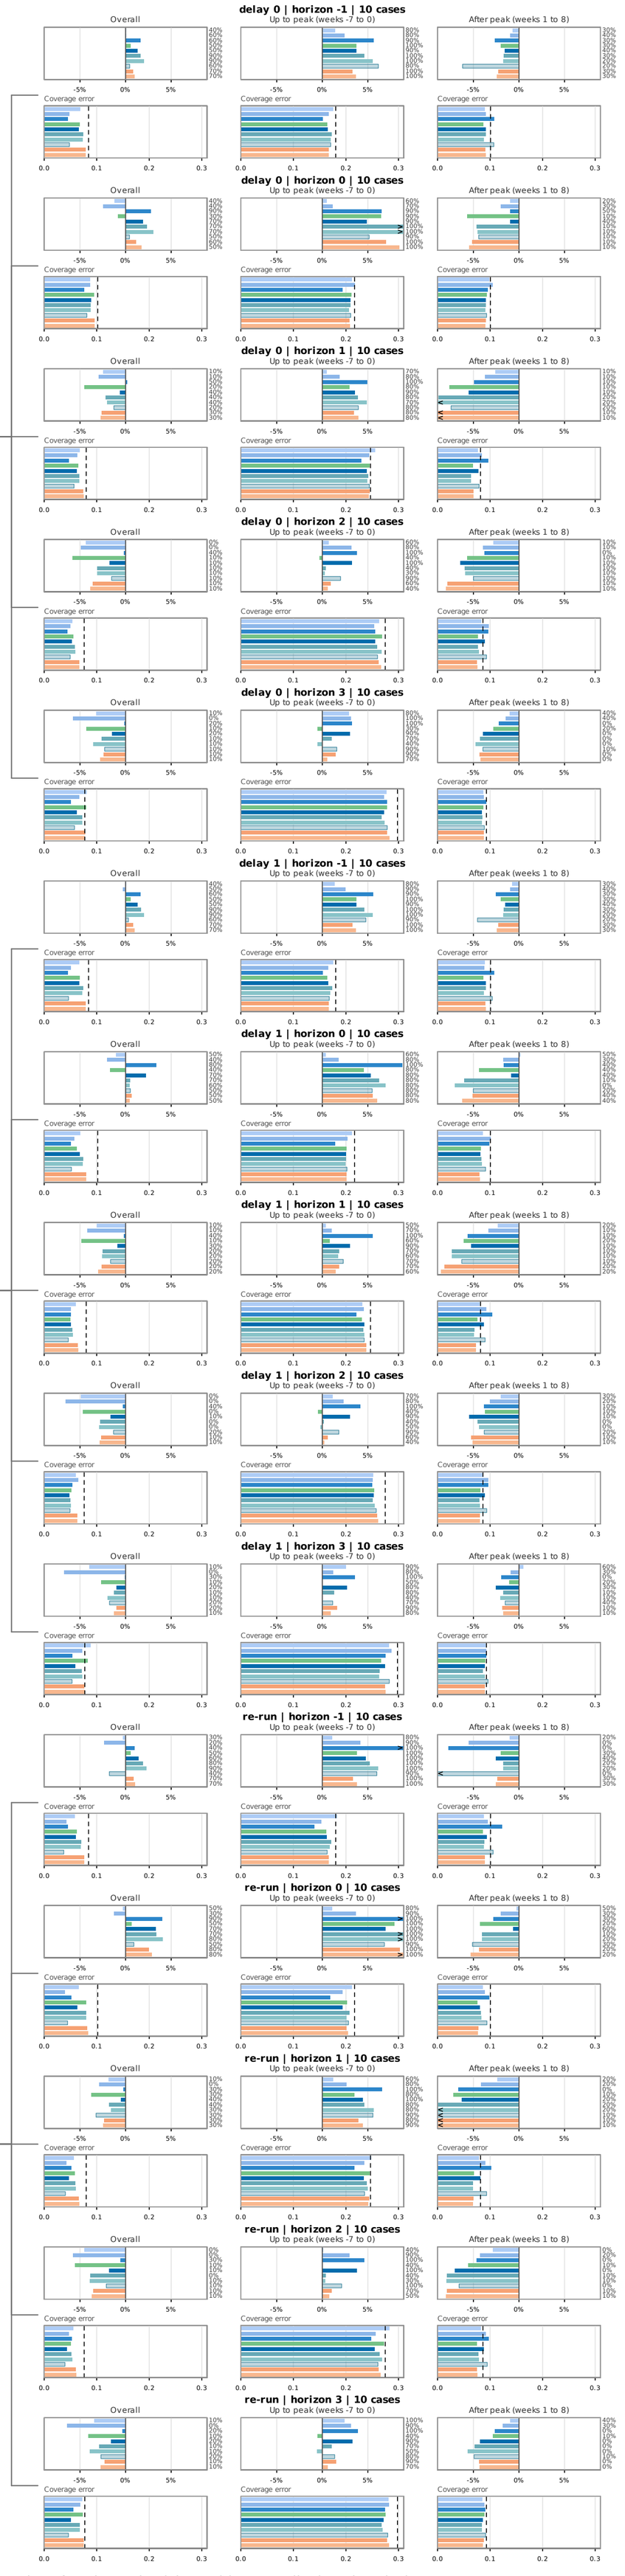
Bars, top to bottom: the 10 best on this page (quantile loss after peak, all designs and horizons)

- Leaky MultiQT, adaptive rate
- Fixed MultiQT, adaptive rate
- Delphi sqrt 0.018 (flu prod)
- Weekly starts: most recent, learned + sharing
- Delphi REF-op (covid prod)
- t-PIT scale: most recent starts, equal
- t-PIT scale: most recent starts, learned + sharing
- t-PIT scale: rate grid, learned, common weights
- Power scale: most recent starts, equal
- Power scale: most recent starts, learned + sharing

Colours as in the trade-off plots of R/04.
Pale: common weights for all levels.







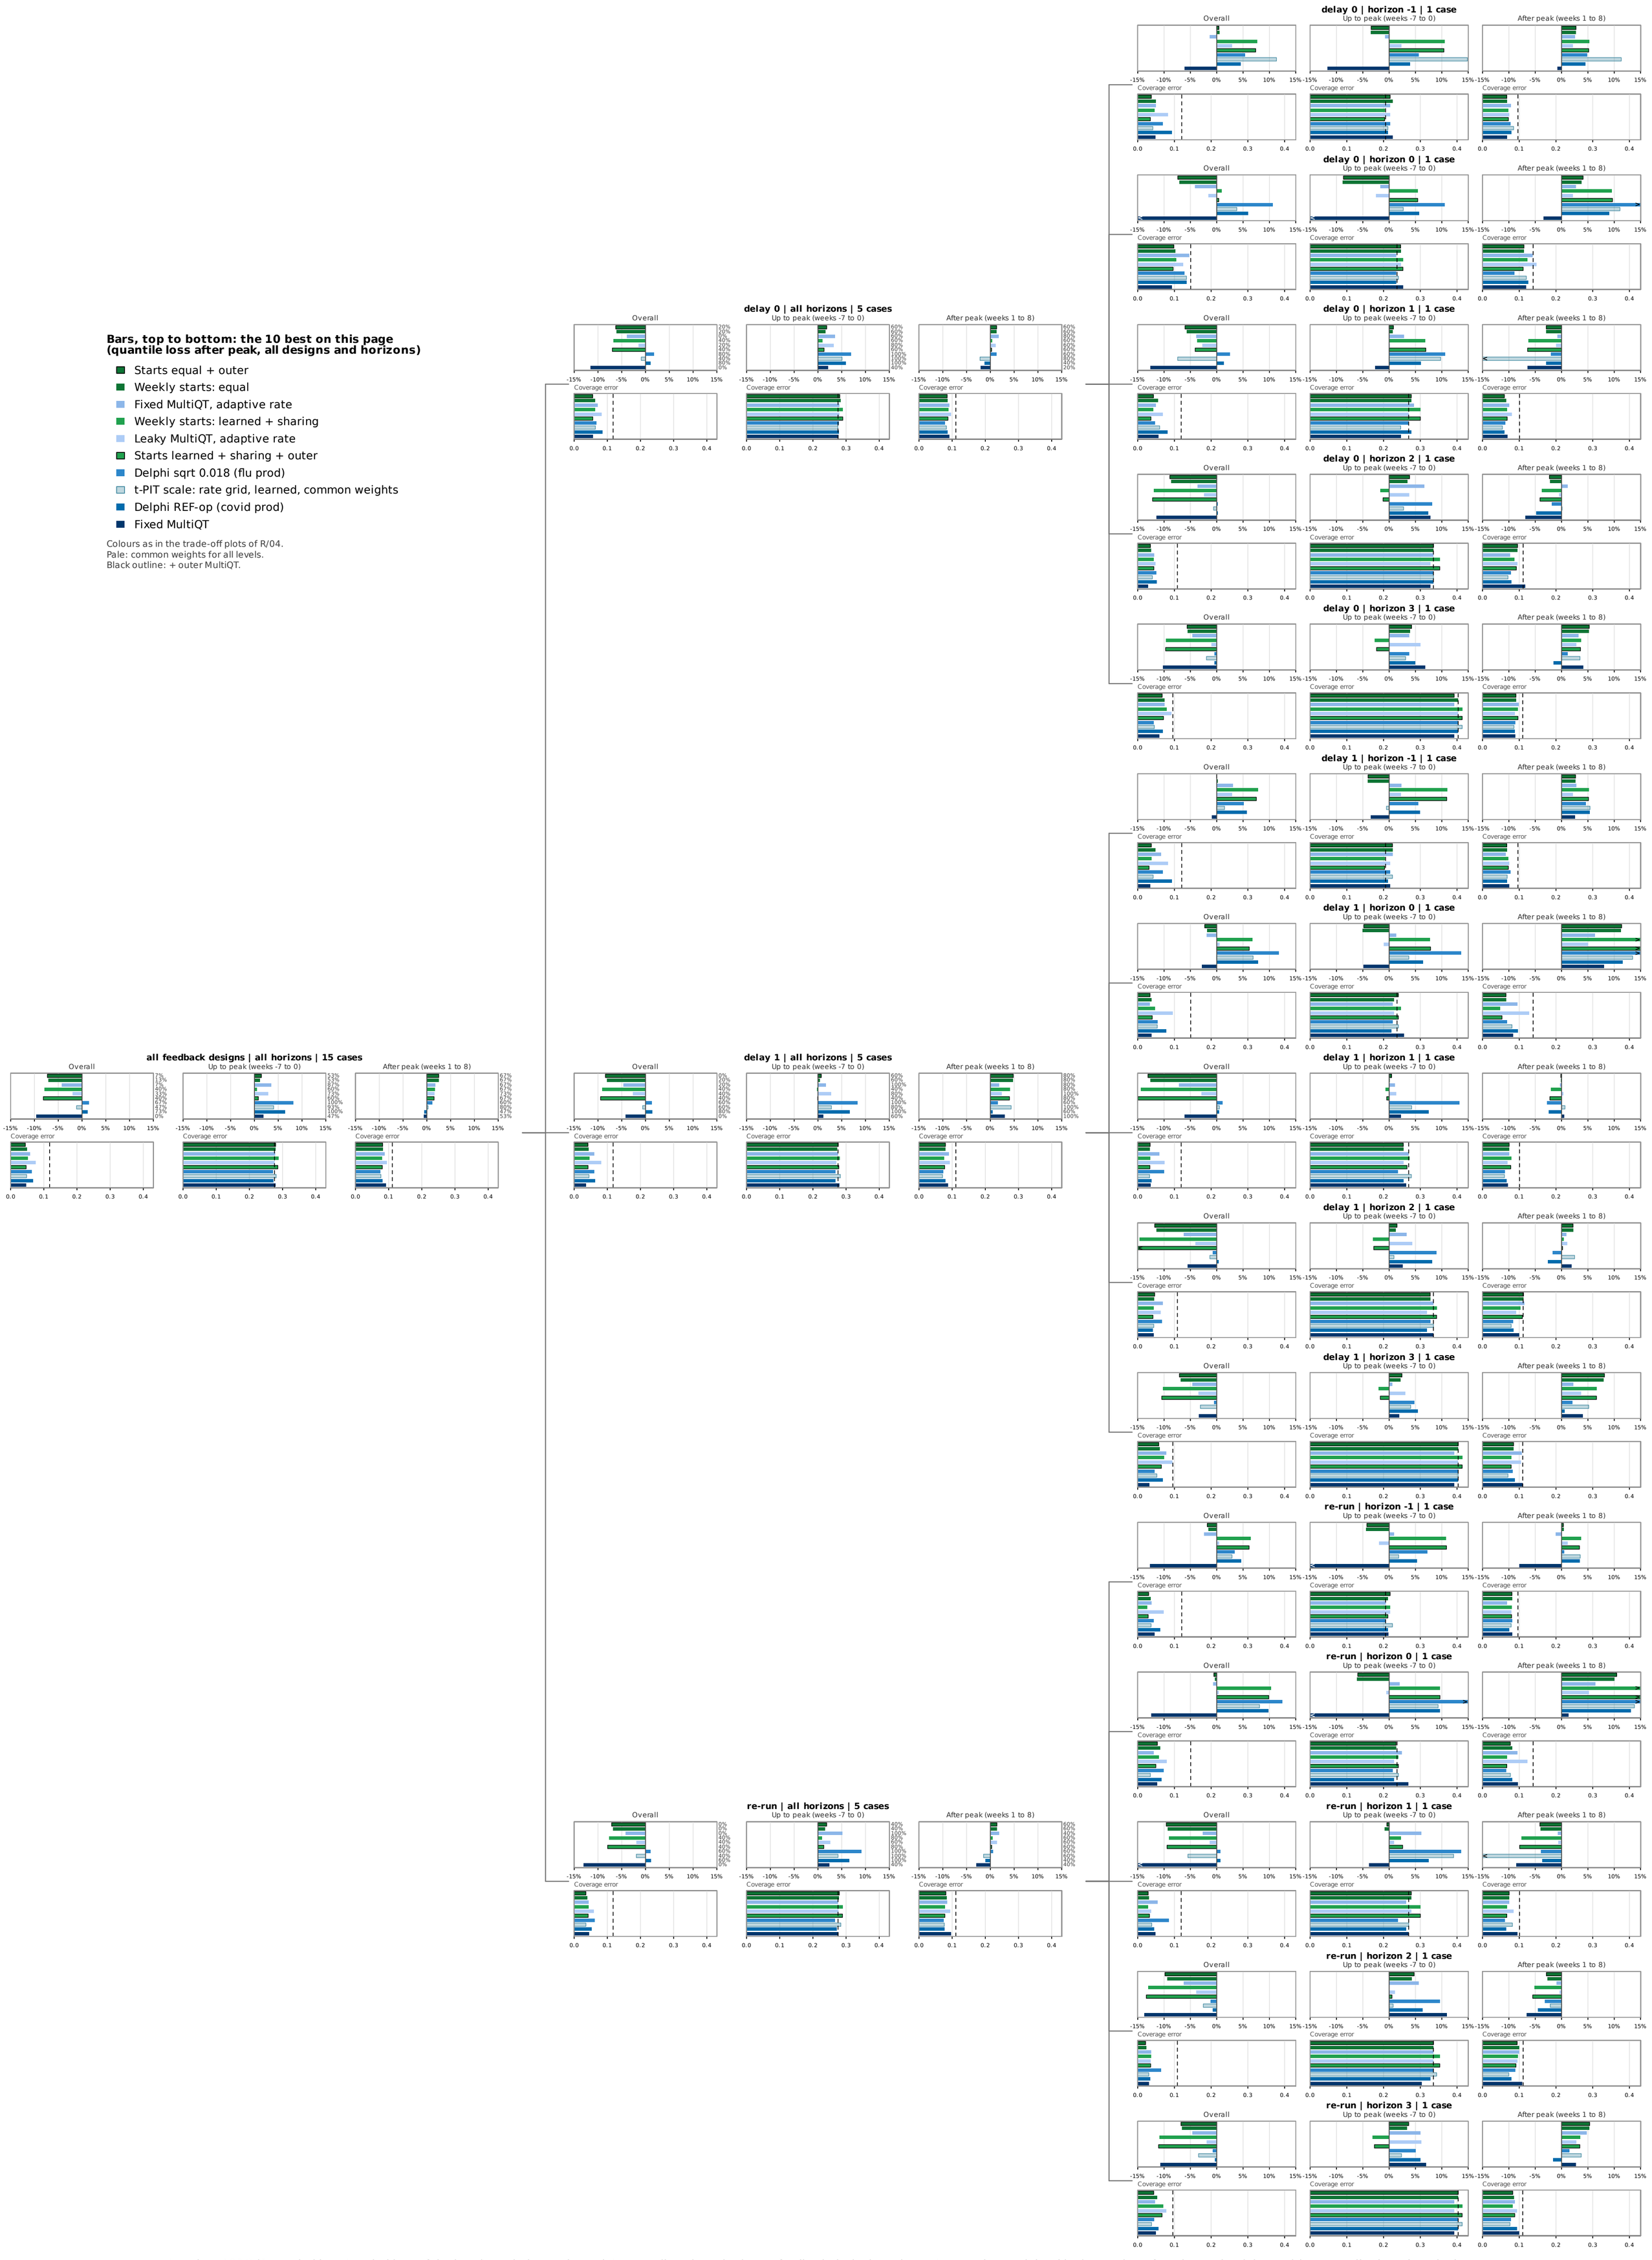
Loss bar: 100 * (1 - pooled loss / pooled loss of the base), pooled over the node's cases (location x horizon x feedback design), each case's mean loss weighted by its number of weeks; to the right: positive = smaller loss than the base.
Numbers right of each loss panel: share of the node's cases in which the method's loss is below the base's (nodes with at least 5 cases). Up to peak / after peak: the weeks around every local peak of the latest counts.
Below each loss panel: coverage error, the mean over the quantile levels of |coverage - alpha| over the weeks of the window, averaged over the cases of the node weighted by their weeks; dashed: the base forecasts.

flu | United States (US) | evaluation: Jun 2025 to Sep 2026 | loss in counts per 100k residents

Bars, top to bottom: the 10 best on this page (quantile loss after peak, all designs and horizons)

- Starts equal + outer
- Weekly starts: equal
- Fixed MultiQT, adaptive rate
- Weekly starts: learned + sharing
- Leaky MultiQT, adaptive rate
- Starts learned + sharing + outer
- Delphi sqrt 0.018 (flu prod)
- t-PIT scale: rate grid, learned, common weights
- Delphi REF-op (covid prod)
- Fixed MultiQT

Colours as in the trade-off plots of R/04.
Pale: common weights for all levels.
Black outline: + outer MultiQT.



Loss bar: 100 * (1 - pooled loss / pooled loss of the base), pooled over the node's cases (location x horizon x feedback design), each case's mean loss weighted by its number of weeks; to the right: positive = smaller loss than the base.
Numbers right of each loss panel: share of the node's cases in which the method's loss is below the base's (nodes with at least 5 cases). Up to peak / after peak: the weeks around every local peak of the latest counts.
Below each loss panel: coverage error, the mean over the quantile levels of |coverage - alpha| over the weeks of the window, averaged over the cases of the node weighted by their weeks; dashed: the base forecasts.

flu | California (06) | evaluation: Jun 2025 to Sep 2026 | loss in counts per 100k residents

Bars, top to bottom: the 10 best on this page (quantile loss after peak, all designs and horizons)

- Spread scale: most recent starts, learned + sharing
- Spread scale: most recent starts, equal
- Power scale: most recent starts, learned + sharing
- Power scale: most recent starts, equal
- Weekly starts: most recent, learned + sharing
- t-PIT scale: most recent starts, learned + sharing
- t-PIT scale: most recent starts, equal
- Leaky MultiQT, adaptive rate
- Leaky MultiQT (log1p tracker), adaptive rate
- Delphi REF-op (covid prod)

Colours as in the trade-off plots of R/04.

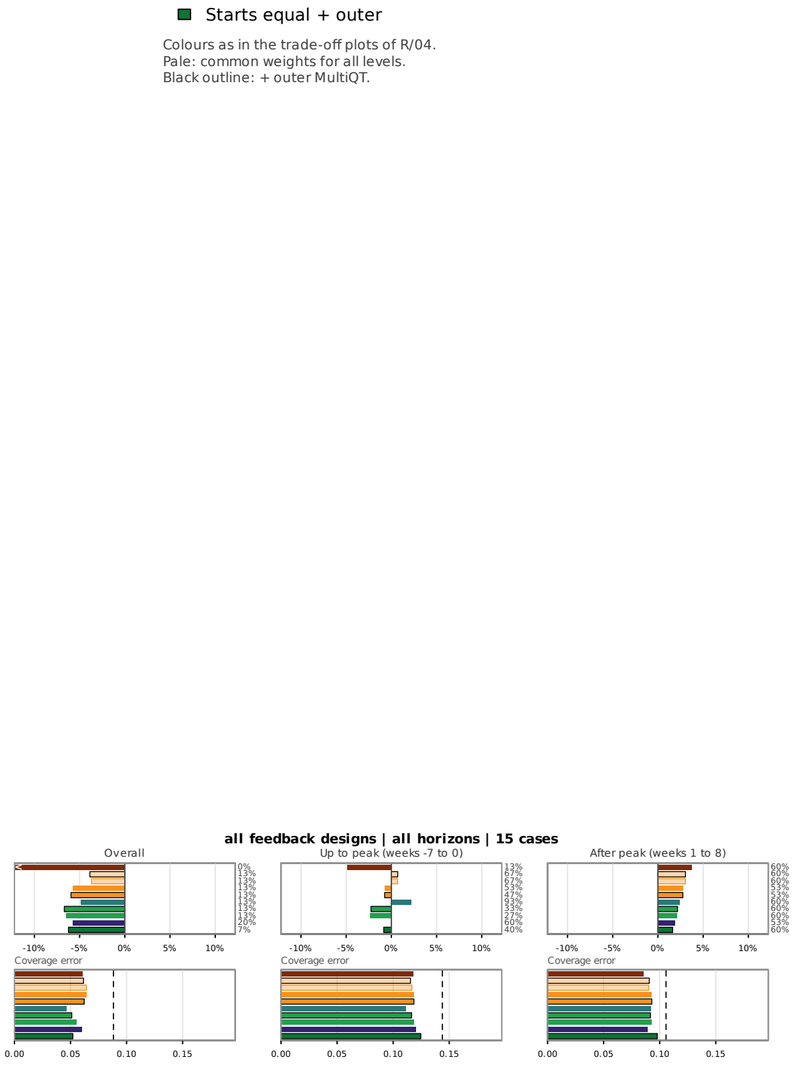
Loss bar: 100 * (1 - pooled loss / pooled loss of the base), pooled over the node's cases (location x horizon x feedback design), each case's mean loss weighted by its number of weeks; to the right: positive = smaller loss than the base.
Numbers right of each loss panel: share of the node's cases in which the method's loss is below the base's (nodes with at least 5 cases). Up to peak / after peak: the weeks around every local peak of the latest counts.
Below each loss panel: coverage error, the mean over the quantile levels of |coverage - alpha| over the weeks of the window, averaged over the cases of the node weighted by their weeks; dashed: the base forecasts.

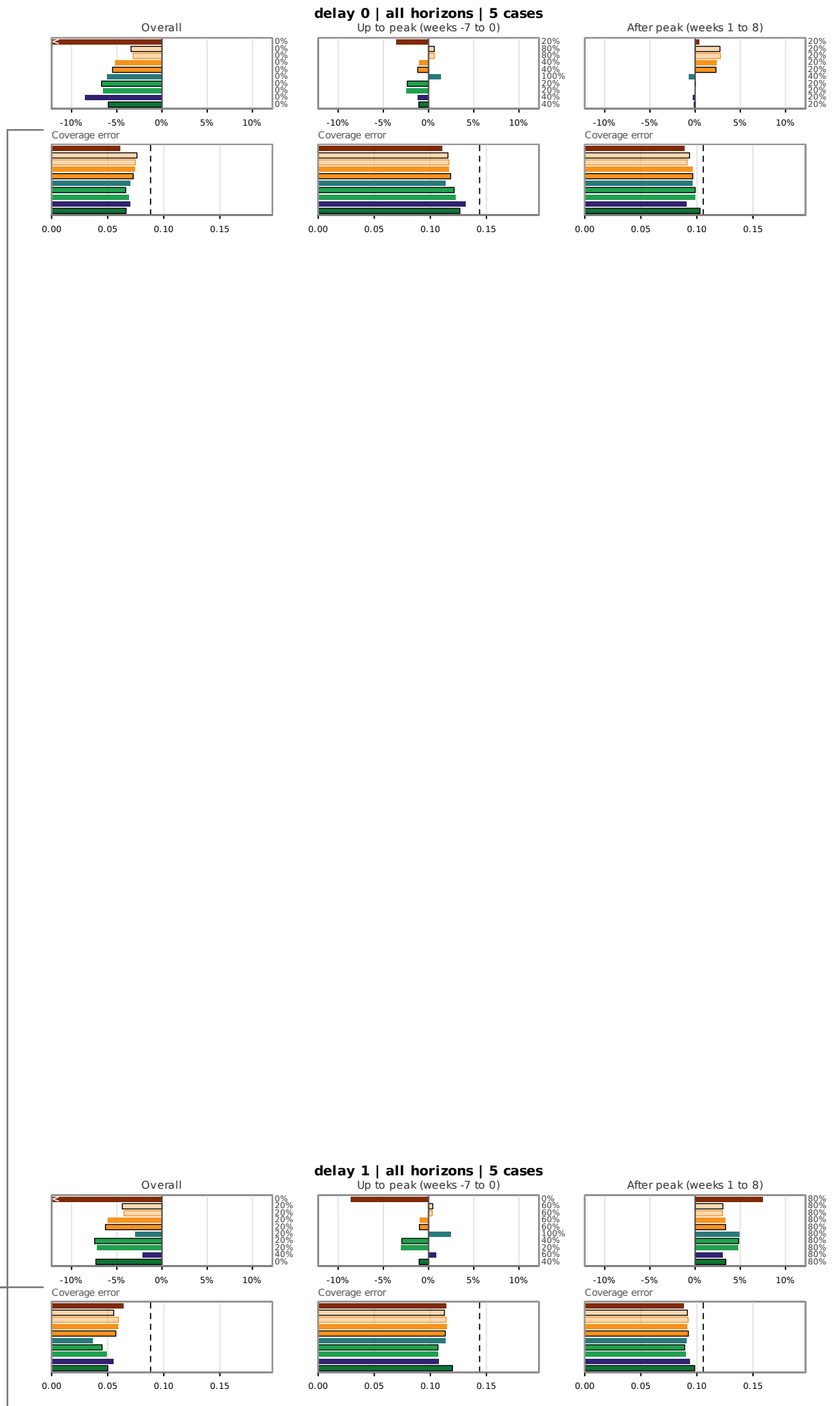
flu | Texas (48) | evaluation: Jun 2025 to Sep 2026 | loss in counts per 100k residents

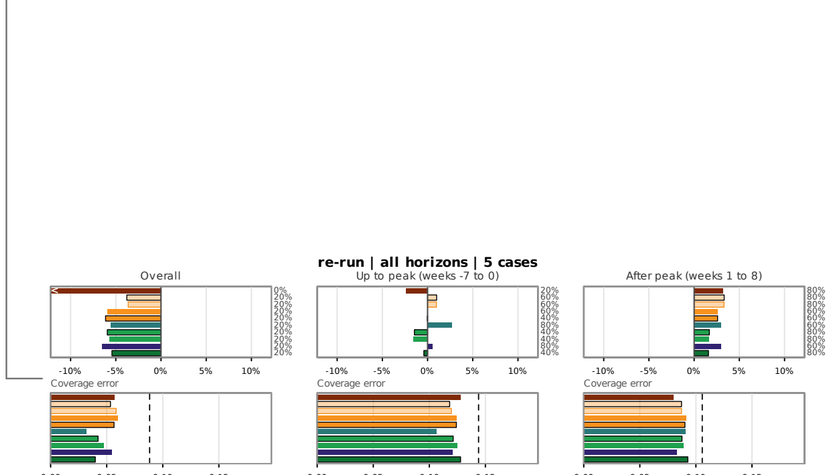
Bars, top to bottom: the 10 best on this page (quantile loss after peak, all designs and horizons)

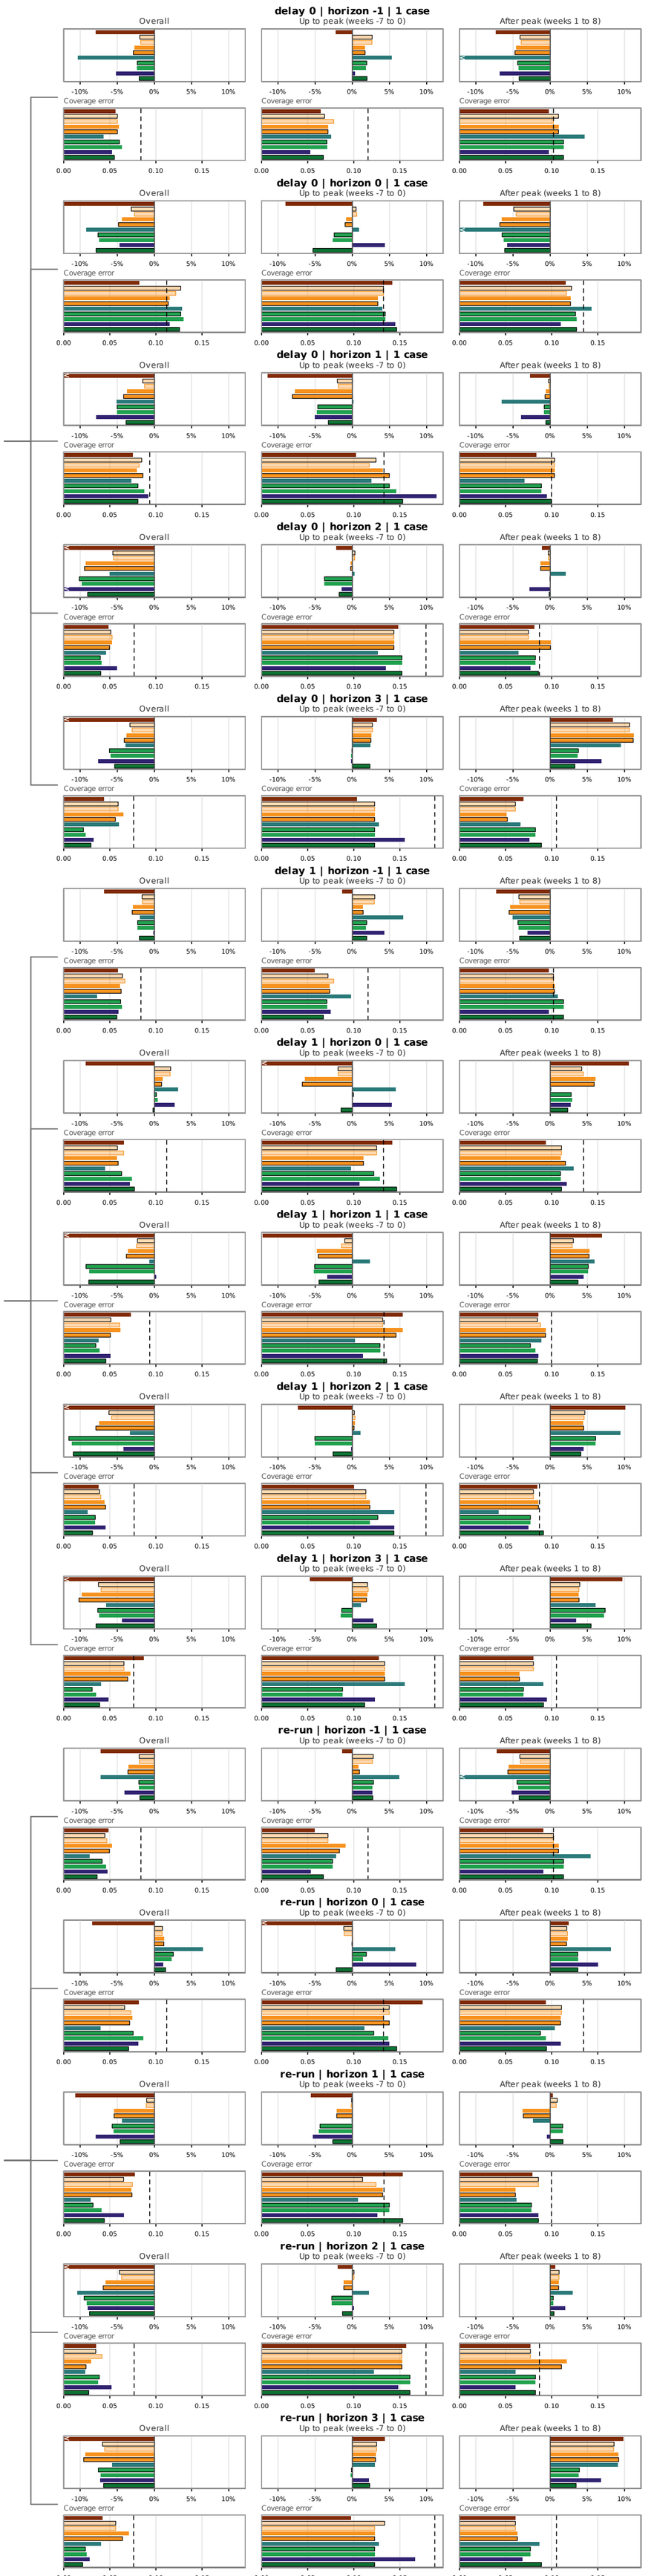
- Rate grid: equal
- Rate grid: sharing, common weights + outer
- Rate grid: learned + sharing, common weights
- Rate grid: learned + sharing
- Rate grid: sharing + outer
- Studentized full base (level RMS)
- Starts learned + sharing + outer
- Weekly starts: learned + sharing
- Score experts: equal weights
- Starts equal + outer

Colours as in the trade-off plots of R/04.
Pale: common weights for all levels.
Black outline: + outer MultiQT.









Loss bar: 100 * (1 - pooled loss / pooled loss of the base), pooled over the node's cases (location x horizon x feedback design), each case's mean loss weighted by its number of weeks; to the right: positive = smaller loss than the base.
Numbers right of each loss panel: share of the node's cases in which the method's loss is below the base's (nodes with at least 5 cases). Up to peak / after peak: the weeks around every local peak of the latest counts.
Below each loss panel: coverage error, the mean over the quantile levels of |coverage - alpha| over the weeks of the window, averaged over the cases of the node weighted by their weeks; dashed: the base forecasts.

flu | New York (36) | evaluation: Jun 2025 to Sep 2026 | loss in counts per 100k residents

Bars, top to bottom: the 10 best on this page (quantile loss after peak, all designs and horizons)

- Spread scale: most recent starts, learned + sharing
- Delphi sqrt 0.018 (flu prod)
- Spread scale: most recent starts, equal
- Leaky MultiQT, adaptive rate
- Weekly starts: most recent, learned + sharing
- Power scale: most recent starts, learned + sharing
- Power scale: most recent starts, equal
- Delphi REF-op (covid prod)
- t-PIT scale: most recent starts, learned + sharing
- Sqrt tracker, adaptive rate

Colours as in the trade-off plots of R/04.

Loss bar: 100 * (1 - pooled loss / pooled loss of the base), pooled over the node's cases (location x horizon x feedback design), each case's mean loss weighted by its number of weeks; to the right: positive = smaller loss than the base.
Numbers right of each loss panel: share of the node's cases in which the method's loss is below the base's (nodes with at least 5 cases). Up to peak / after peak: the weeks around every local peak of the latest counts.
Below each loss panel: coverage error, the mean over the quantile levels of |coverage - alpha| over the weeks of the window, averaged over the cases of the node weighted by their weeks; dashed: the base forecasts.

flu | Pennsylvania (42) | evaluation: Jun 2025 to Sep 2026 | loss in counts per 100k residents

Bars, top to bottom: the 10 best on this page (quantile loss after peak, all designs and horizons)

- Leaky MultiQT, adaptive rate
- Fixed MultiQT, adaptive rate
- Delphi sqrt 0.018 (flu prod)
- t-PIT scale: most recent starts, equal
- t-PIT scale: most recent starts, learned + sharing
- Rate grid: learned, common weights
- Weekly starts: most recent, learned + sharing
- Delphi REF-op (covid prod)
- Power scale: most recent starts, equal
- Power scale: most recent starts, learned + sharing

Colours as in the trade-off plots of R/04.
Pale: common weights for all levels.