Frozen rate-grid weights

covid | horizon 1 | United States (US) | evaluation: Jun 2025 to Sep 2026 | raw-count updates | feedback re-run

Equal weights | **Learned, no sharing** | **Learned + sharing**

alpha = 0.1 | alpha = 0.2 | alpha = 0.3 | alpha = 0.4 | alpha = 0.5 | alpha = 0.6 | alpha = 0.7 | alpha = 0.8 | alpha = 0.9 | All levels: weights

Forecast, alpha = 0.2 (counts per 100k residents)

Forecast, alpha = 0.5 (counts per 100k residents)

Forecast, alpha = 0.8 (counts per 100k residents)

All levels: forecast, alpha = 0.5 (counts per 100k residents)

Rate grid: equal

Rate grid: learned, no sharing

Rate grid: learned + sharing

Rate grid: learned, common weights

Rate grid: learned + sharing, common weights

Jan 2025 | Jul 2025 | Jan 2026 | Jul 2026

Experts: weights and forecasts

m=0 | m=0.01 | m=0.015 | m=0.02 | m=0.03 | m=0.05 | m=0.075 | m=0.1 | m=0.15 | m=0.2 | m=0.3 | m=0.45

Latest reported counts

Each expert, while weighted: its weight colour

Mixture: dashed, in the colour of its method

Target date | Weight at issuance for each target | Vertical dotted lines: peaks

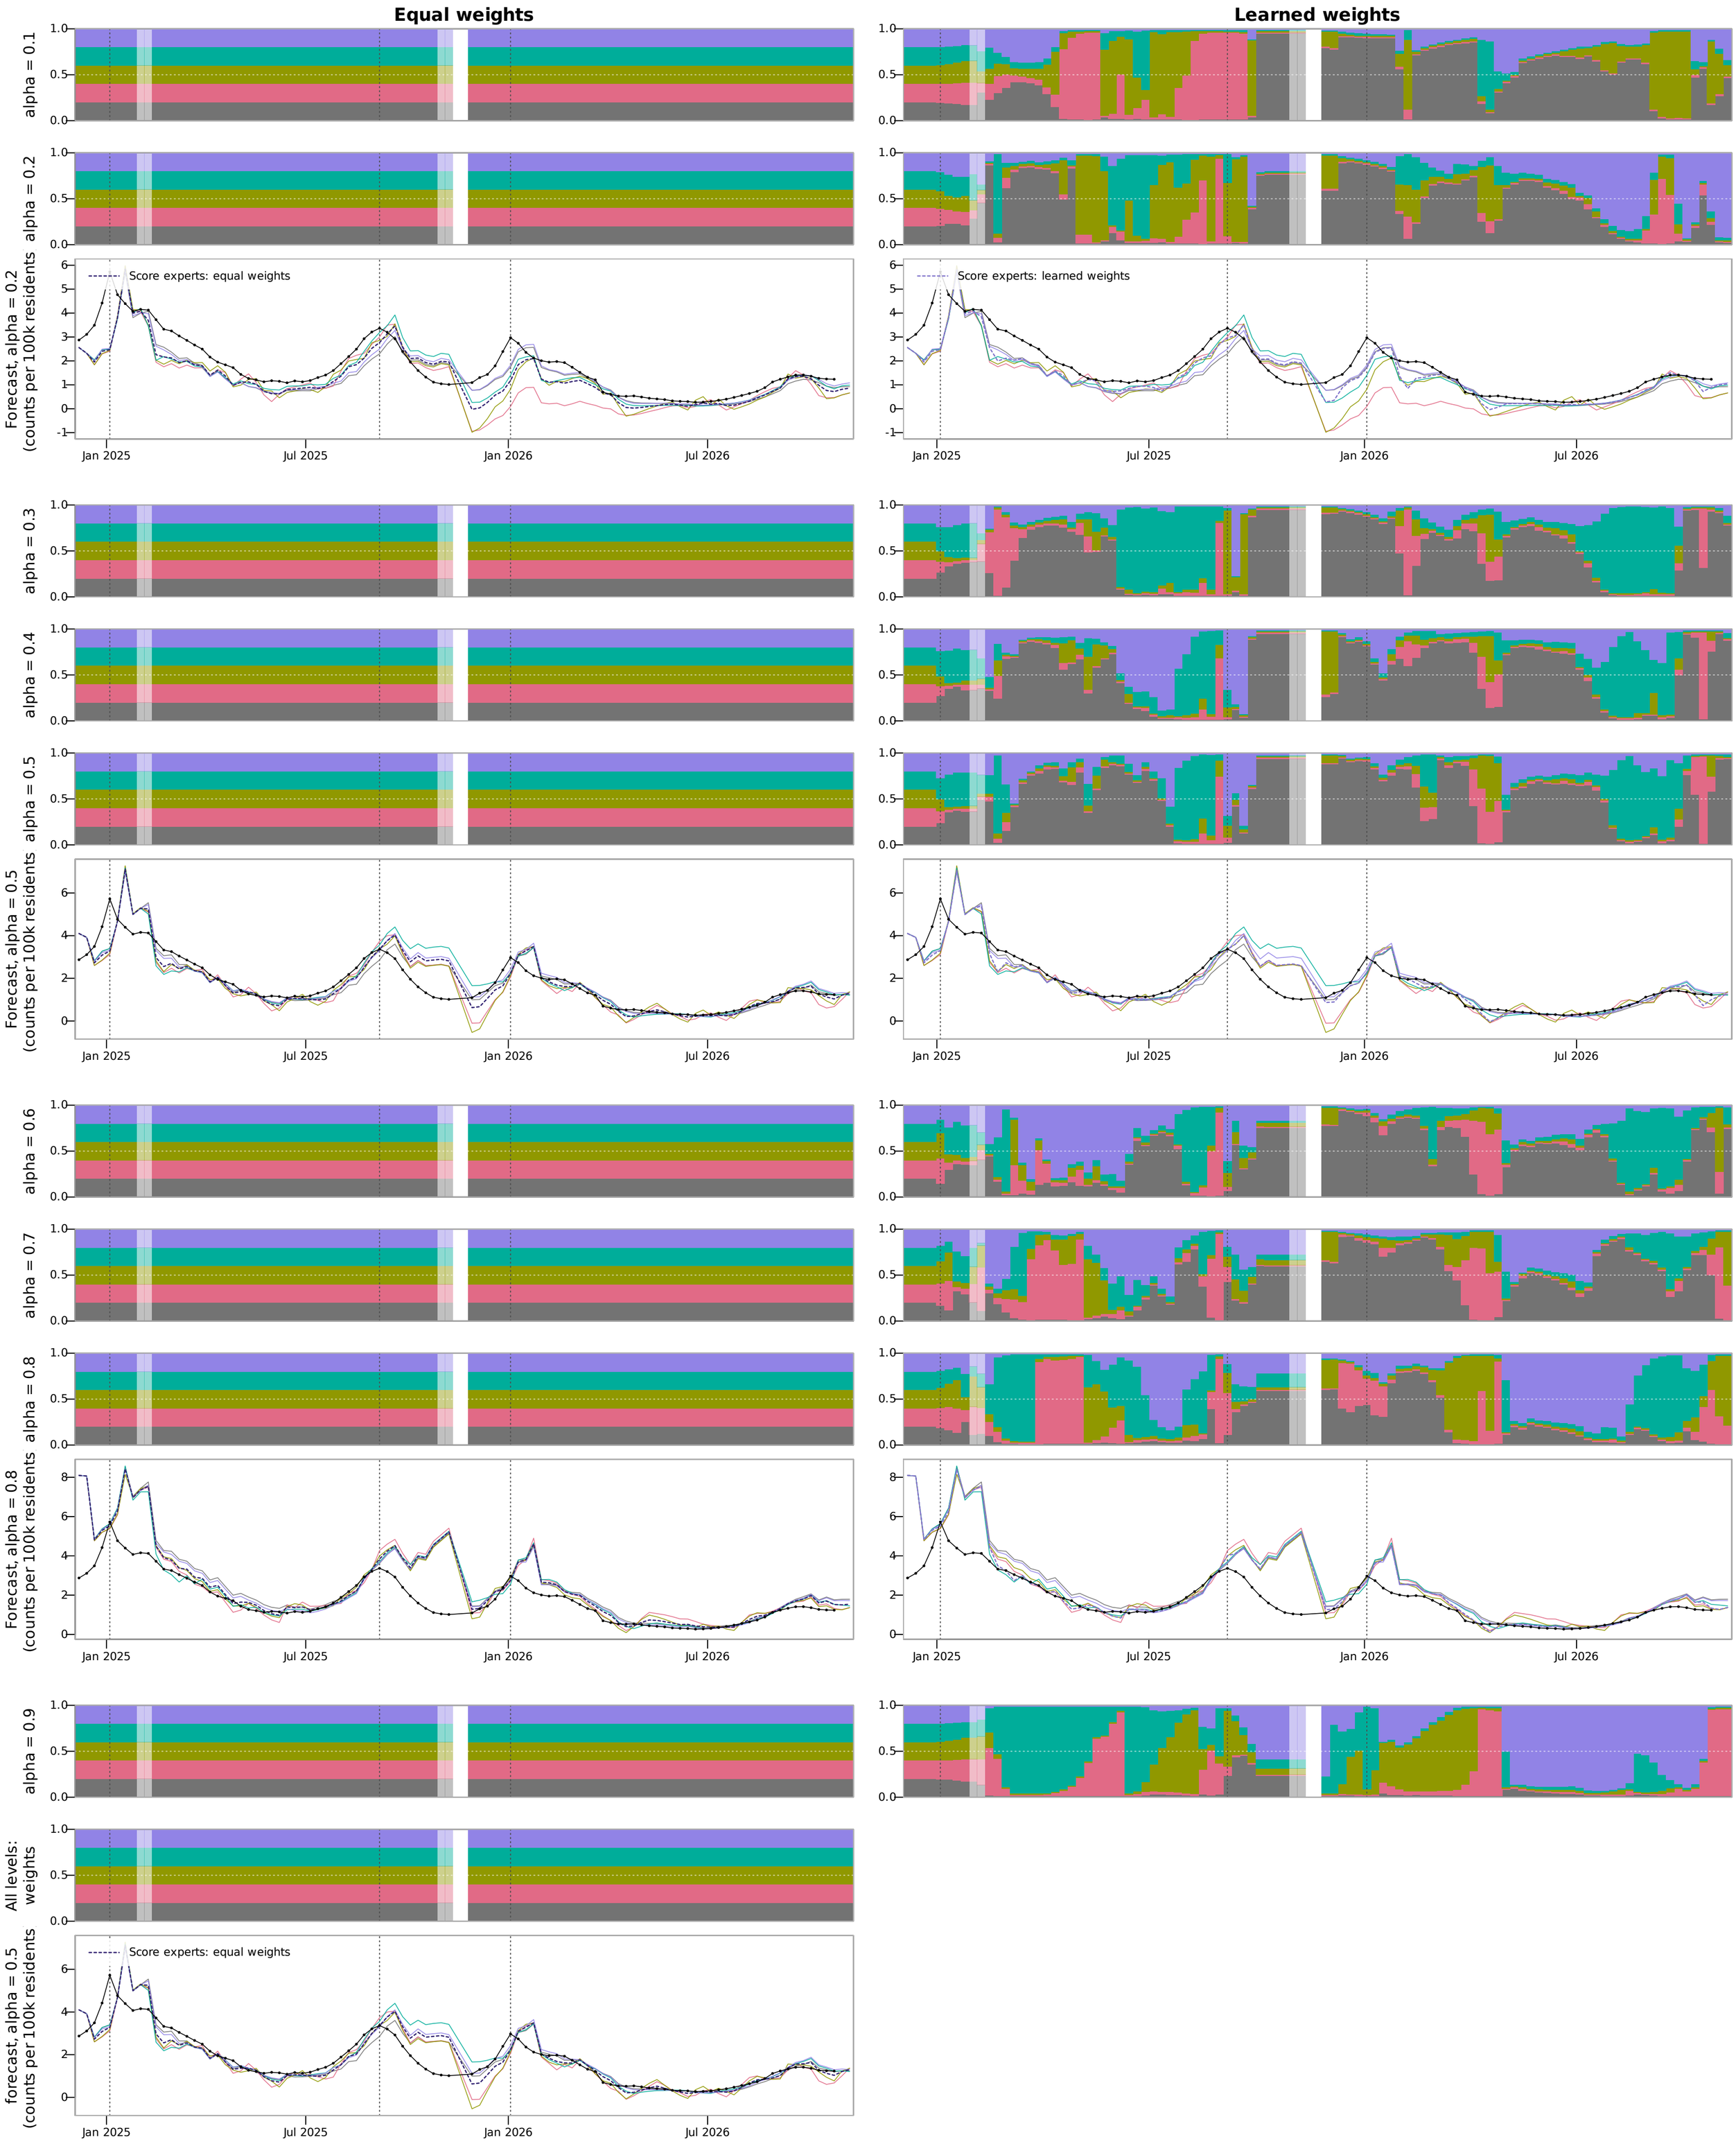
Rows alpha = 0.1-0.9: the weights of each level (one weight vector for all levels: the same at every level); below alpha = 0.2, 0.5 and 0.8 the forecasts at that level (after crossing repair, before any outer correction).

Rows 'All levels': the mixture with one weight vector for all levels (the column's mixture, or its common-weight variant), learned from the mean quantile loss over the levels (equal weights: not learned); its weights and its forecasts at alpha = 0.5.

Lighter bars: rounds without a CMU forecast, filled with the latest earlier CMU forecast of the same target week.

Score-based expert weights

covid | horizon 1 | United States (US) | evaluation: Jun 2025 to Sep 2026 | raw-count updates | feedback re-run



Experts: weights and forecasts

- Base
- Fixed MultiQT
- Leaky MultiQT
- Studentized full base (level RMS)
- Delphi sqrt 0.018 (flu prod)

- Latest reported counts
- Each expert, while weighted: its weight colour
- Mixture: dashed, in the colour of its method

Target date | Weight at issuance for each target | Vertical dotted lines: peaks

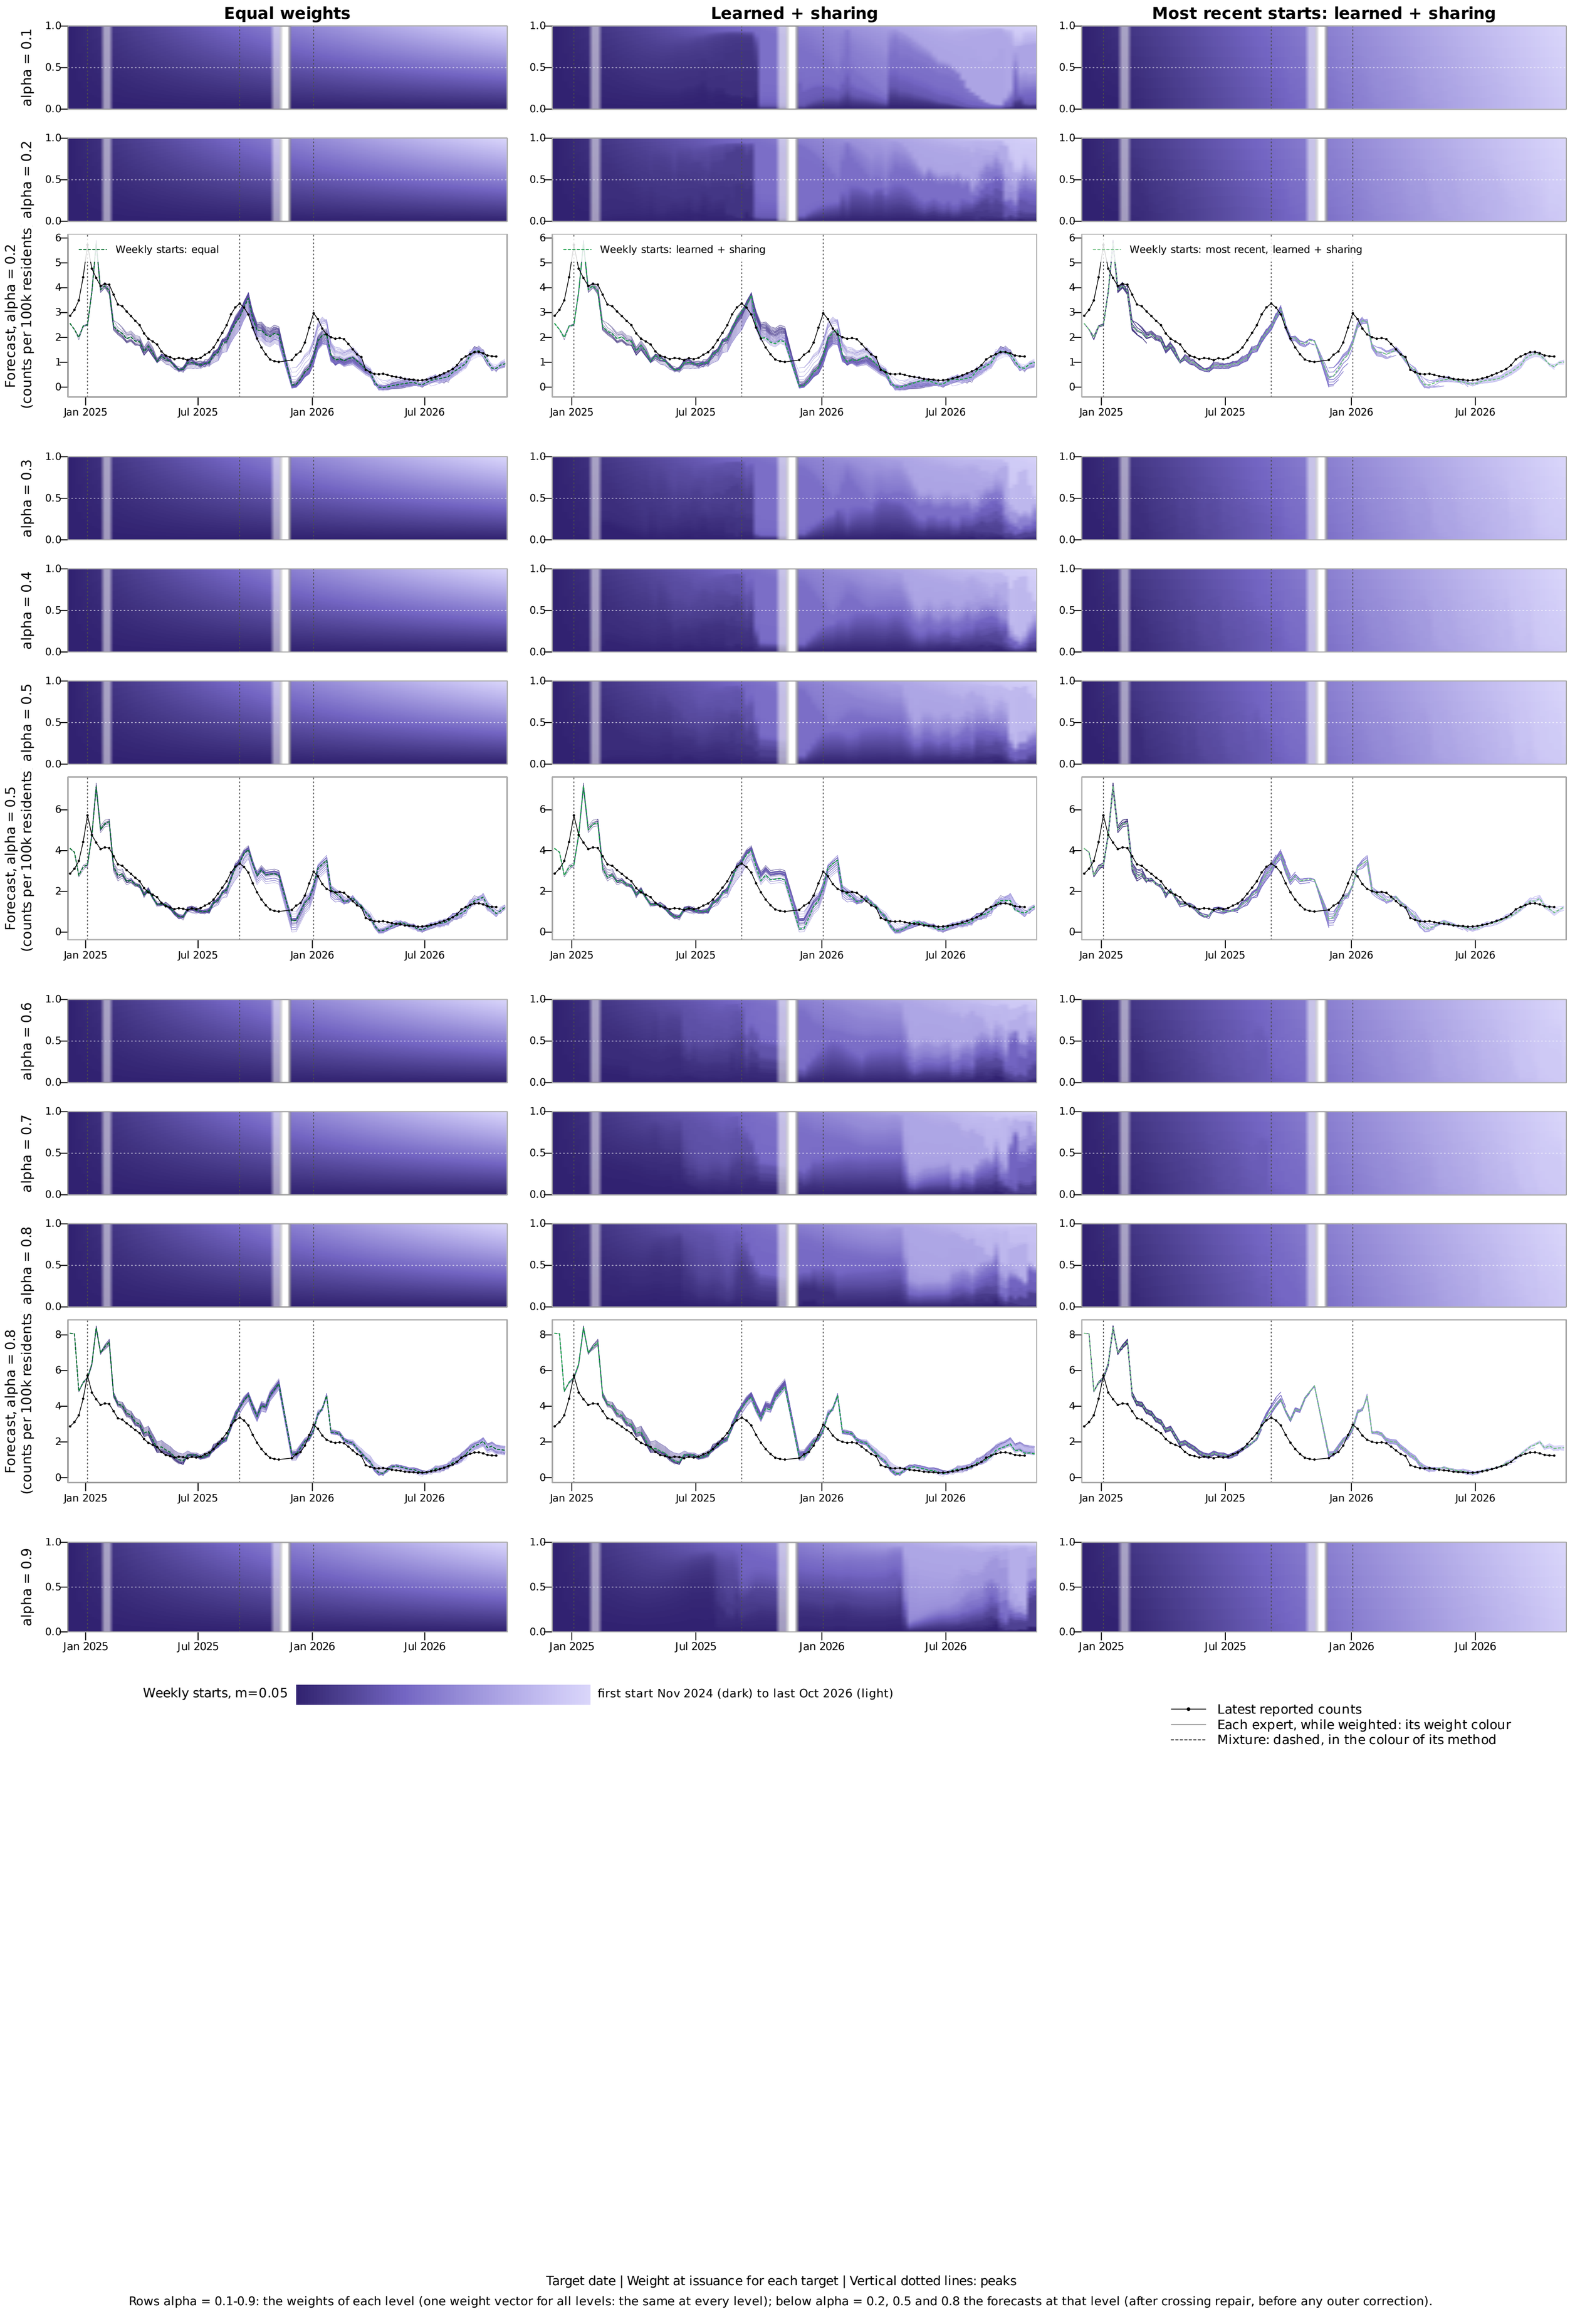
Rows alpha = 0.1-0.9: the weights of each level (one weight vector for all levels: the same at every level); below alpha = 0.2, 0.5 and 0.8 the forecasts at that level (after crossing repair, before any outer correction).

Rows 'All levels': the mixture with one weight vector for all levels (the column's mixture, or its common-weight variant), learned from the mean quantile loss over the levels (equal weights: not learned); its weights and its forecasts at alpha = 0.5.

Lighter bars: rounds without a CMU forecast, filled with the latest earlier CMU forecast of the same target week.

Weekly-start weights

covid | horizon 1 | United States (US) | evaluation: Jun 2025 to Sep 2026 | raw-count updates | feedback re-run



Target date | Weight at issuance for each target | Vertical dotted lines: peaks

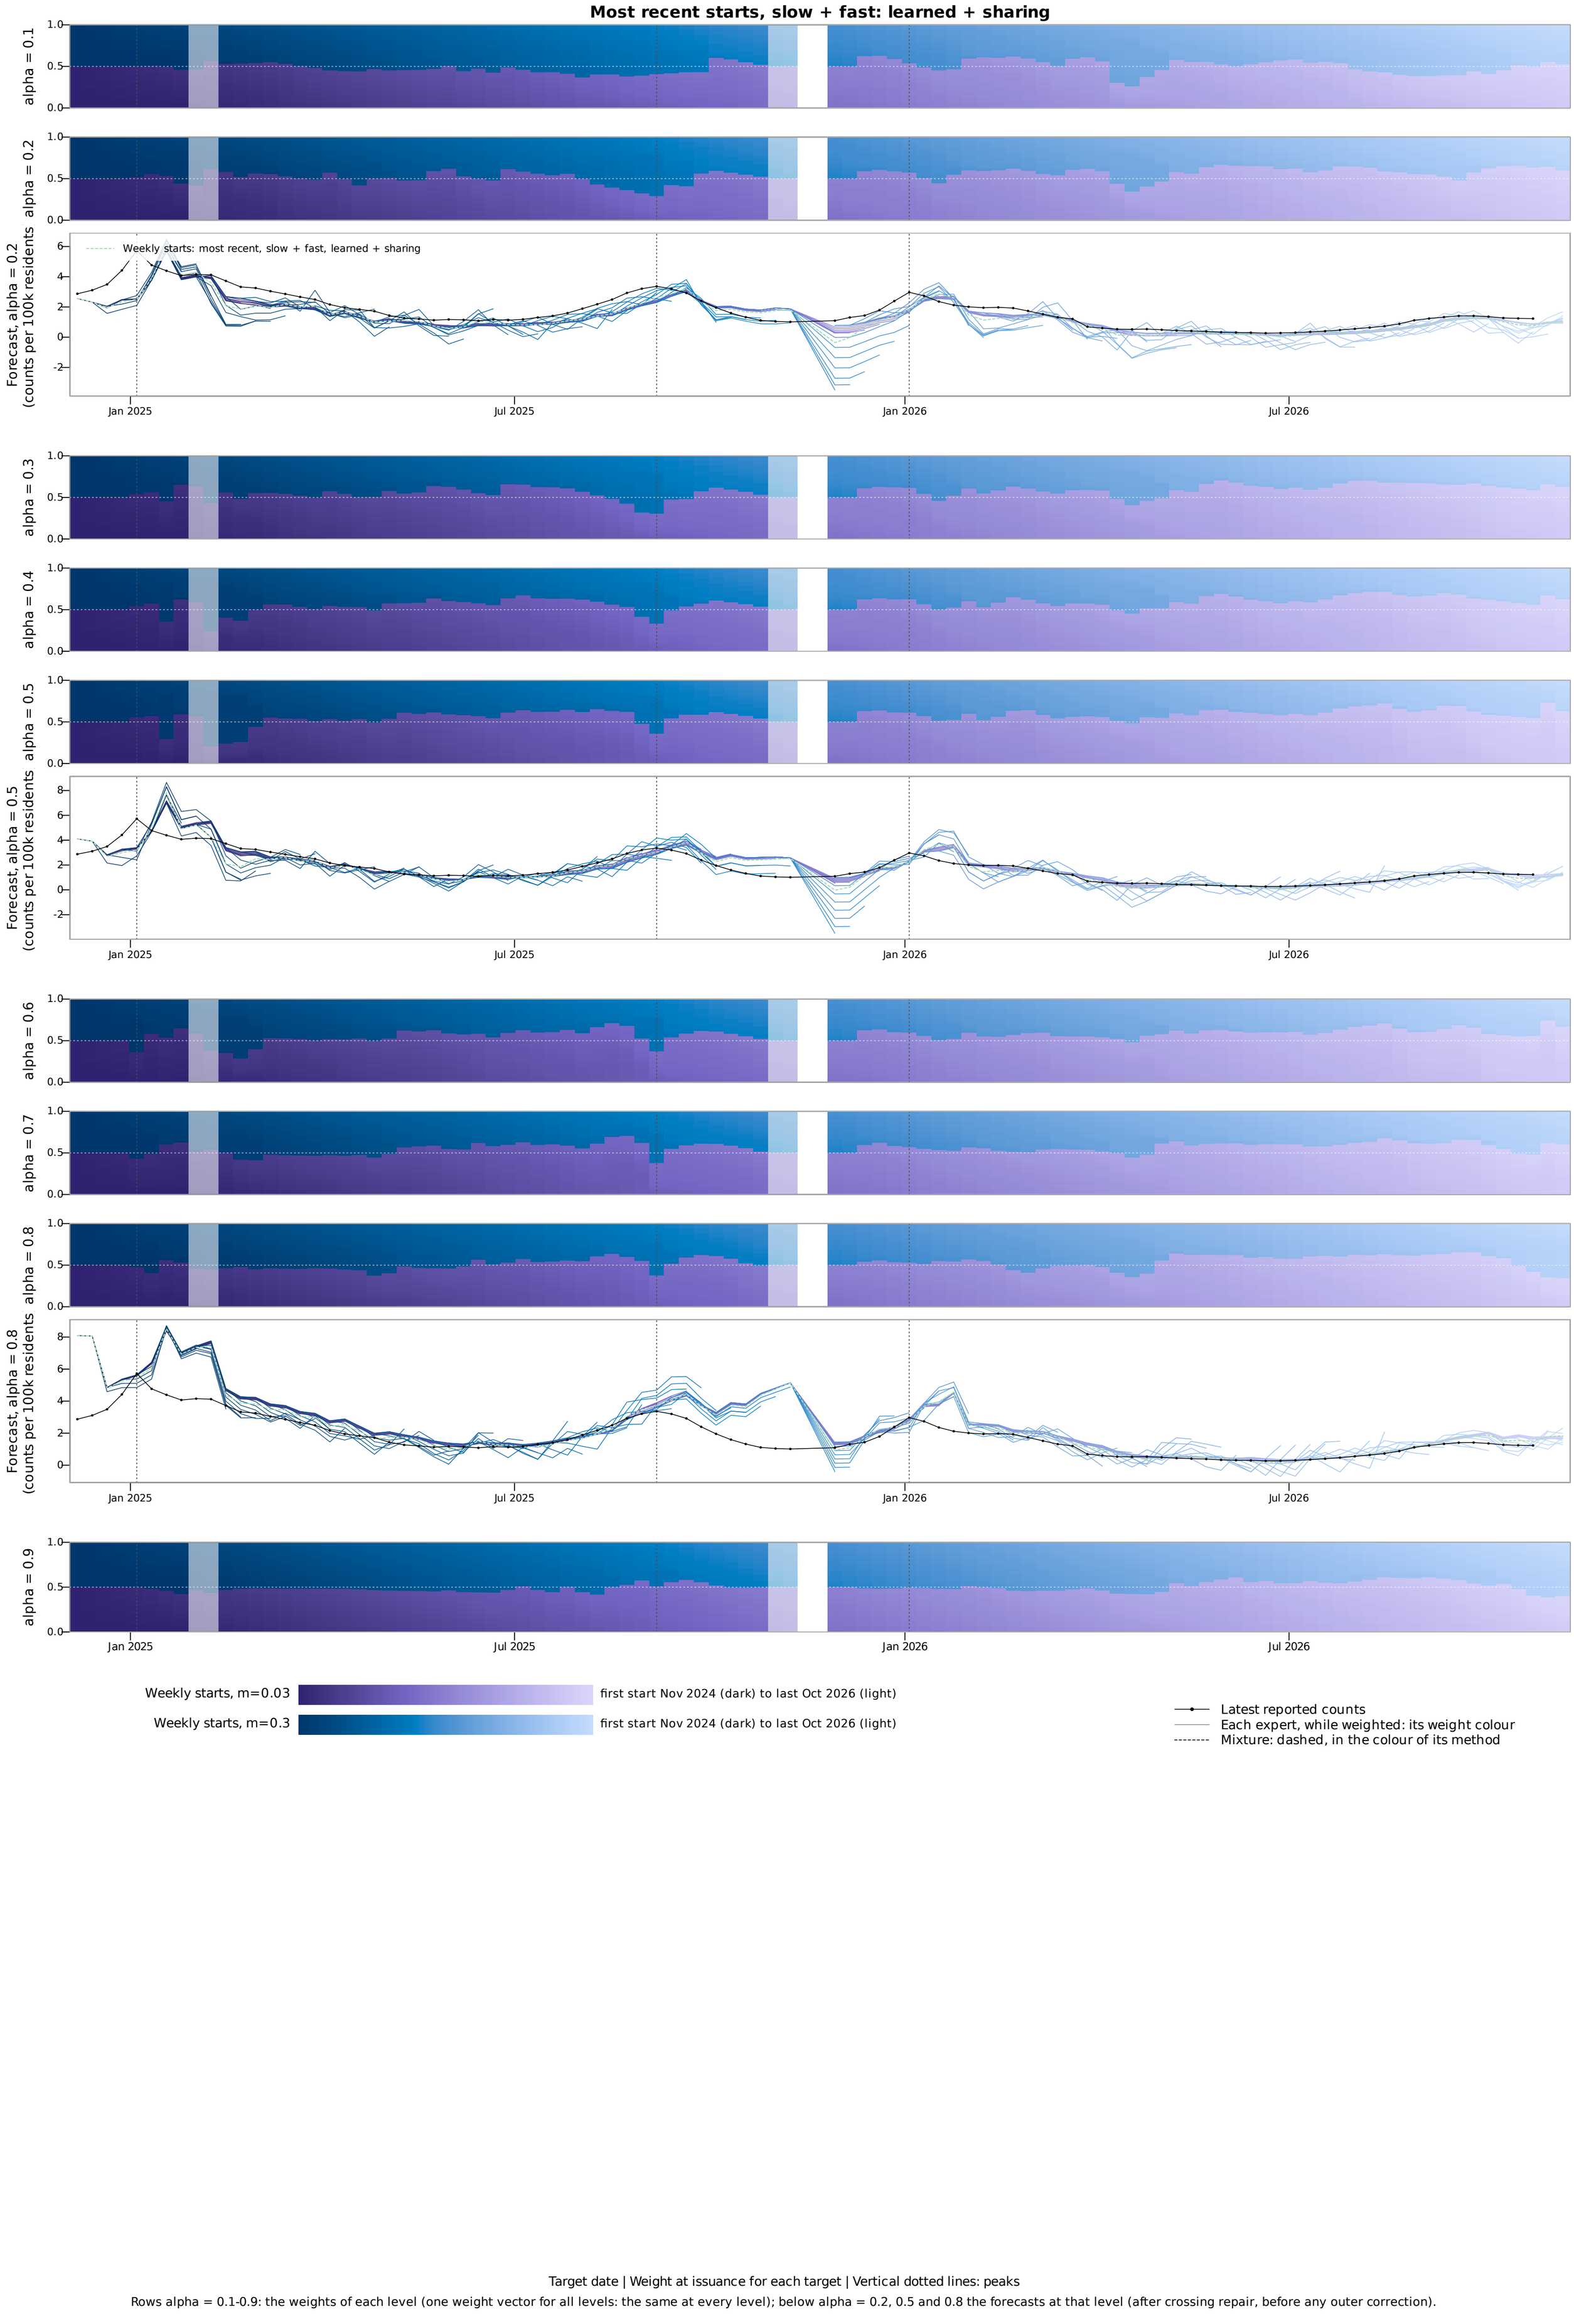
Rows alpha = 0.1-0.9: the weights of each level (one weight vector for all levels: the same at every level); below alpha = 0.2, 0.5 and 0.8 the forecasts at that level (after crossing repair, before any outer correction).

Lighter bars: rounds without a CMU forecast, filled with the latest earlier CMU forecast of the same target week.

Weekly-start weights: slow + fast copies

covid | horizon 1 | United States (US) | evaluation: Jun 2025 to Sep 2026 | raw-count updates | feedback re-run

Most recent starts, slow + fast: learned + sharing



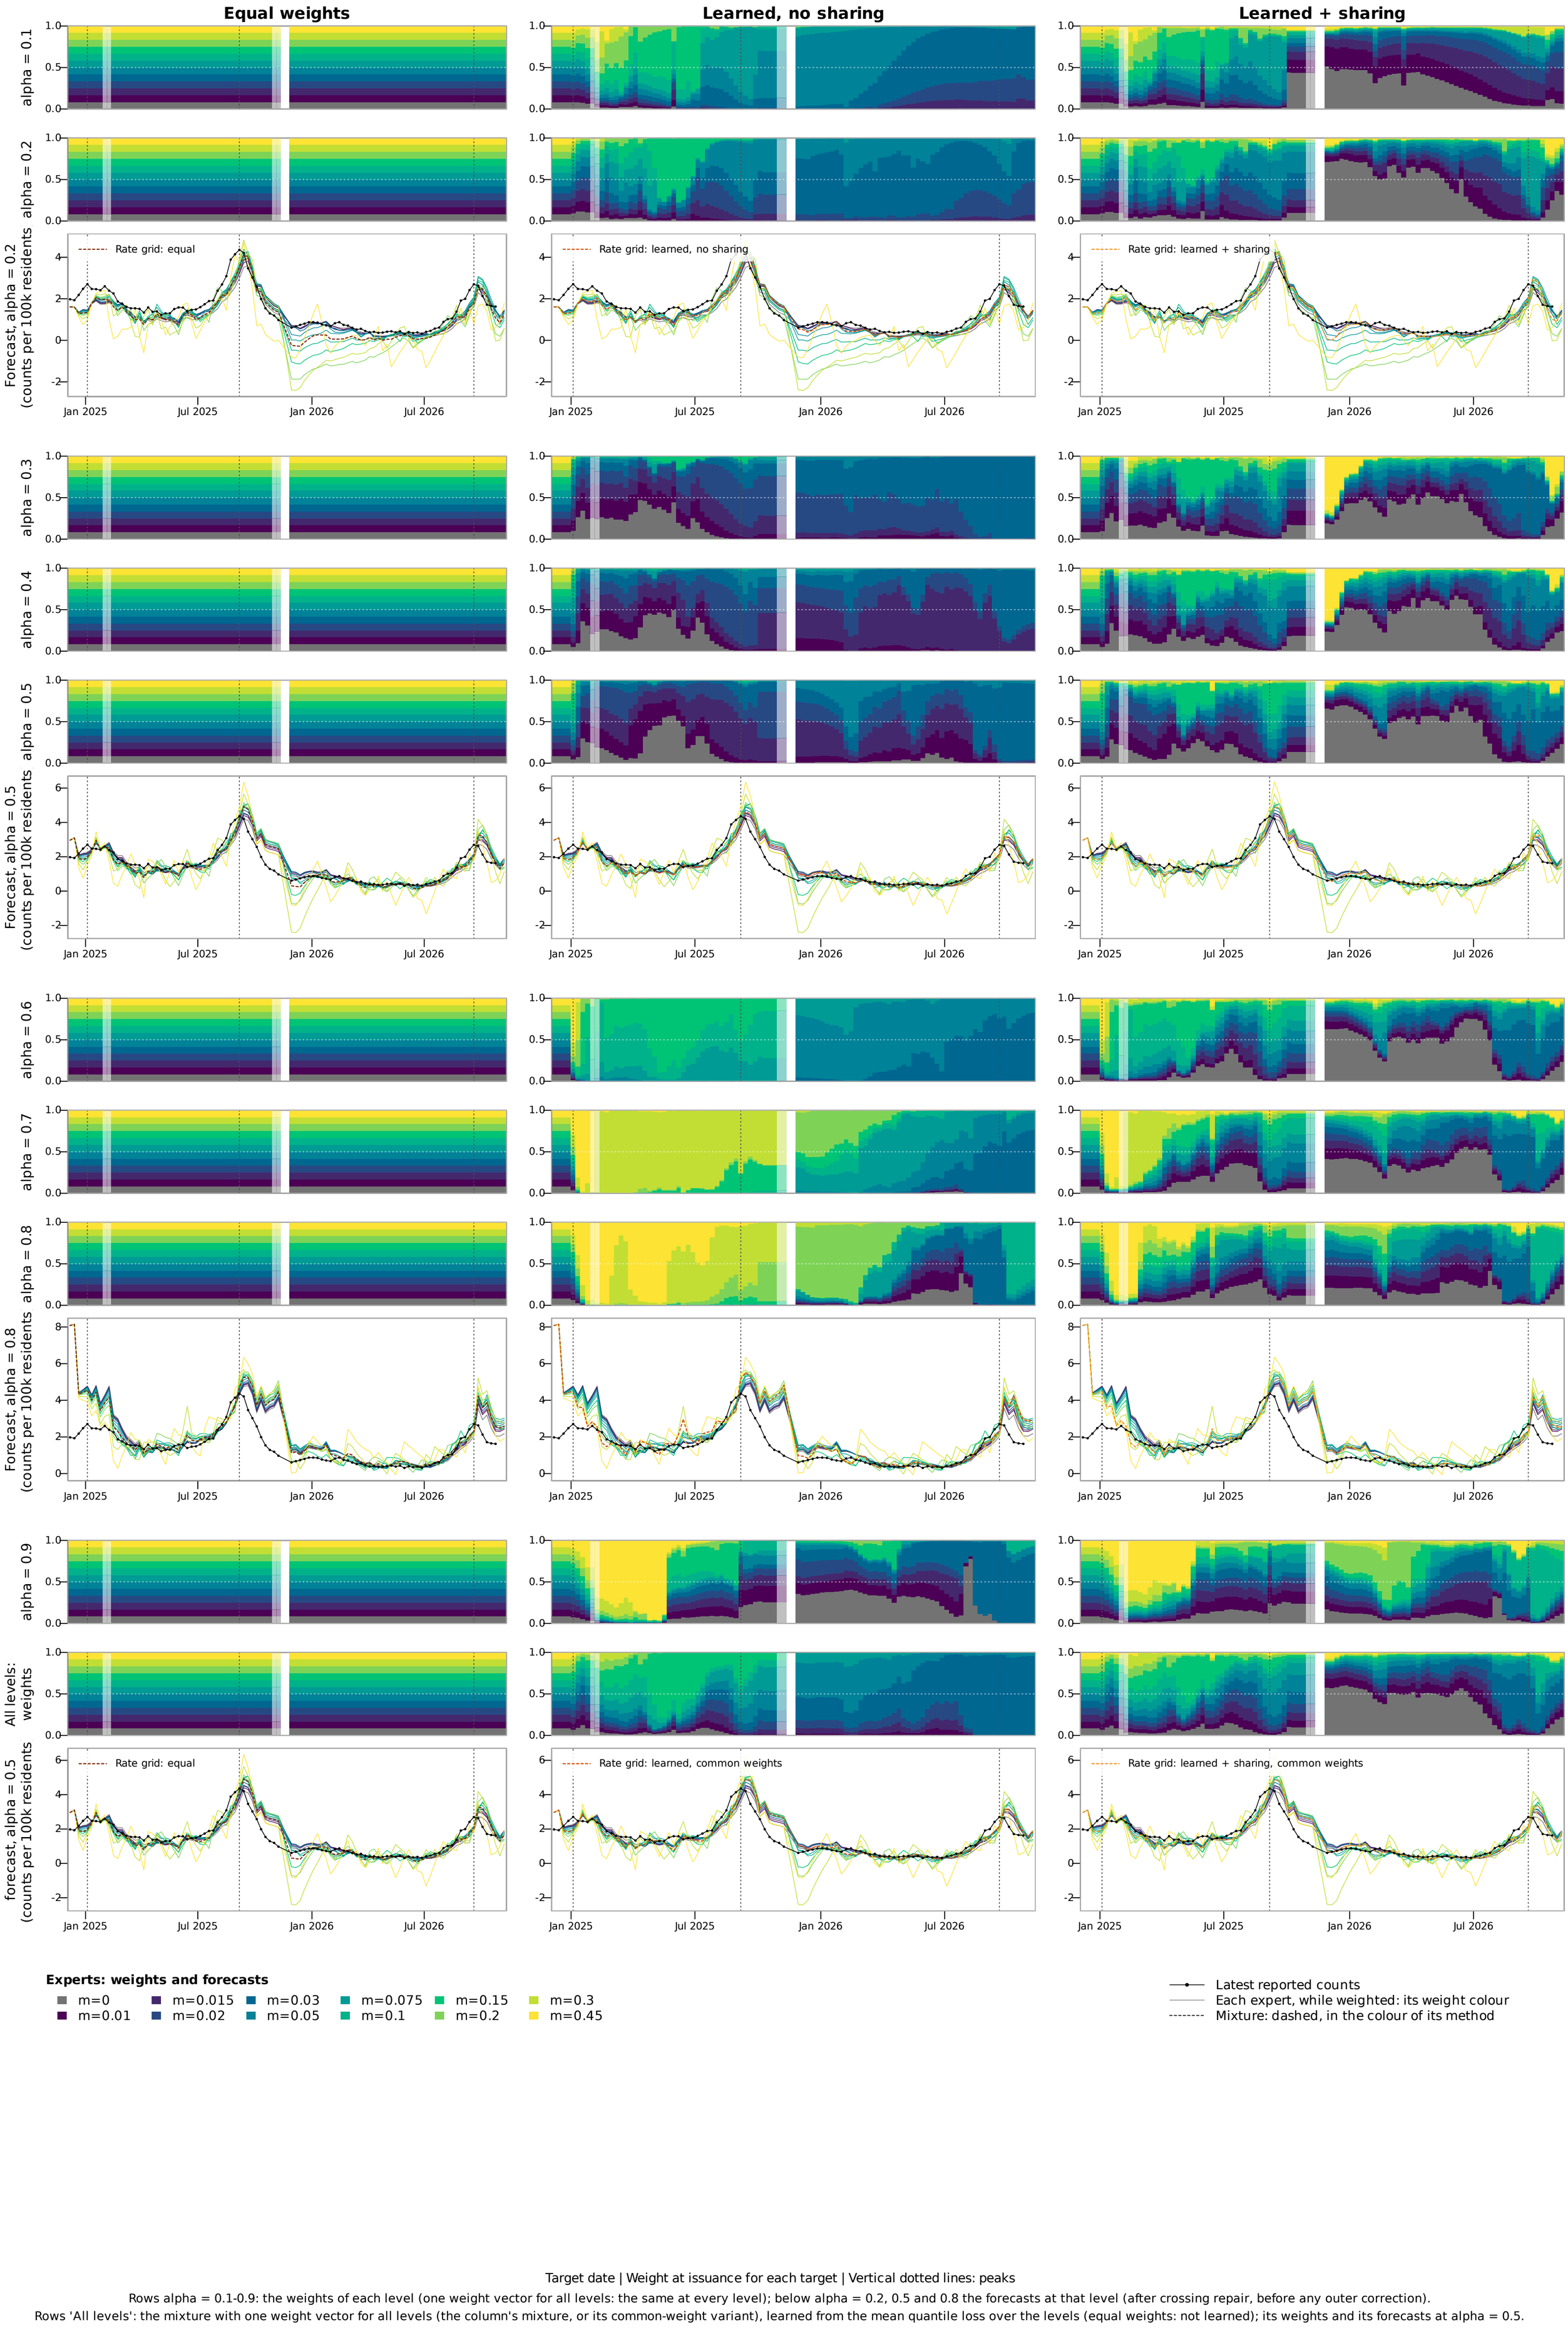
Target date | Weight at issuance for each target | Vertical dotted lines: peaks

Rows alpha = 0.1-0.9: the weights of each level (one weight vector for all levels: the same at every level); below alpha = 0.2, 0.5 and 0.8 the forecasts at that level (after crossing repair, before any outer correction).

Lighter bars: rounds without a CMU forecast, filled with the latest earlier CMU forecast of the same target week.

Frozen rate-grid weights

covid | horizon 1 | California (06) | evaluation: Jun 2025 to Sep 2026 | raw-count updates | feedback re-run



Target date | Weight at issuance for each target | Vertical dotted lines: peaks

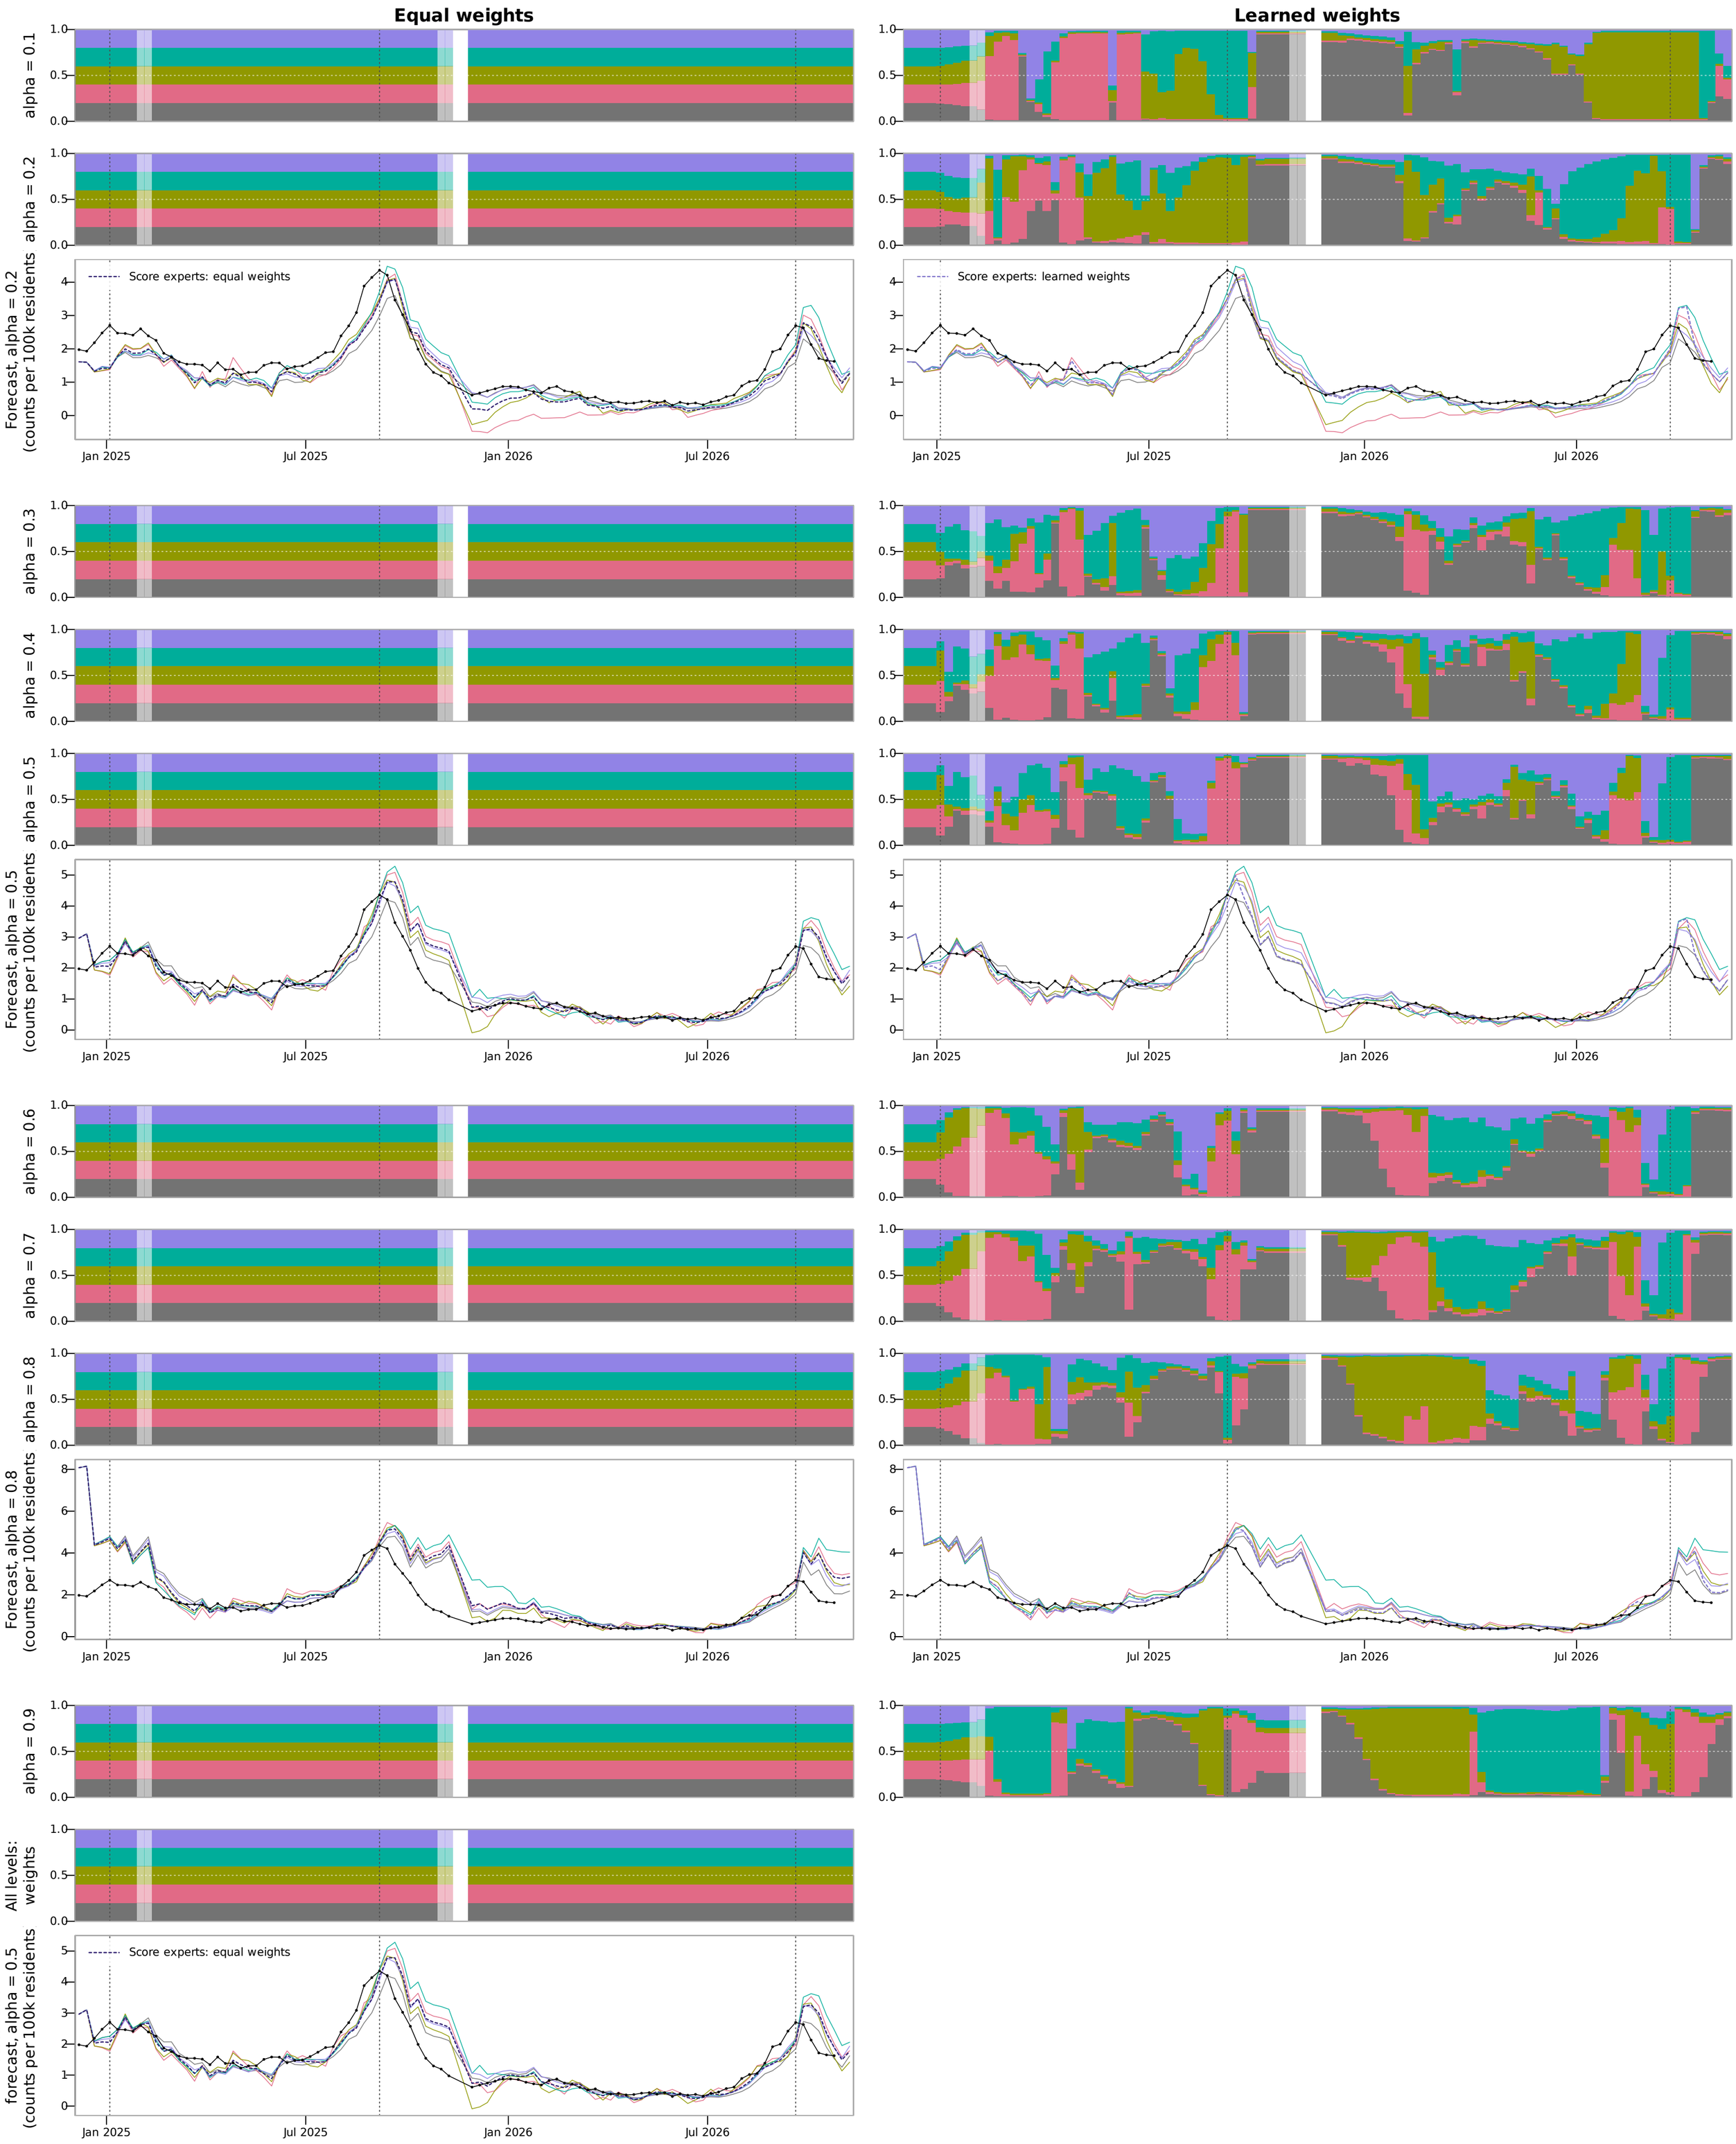
Rows alpha = 0.1-0.9: the weights of each level (one weight vector for all levels: the same at every level); below alpha = 0.2, 0.5 and 0.8 the forecasts at that level (after crossing repair, before any outer correction).

Rows 'All levels': the mixture with one weight vector for all levels (the column's mixture, or its common-weight variant), learned from the mean quantile loss over the levels (equal weights: not learned); its weights and its forecasts at alpha = 0.5.

Lighter bars: rounds without a CMU forecast, filled with the latest earlier CMU forecast of the same target week.

Score-based expert weights

covid | horizon 1 | California (06) | evaluation: Jun 2025 to Sep 2026 | raw-count updates | feedback re-run



Target date | Weight at issuance for each target | Vertical dotted lines: peaks

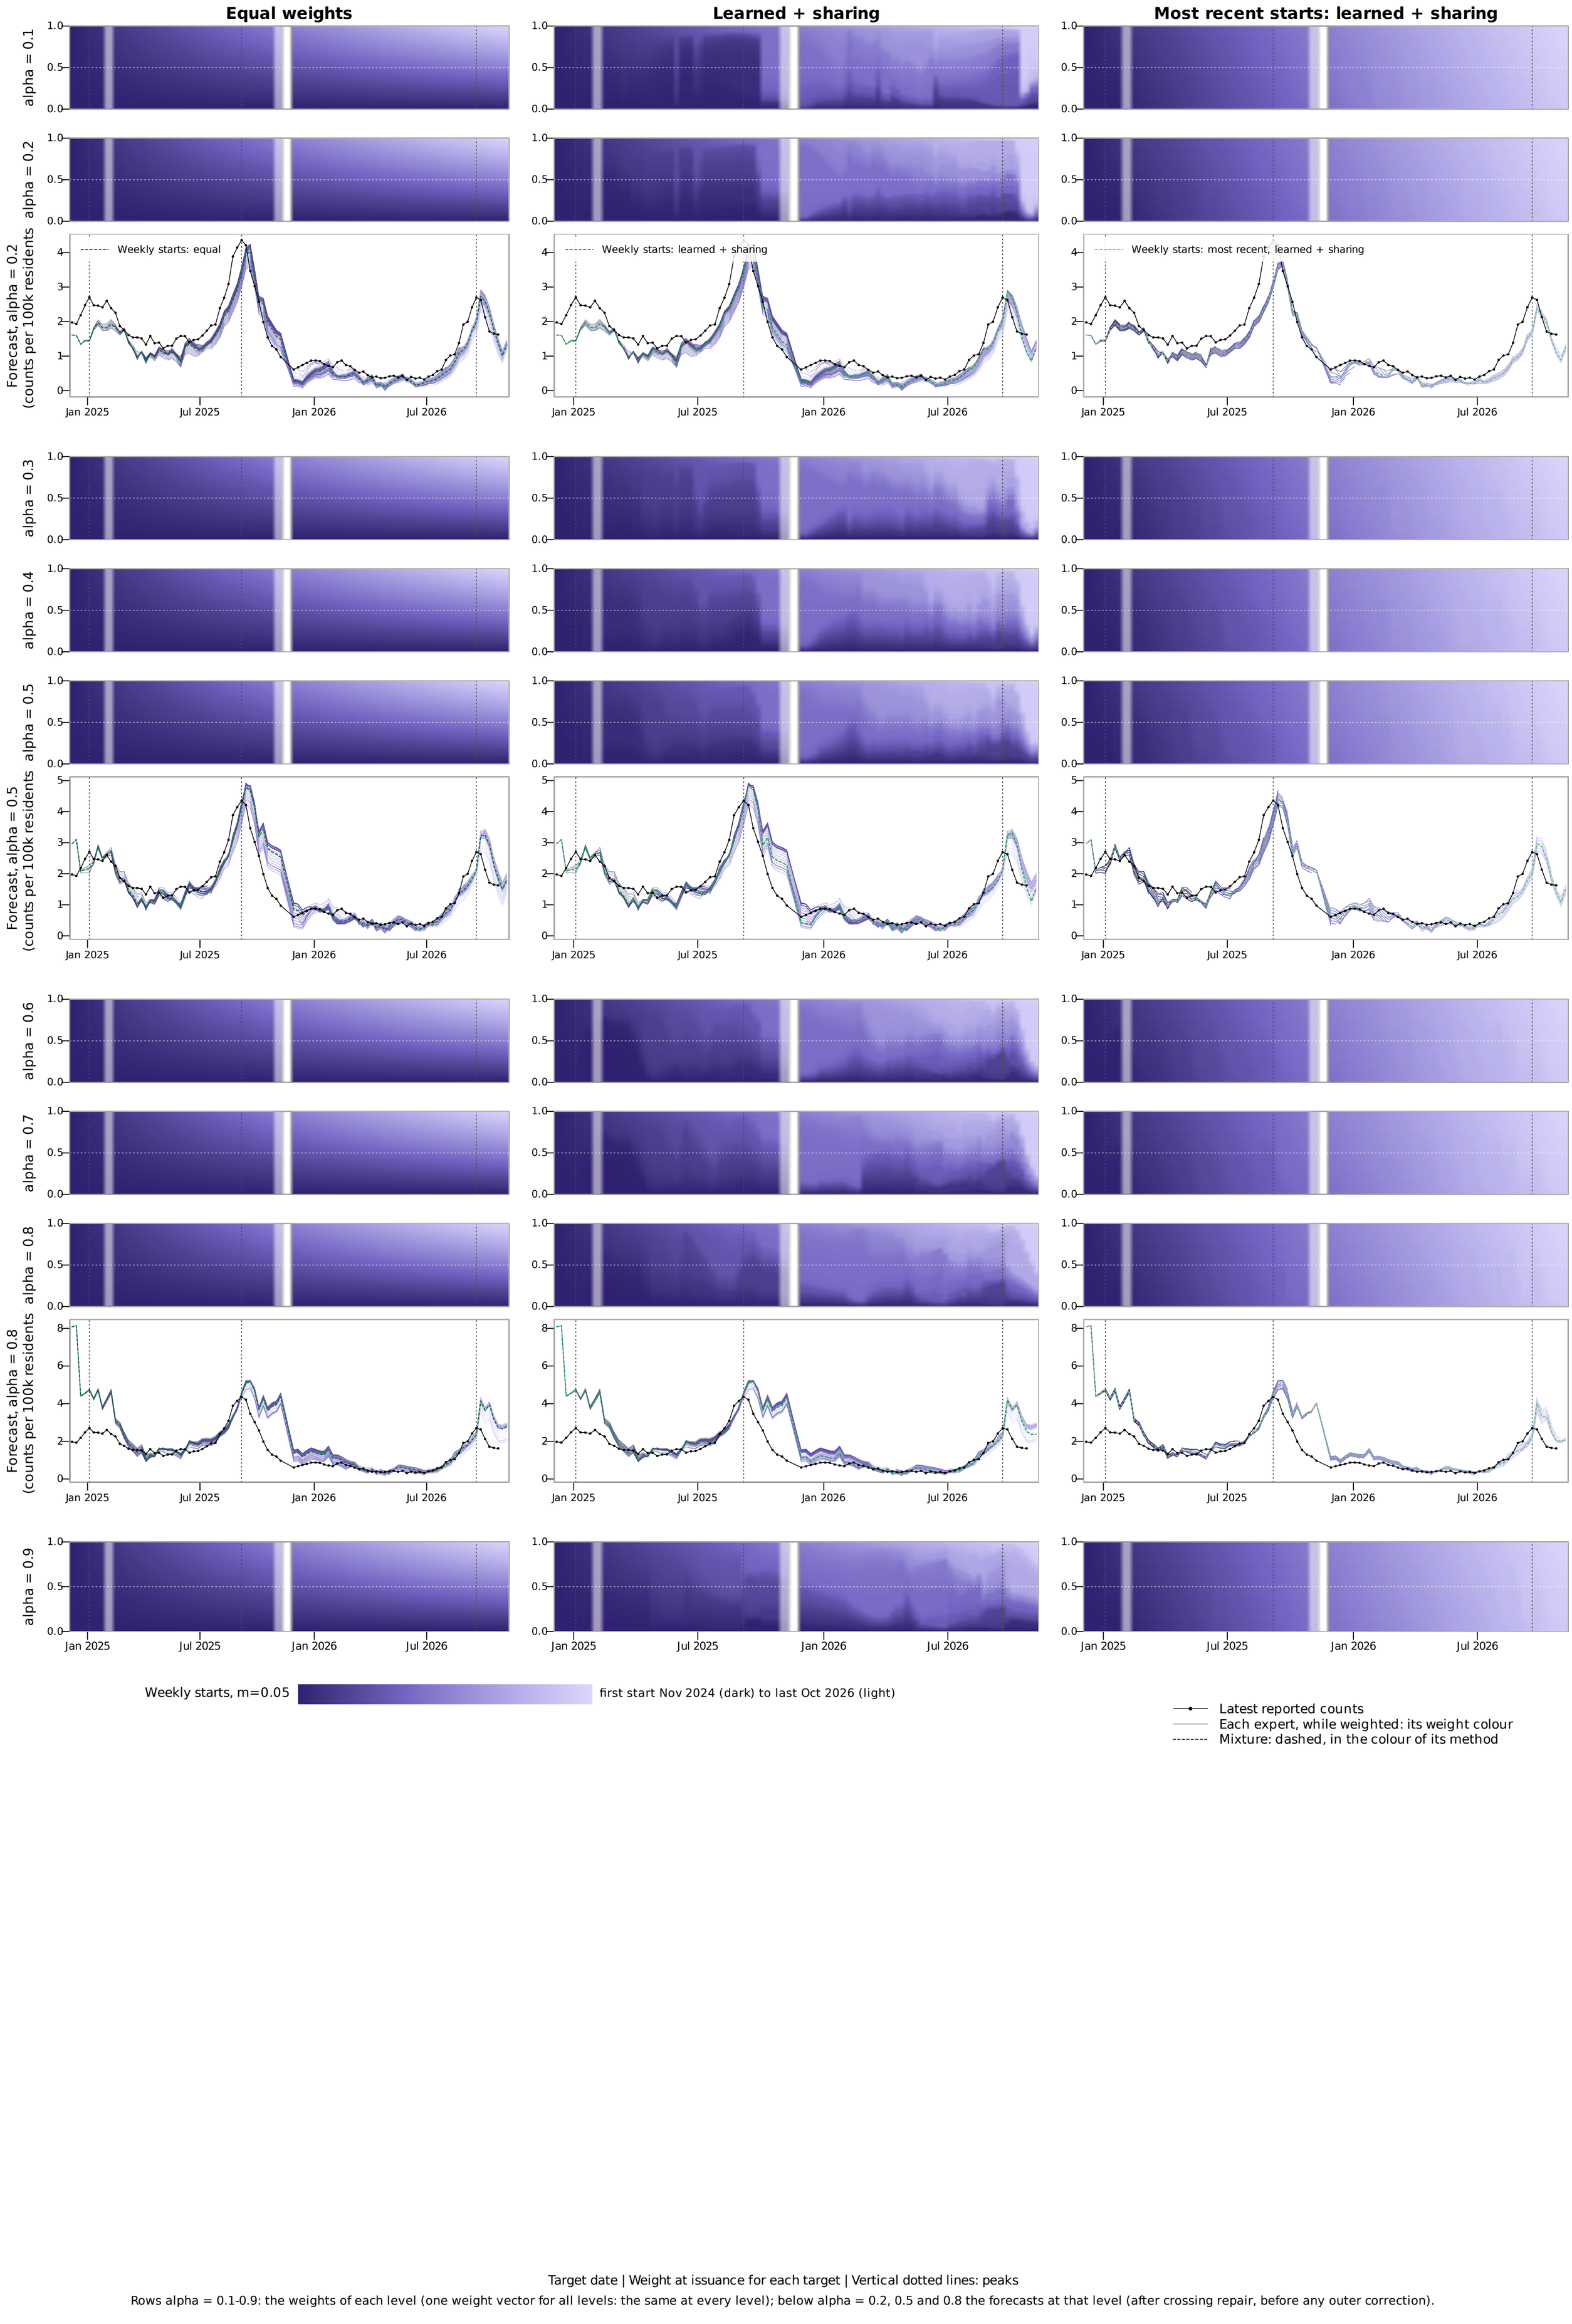
Rows alpha = 0.1-0.9: the weights of each level (one weight vector for all levels: the same at every level); below alpha = 0.2, 0.5 and 0.8 the forecasts at that level (after crossing repair, before any outer correction).

Rows 'All levels': the mixture with one weight vector for all levels (the column's mixture, or its common-weight variant), learned from the mean quantile loss over the levels (equal weights: not learned); its weights and its forecasts at alpha = 0.5.

Lighter bars: rounds without a CMU forecast, filled with the latest earlier CMU forecast of the same target week.

Weekly-start weights

covid | horizon 1 | California (06) | evaluation: Jun 2025 to Sep 2026 | raw-count updates | feedback re-run



Target date | Weight at issuance for each target | Vertical dotted lines: peaks

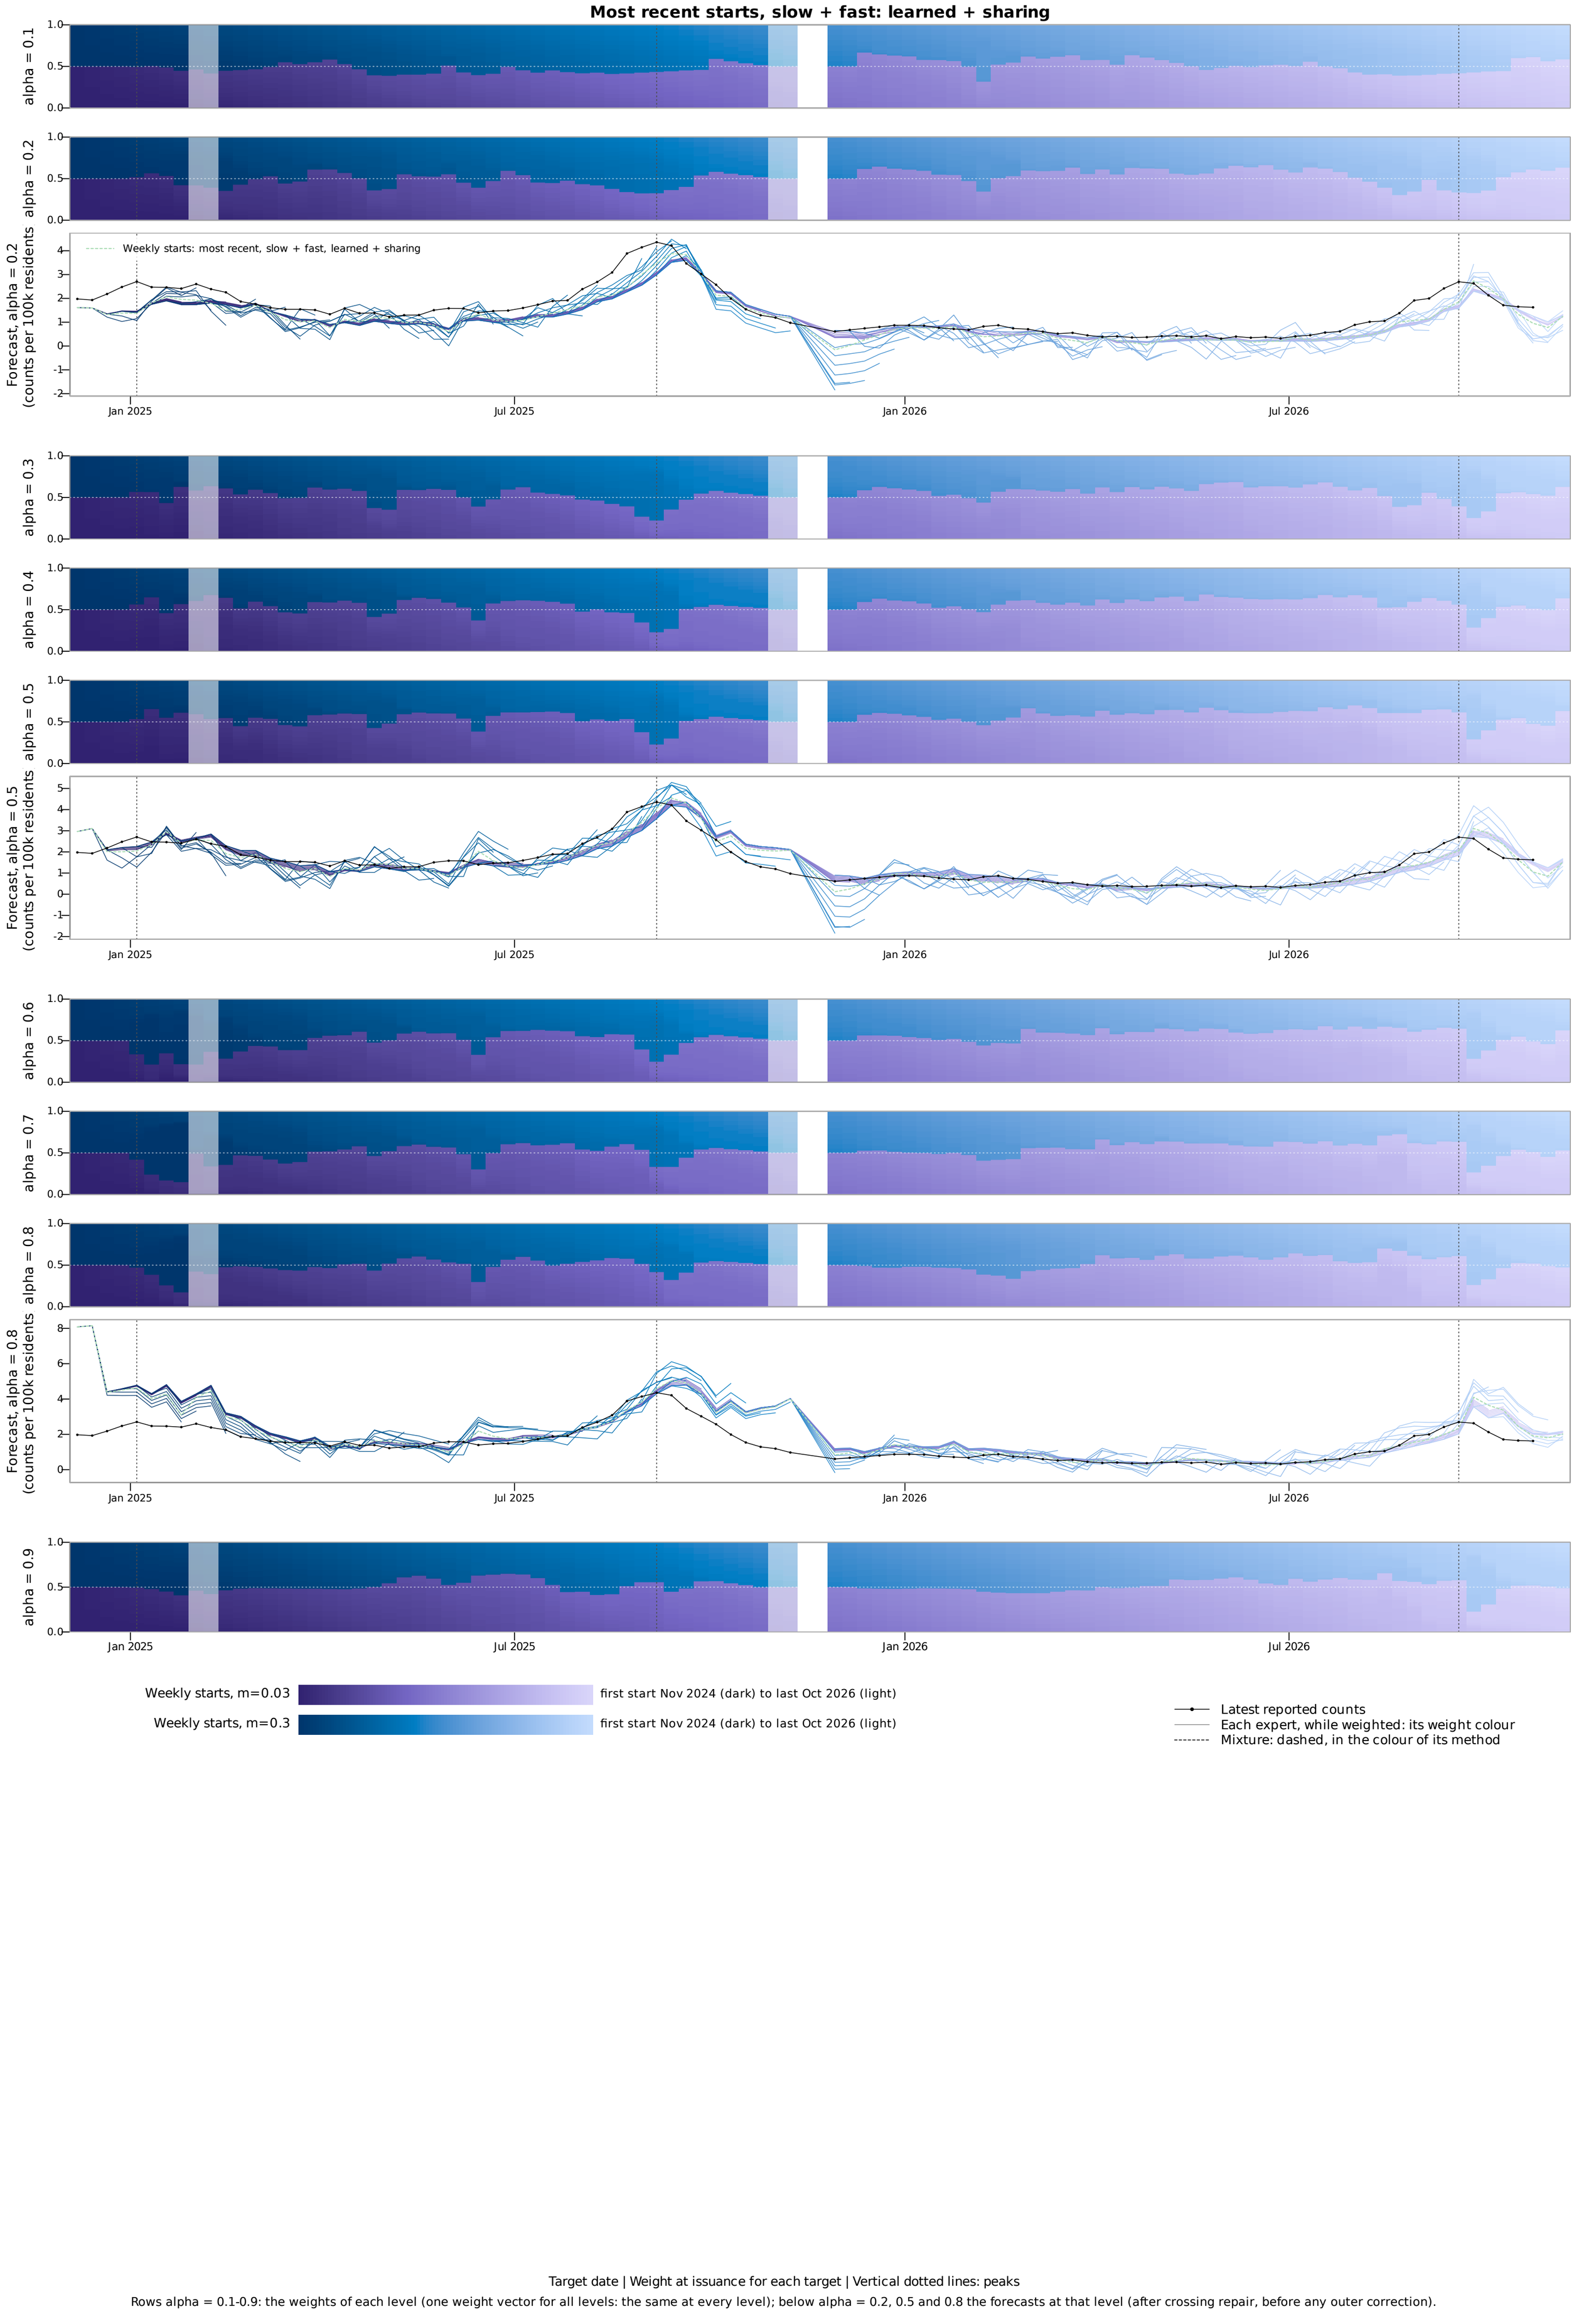
Rows alpha = 0.1-0.9: the weights of each level (one weight vector for all levels: the same at every level); below alpha = 0.2, 0.5 and 0.8 the forecasts at that level (after crossing repair, before any outer correction).

Lighter bars: rounds without a CMU forecast, filled with the latest earlier CMU forecast of the same target week.

Weekly-start weights: slow + fast copies

covid | horizon 1 | California (06) | evaluation: Jun 2025 to Sep 2026 | raw-count updates | feedback re-run

Most recent starts, slow + fast: learned + sharing



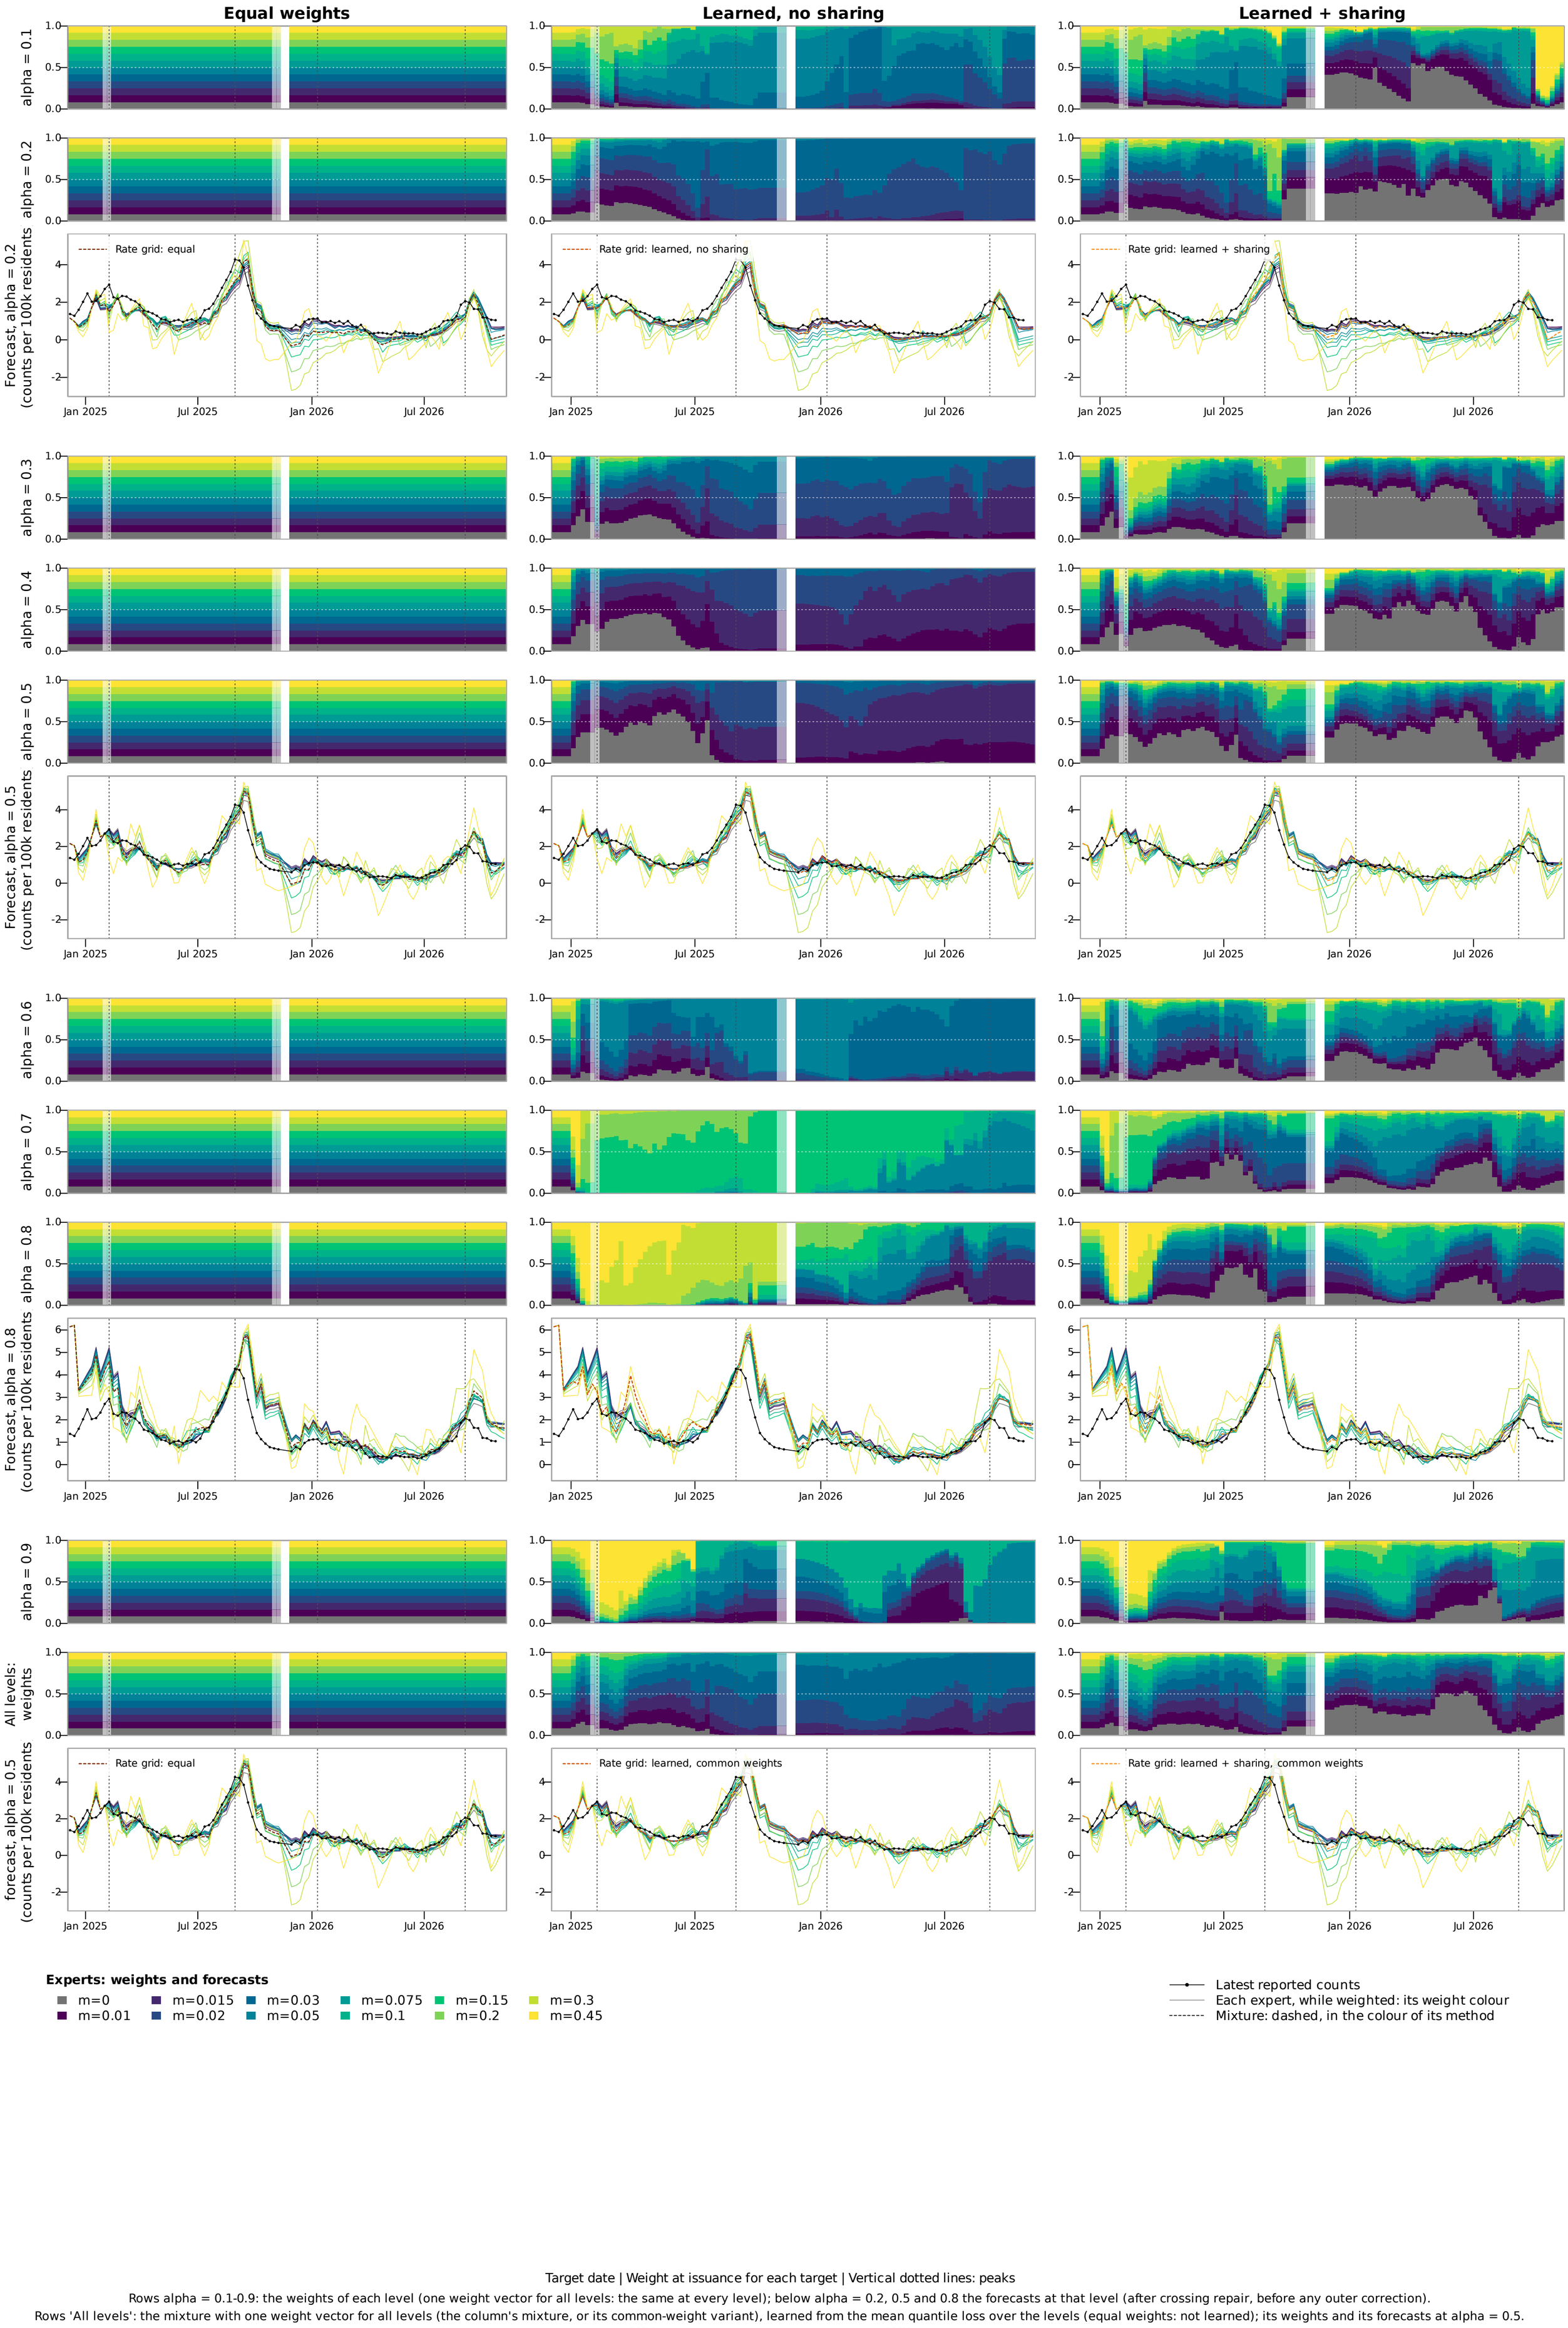
Target date | Weight at issuance for each target | Vertical dotted lines: peaks

Rows alpha = 0.1-0.9: the weights of each level (one weight vector for all levels: the same at every level); below alpha = 0.2, 0.5 and 0.8 the forecasts at that level (after crossing repair, before any outer correction).

Lighter bars: rounds without a CMU forecast, filled with the latest earlier CMU forecast of the same target week.

Frozen rate-grid weights

covid | horizon 1 | Texas (48) | evaluation: Jun 2025 to Sep 2026 | raw-count updates | feedback re-run



Target date | Weight at issuance for each target | Vertical dotted lines: peaks

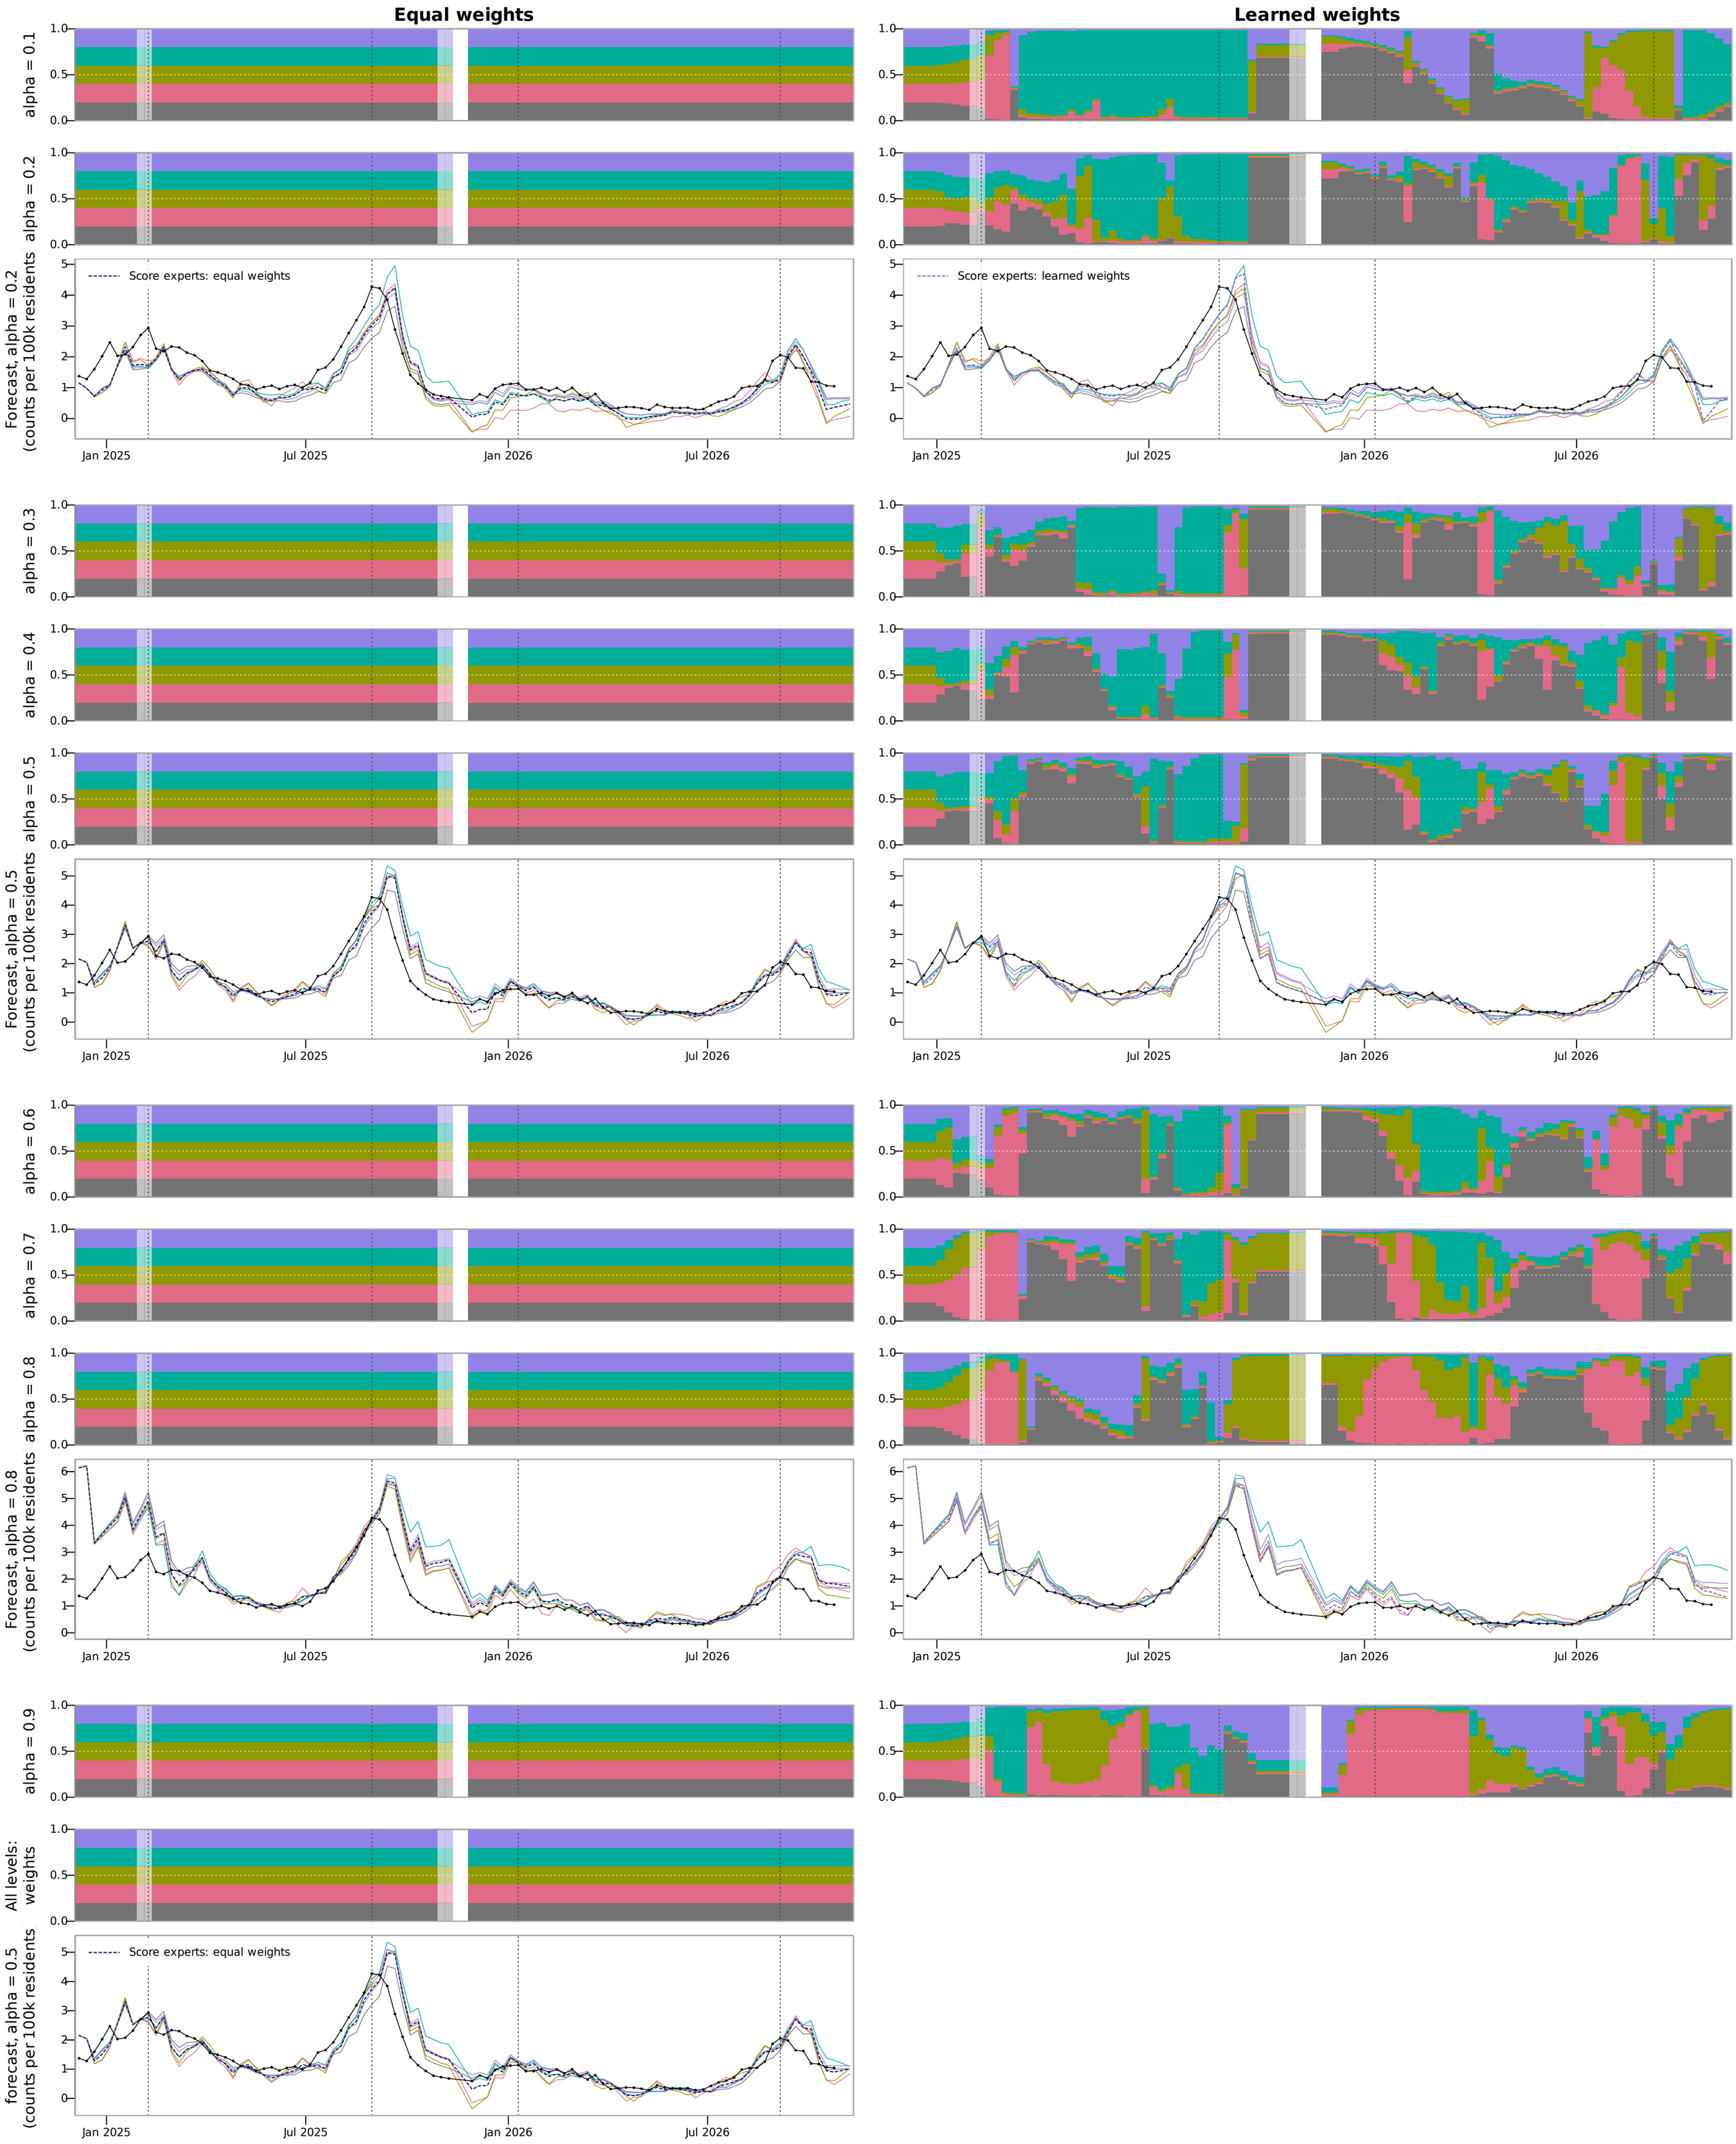
Rows alpha = 0.1-0.9: the weights of each level (one weight vector for all levels: the same at every level); below alpha = 0.2, 0.5 and 0.8 the forecasts at that level (after crossing repair, before any outer correction).

Rows 'All levels': the mixture with one weight vector for all levels (the column's mixture, or its common-weight variant), learned from the mean quantile loss over the levels (equal weights: not learned); its weights and its forecasts at alpha = 0.5.

Lighter bars: rounds without a CMU forecast, filled with the latest earlier CMU forecast of the same target week.

Score-based expert weights

covid | horizon 1 | Texas (48) | evaluation: Jun 2025 to Sep 2026 | raw-count updates | feedback re-run



Target date | Weight at issuance for each target | Vertical dotted lines: peaks

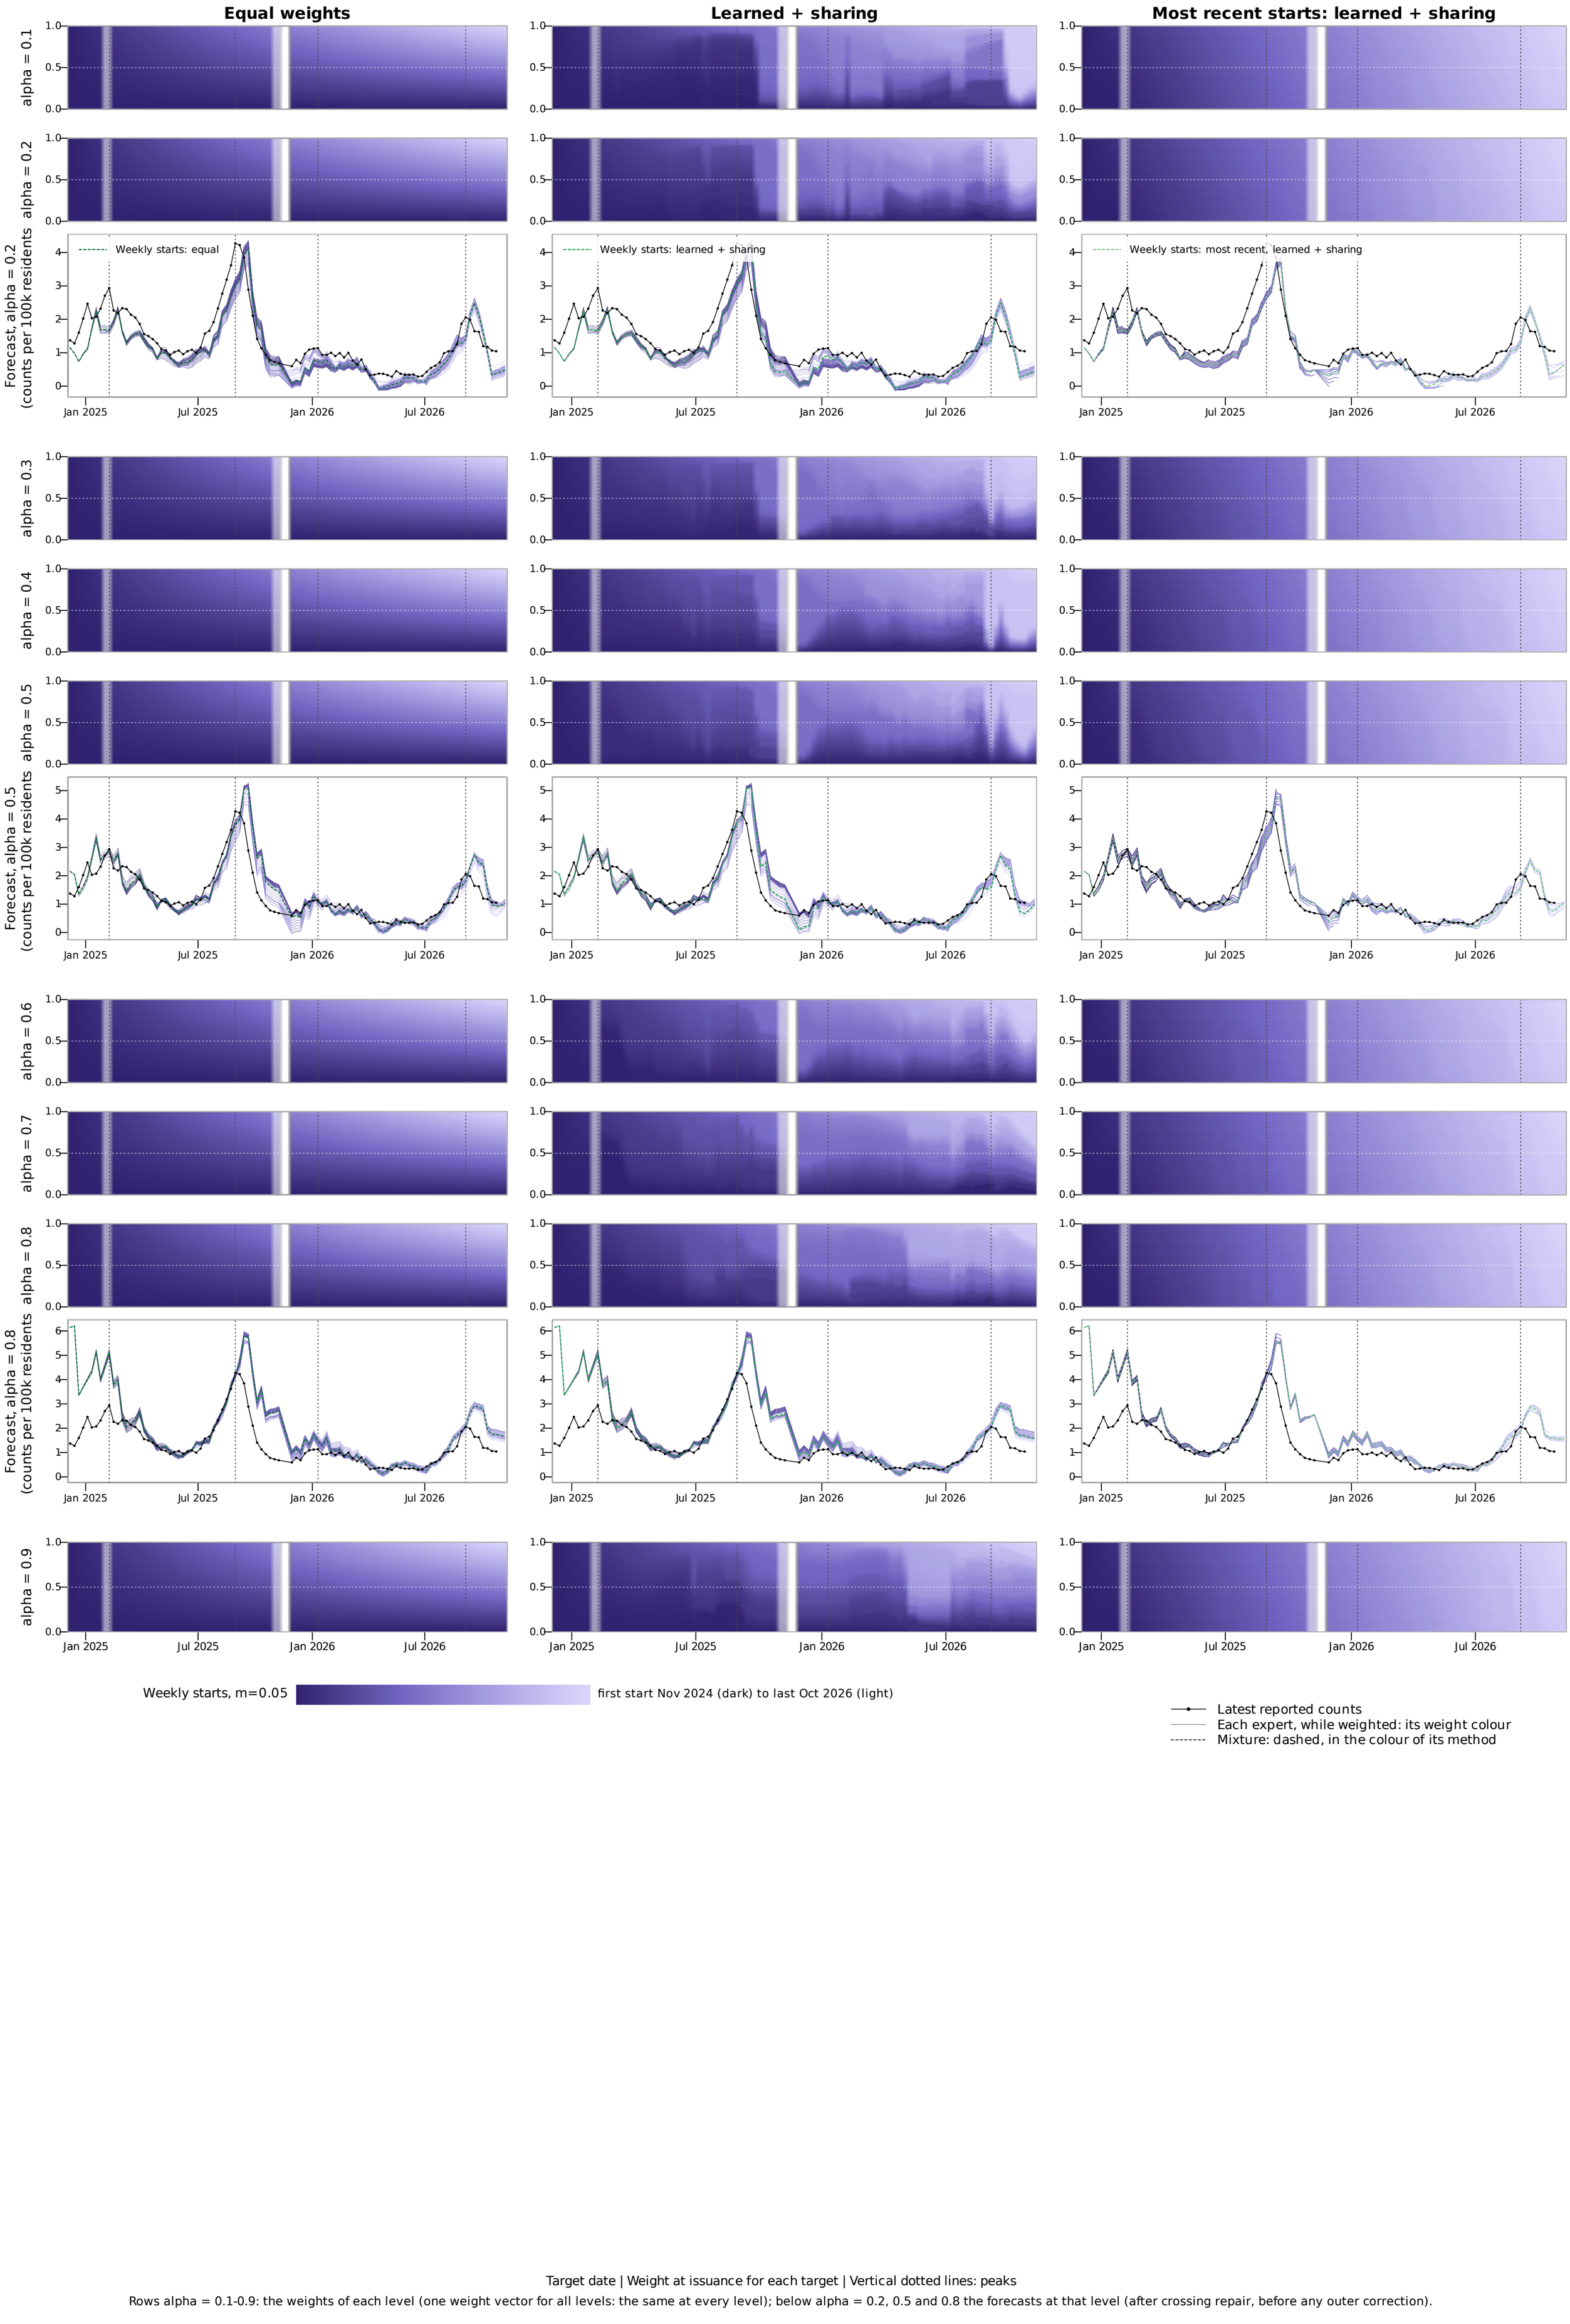
Rows alpha = 0.1-0.9: the weights of each level (one weight vector for all levels: the same at every level); below alpha = 0.2, 0.5 and 0.8 the forecasts at that level (after crossing repair, before any outer correction).

Rows 'All levels': the mixture with one weight vector for all levels (the column's mixture, or its common-weight variant), learned from the mean quantile loss over the levels (equal weights: not learned); its weights and its forecasts at alpha = 0.5.

Lighter bars: rounds without a CMU forecast, filled with the latest earlier CMU forecast of the same target week.

Weekly-start weights

covid | horizon 1 | Texas (48) | evaluation: Jun 2025 to Sep 2026 | raw-count updates | feedback re-run



Target date | Weight at issuance for each target | Vertical dotted lines: peaks

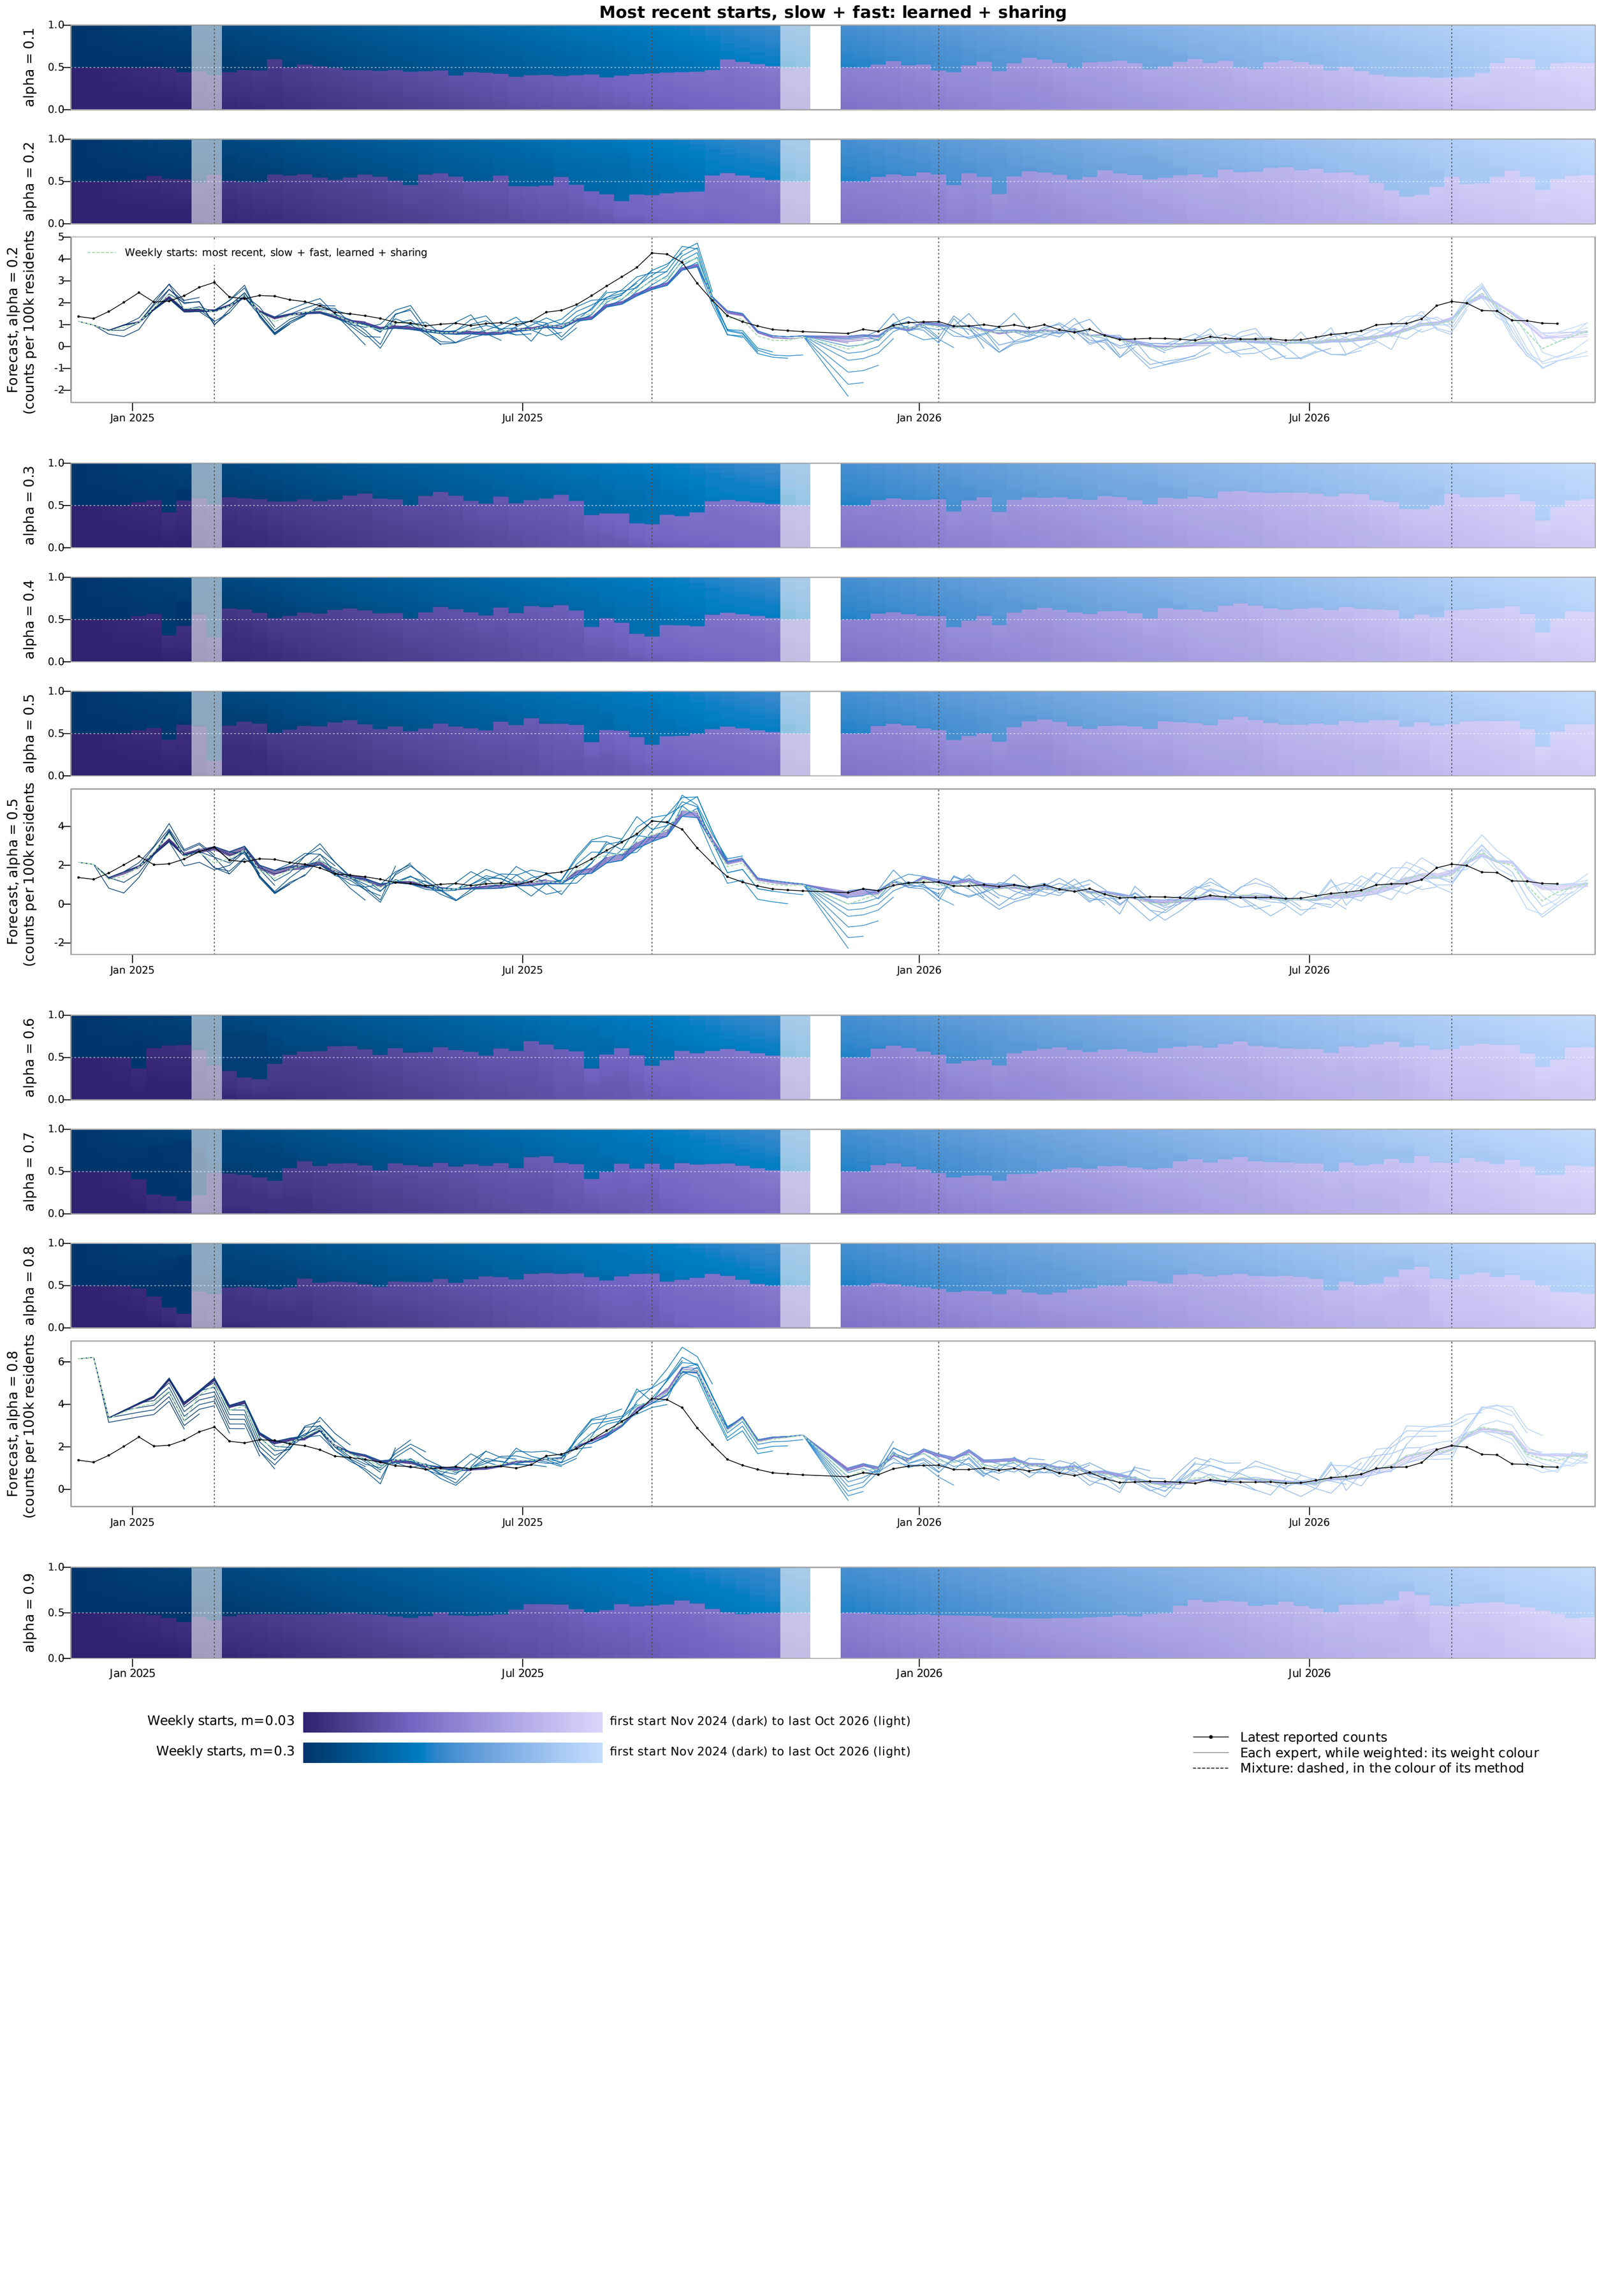
Rows alpha = 0.1-0.9: the weights of each level (one weight vector for all levels: the same at every level); below alpha = 0.2, 0.5 and 0.8 the forecasts at that level (after crossing repair, before any outer correction).

Lighter bars: rounds without a CMU forecast, filled with the latest earlier CMU forecast of the same target week.

Weekly-start weights: slow + fast copies

covid | horizon 1 | Texas (48) | evaluation: Jun 2025 to Sep 2026 | raw-count updates | feedback re-run

Most recent starts, slow + fast: learned + sharing



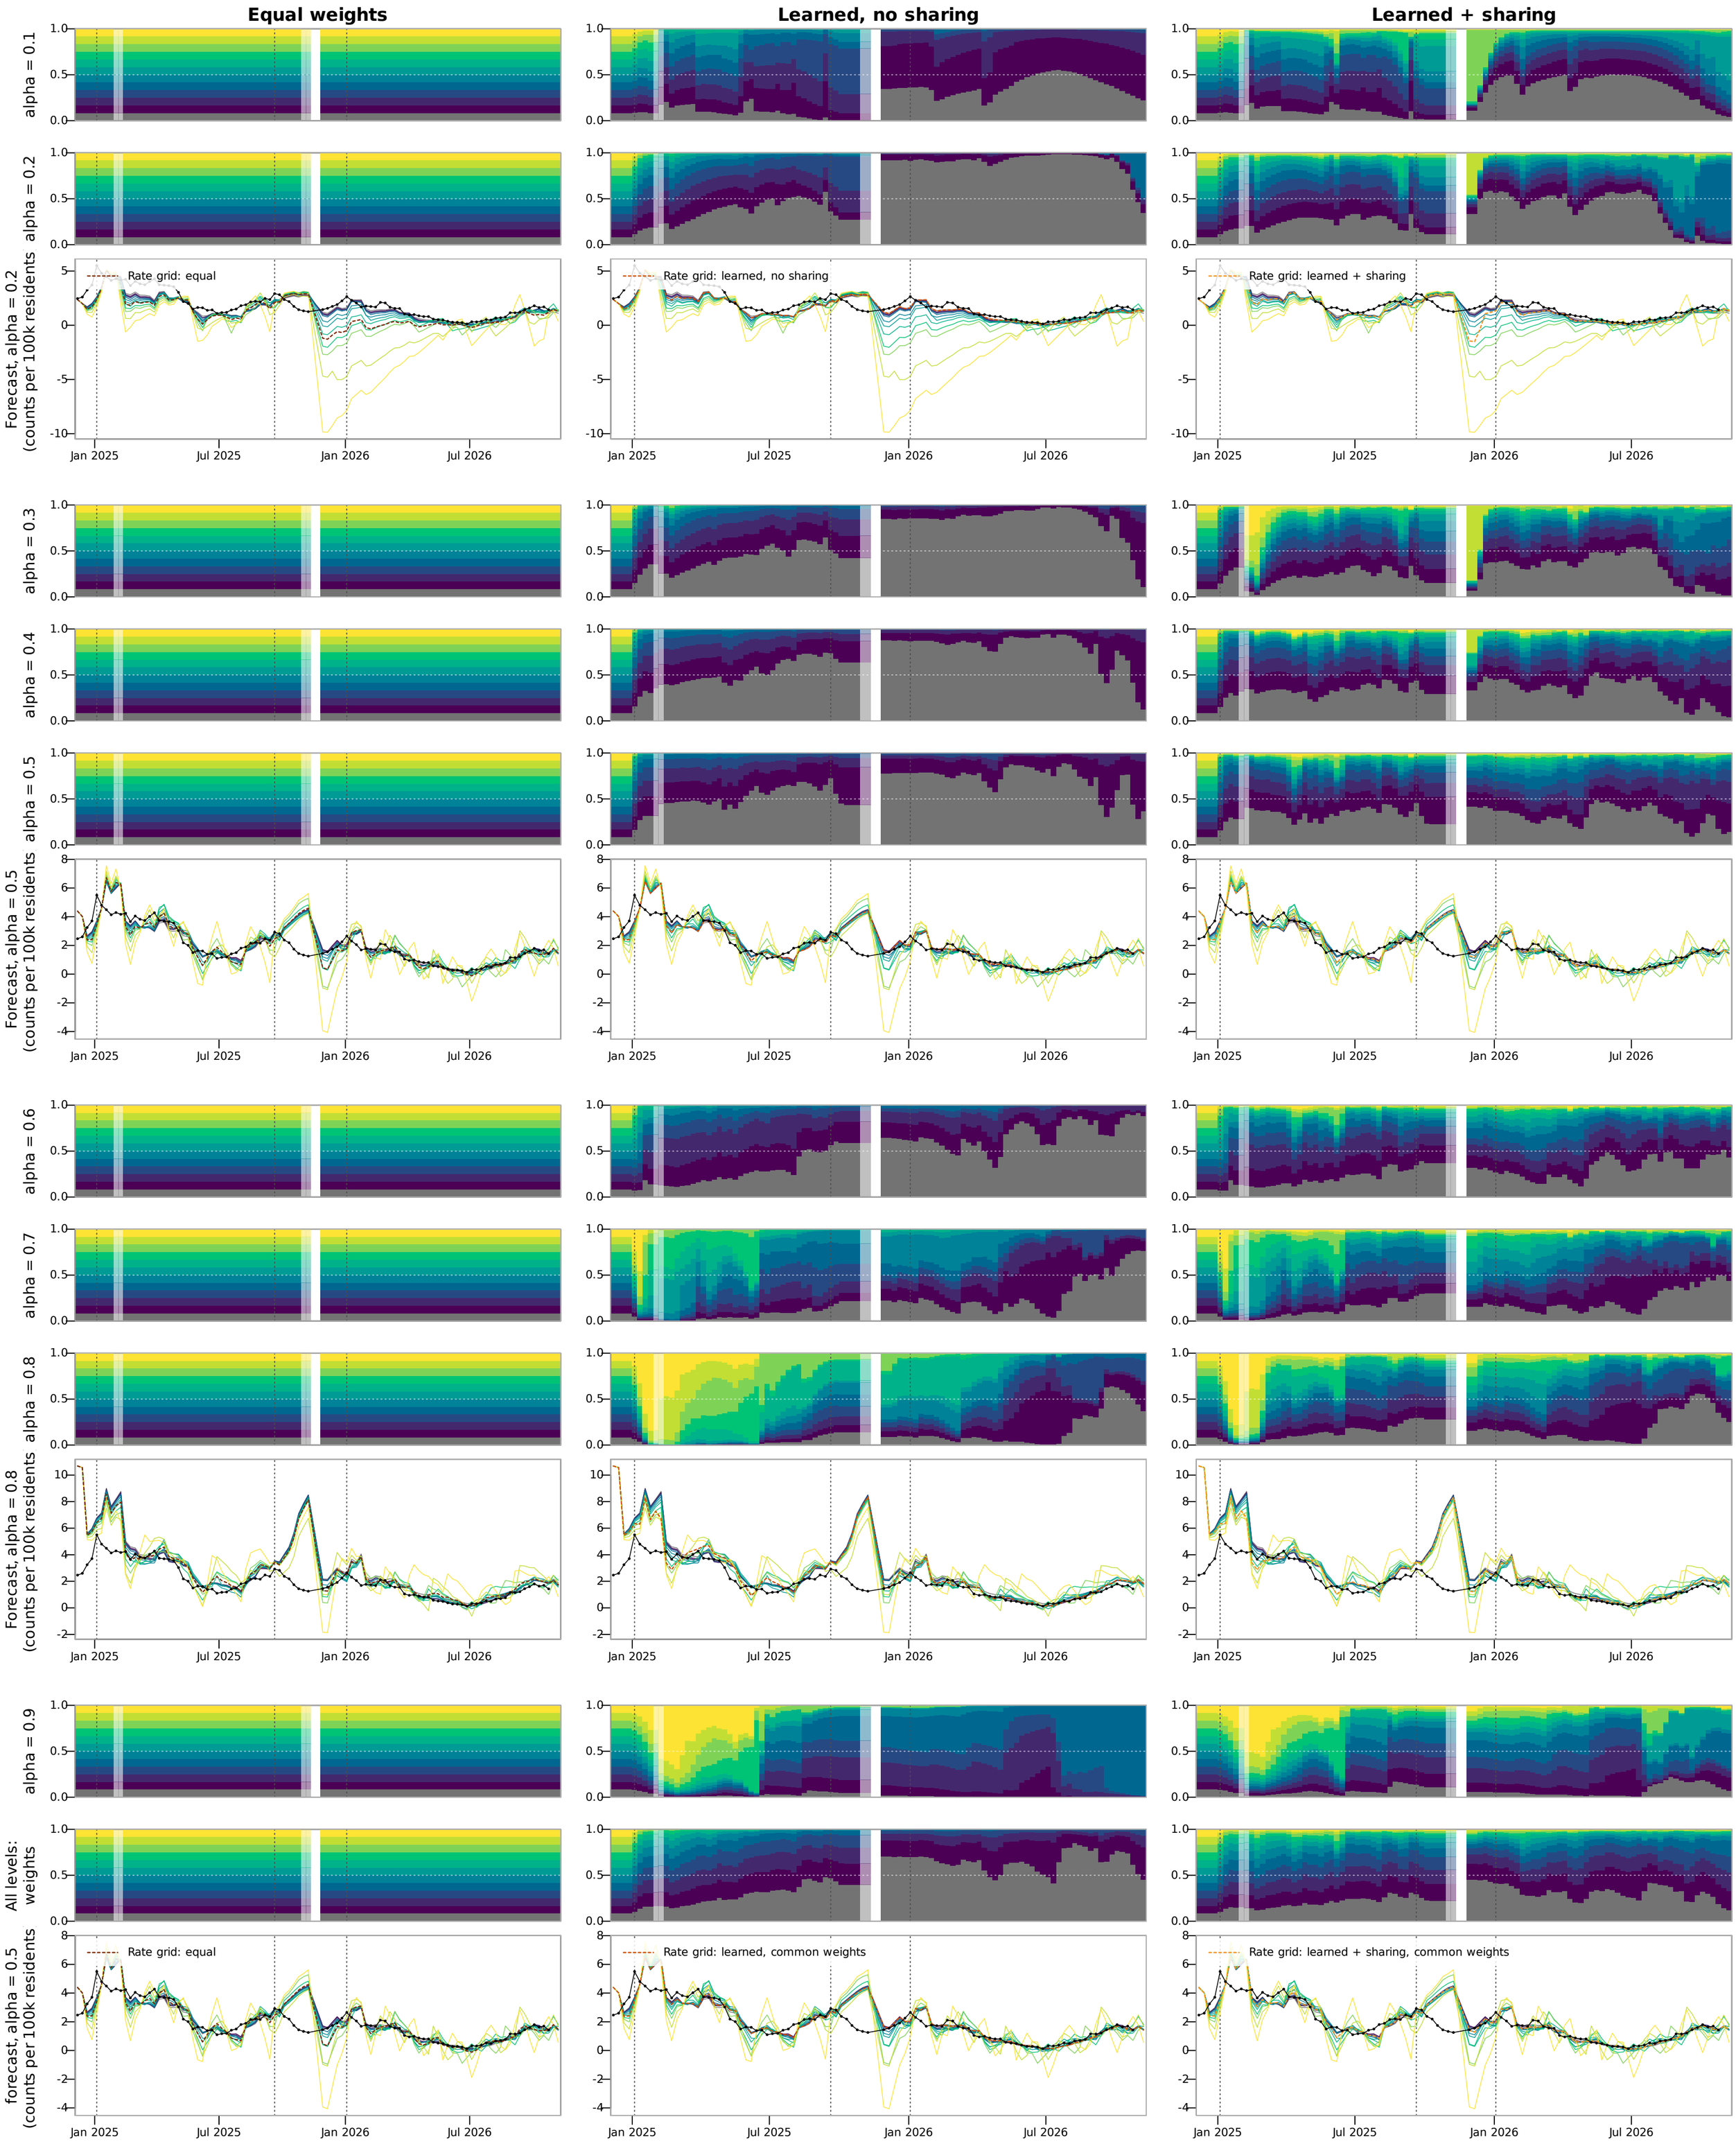
Target date | Weight at issuance for each target | Vertical dotted lines: peaks

Rows alpha = 0.1-0.9: the weights of each level (one weight vector for all levels: the same at every level); below alpha = 0.2, 0.5 and 0.8 the forecasts at that level (after crossing repair, before any outer correction).

Lighter bars: rounds without a CMU forecast, filled with the latest earlier CMU forecast of the same target week.

Frozen rate-grid weights

covid | horizon 1 | New York (36) | evaluation: Jun 2025 to Sep 2026 | raw-count updates | feedback re-run



Target date | Weight at issuance for each target | Vertical dotted lines: peaks

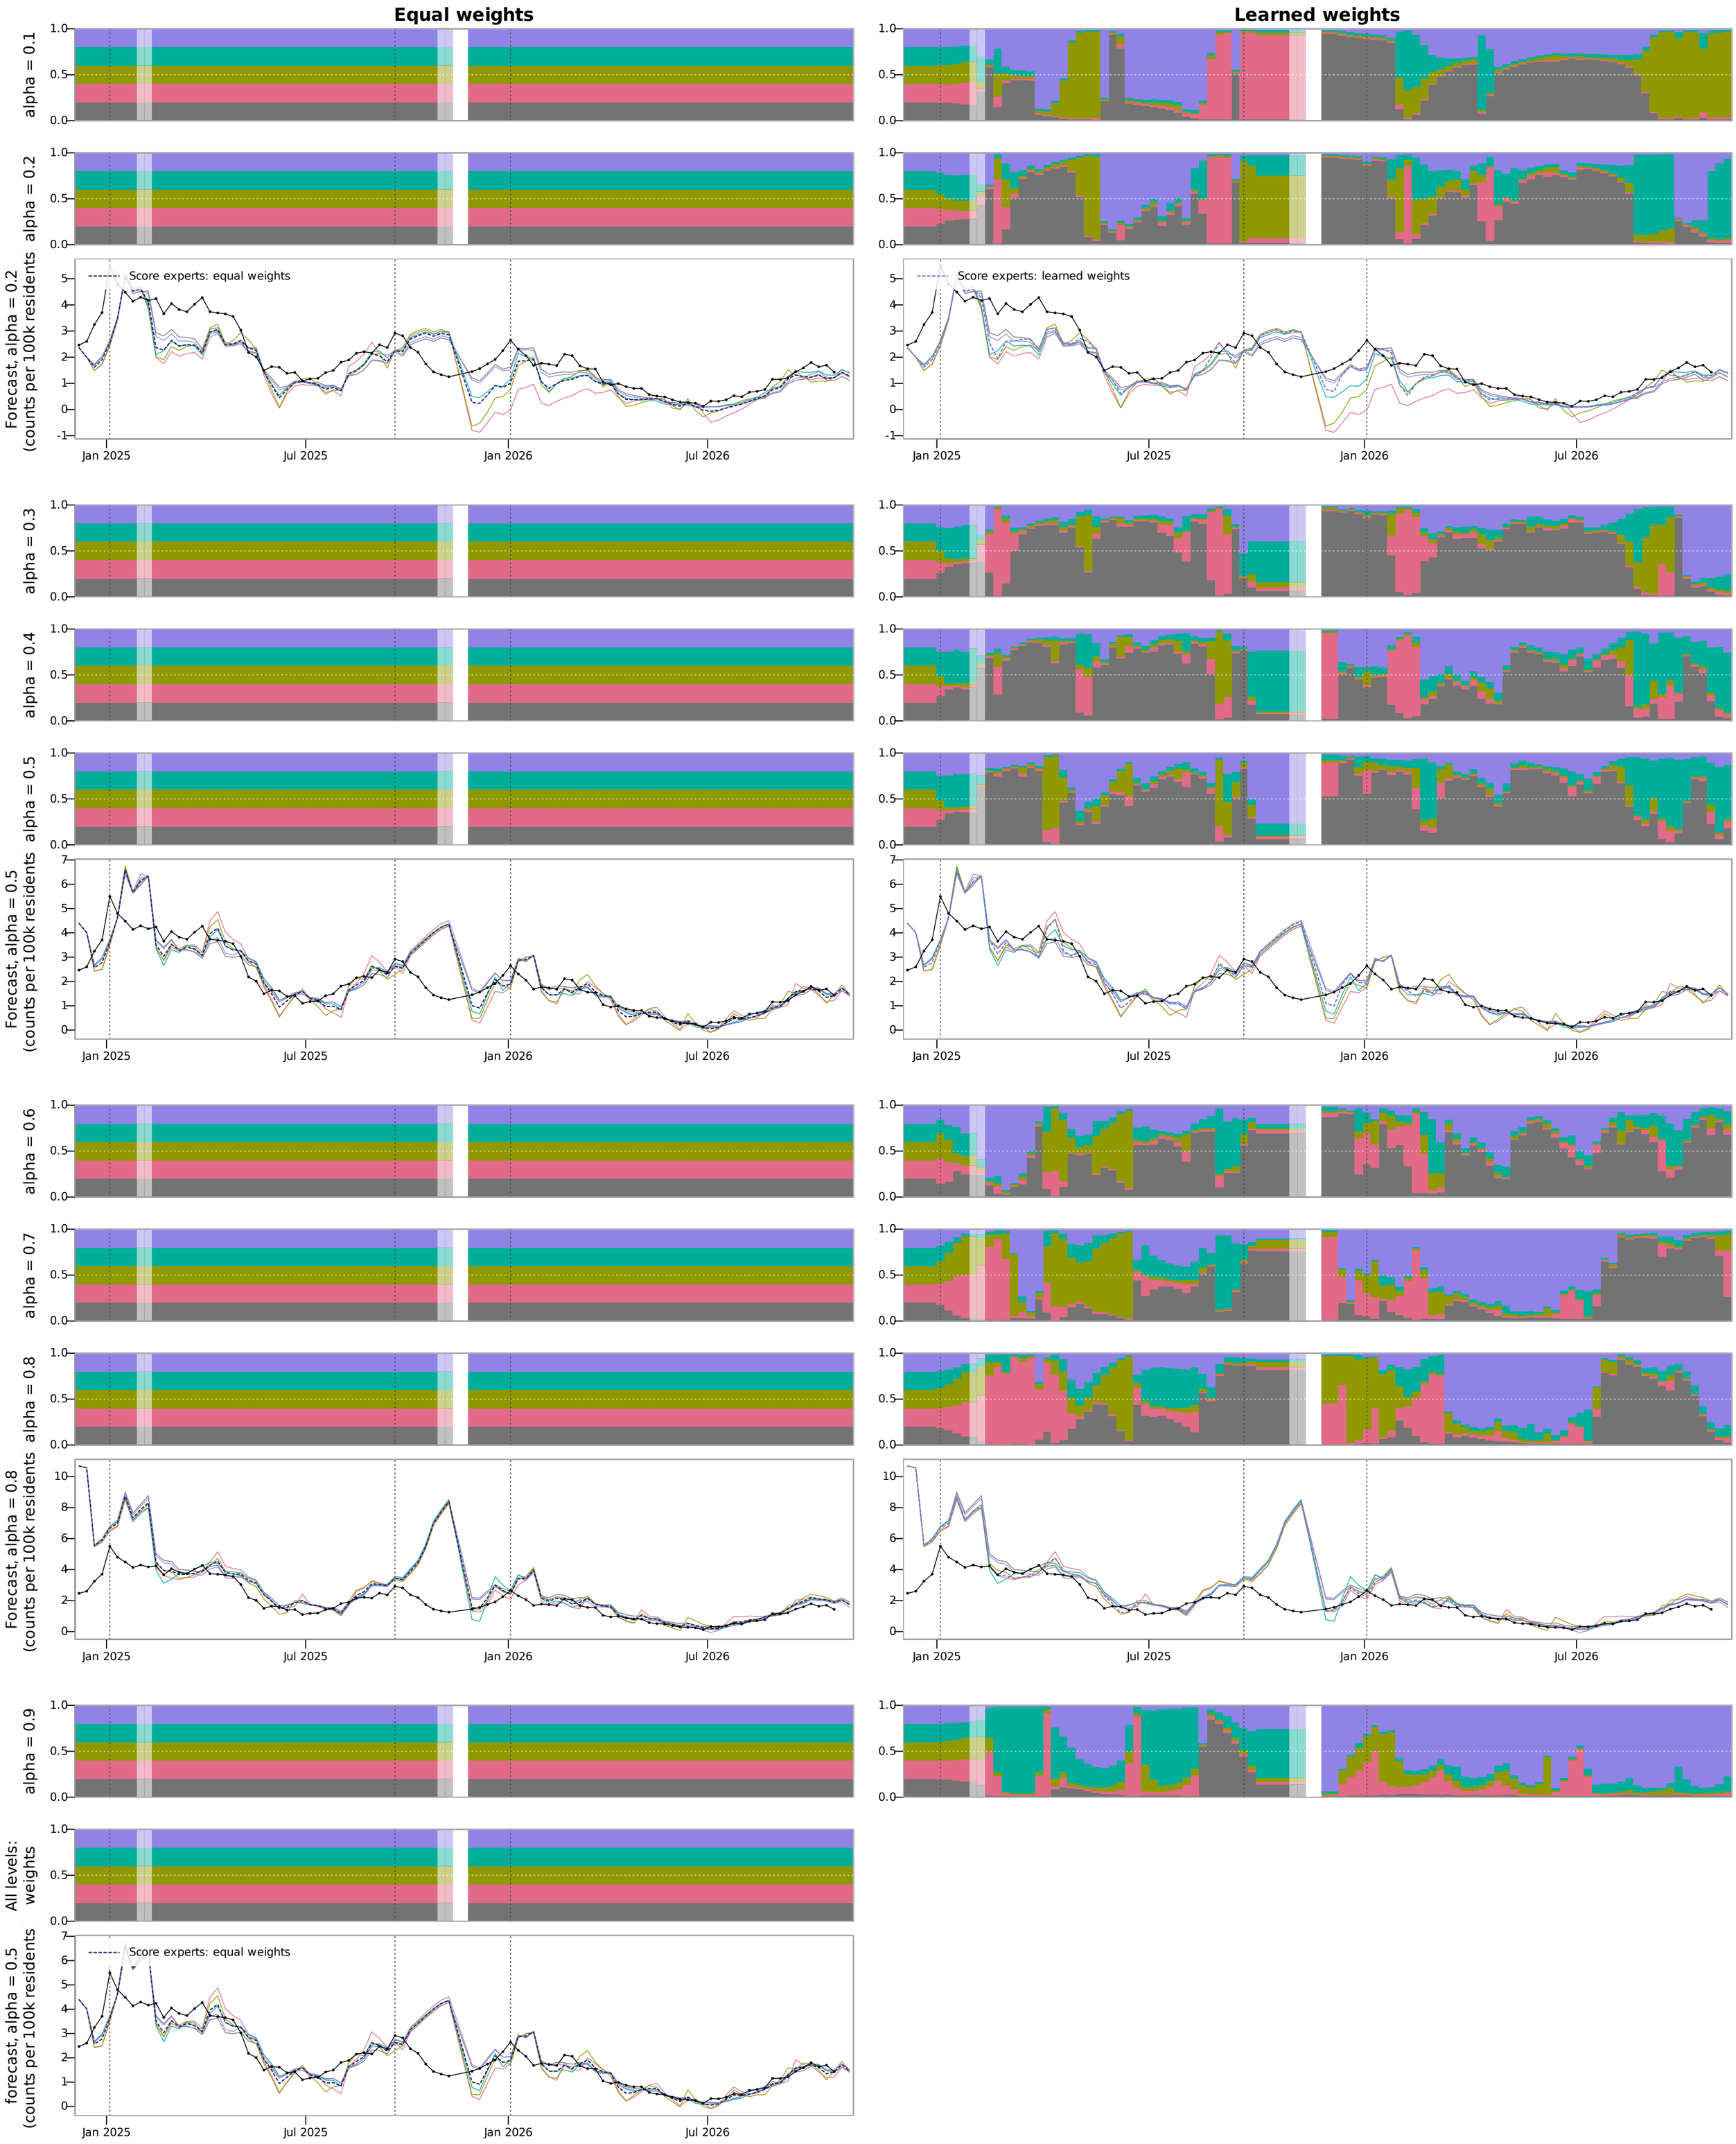
Rows alpha = 0.1-0.9: the weights of each level (one weight vector for all levels: the same at every level); below alpha = 0.2, 0.5 and 0.8 the forecasts at that level (after crossing repair, before any outer correction).

Rows 'All levels': the mixture with one weight vector for all levels (the column's mixture, or its common-weight variant), learned from the mean quantile loss over the levels (equal weights: not learned); its weights and its forecasts at alpha = 0.5.

Lighter bars: rounds without a CMU forecast, filled with the latest earlier CMU forecast of the same target week.

Score-based expert weights

covid | horizon 1 | New York (36) | evaluation: Jun 2025 to Sep 2026 | raw-count updates | feedback re-run



Target date | Weight at issuance for each target | Vertical dotted lines: peaks

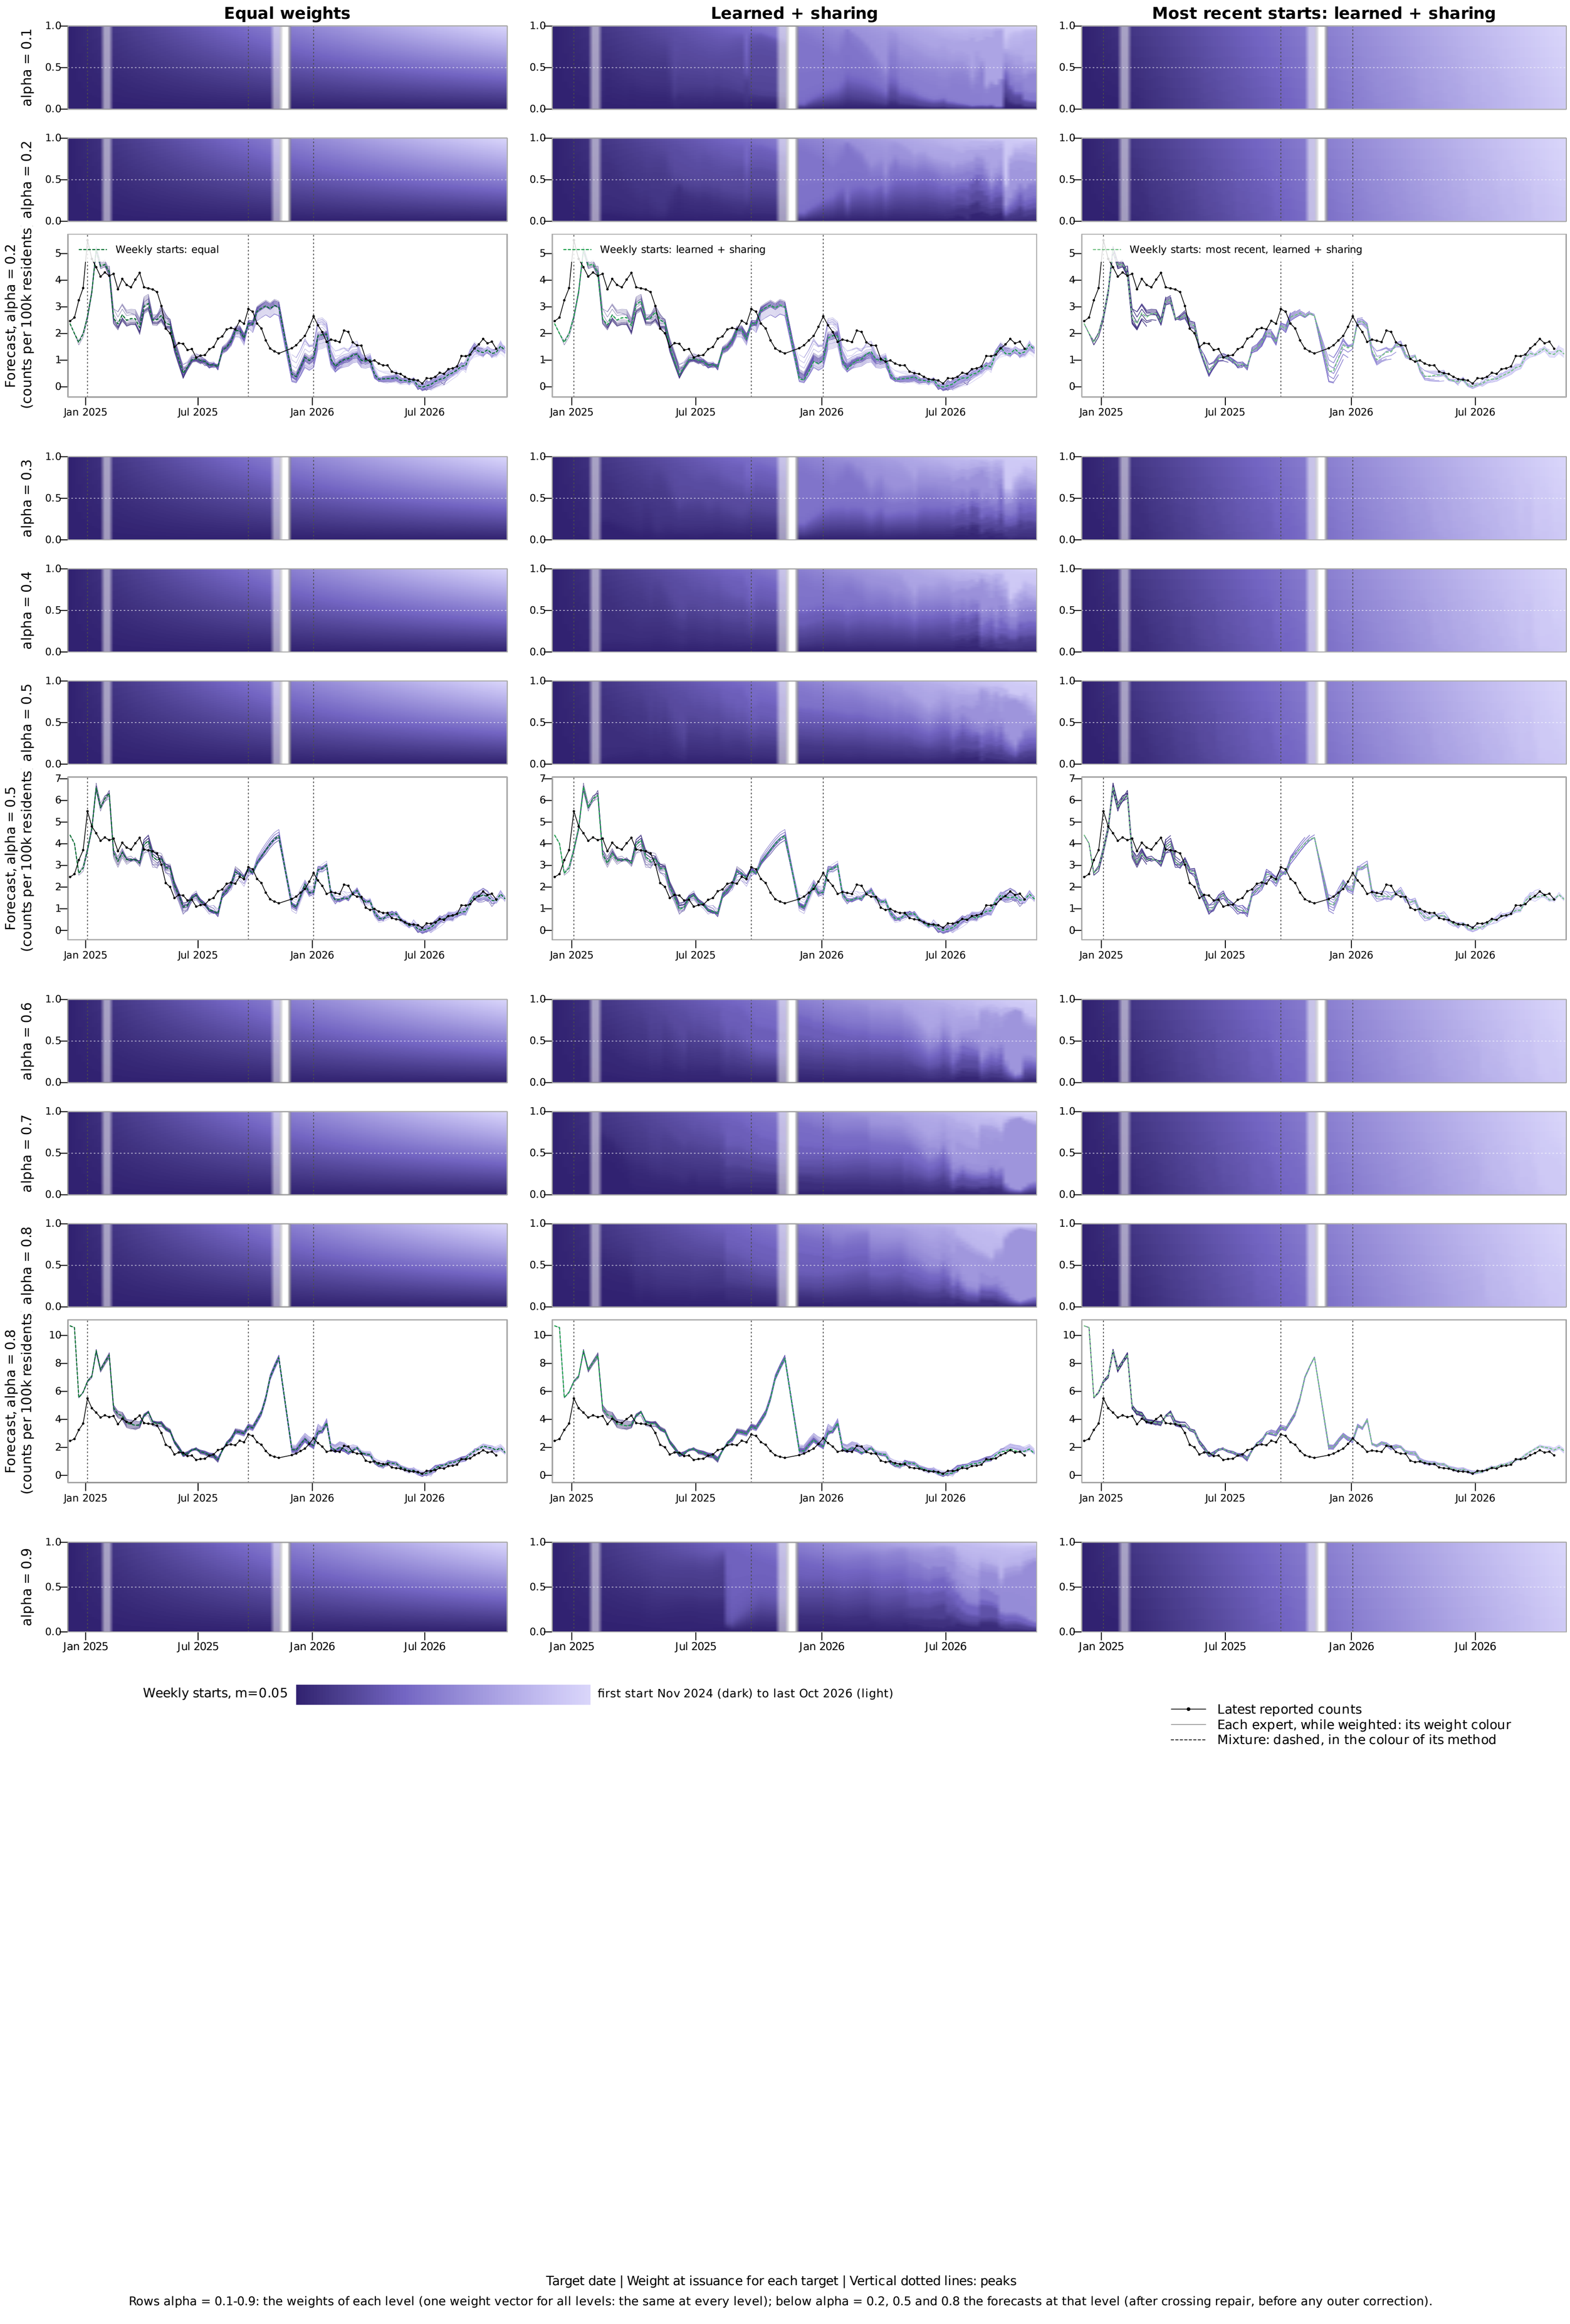
Rows alpha = 0.1-0.9: the weights of each level (one weight vector for all levels: the same at every level); below alpha = 0.2, 0.5 and 0.8 the forecasts at that level (after crossing repair, before any outer correction).

Rows 'All levels': the mixture with one weight vector for all levels (the column's mixture, or its common-weight variant), learned from the mean quantile loss over the levels (equal weights: not learned); its weights and its forecasts at alpha = 0.5.

Lighter bars: rounds without a CMU forecast, filled with the latest earlier CMU forecast of the same target week.

Weekly-start weights

covid | horizon 1 | New York (36) | evaluation: Jun 2025 to Sep 2026 | raw-count updates | feedback re-run



Target date | Weight at issuance for each target | Vertical dotted lines: peaks

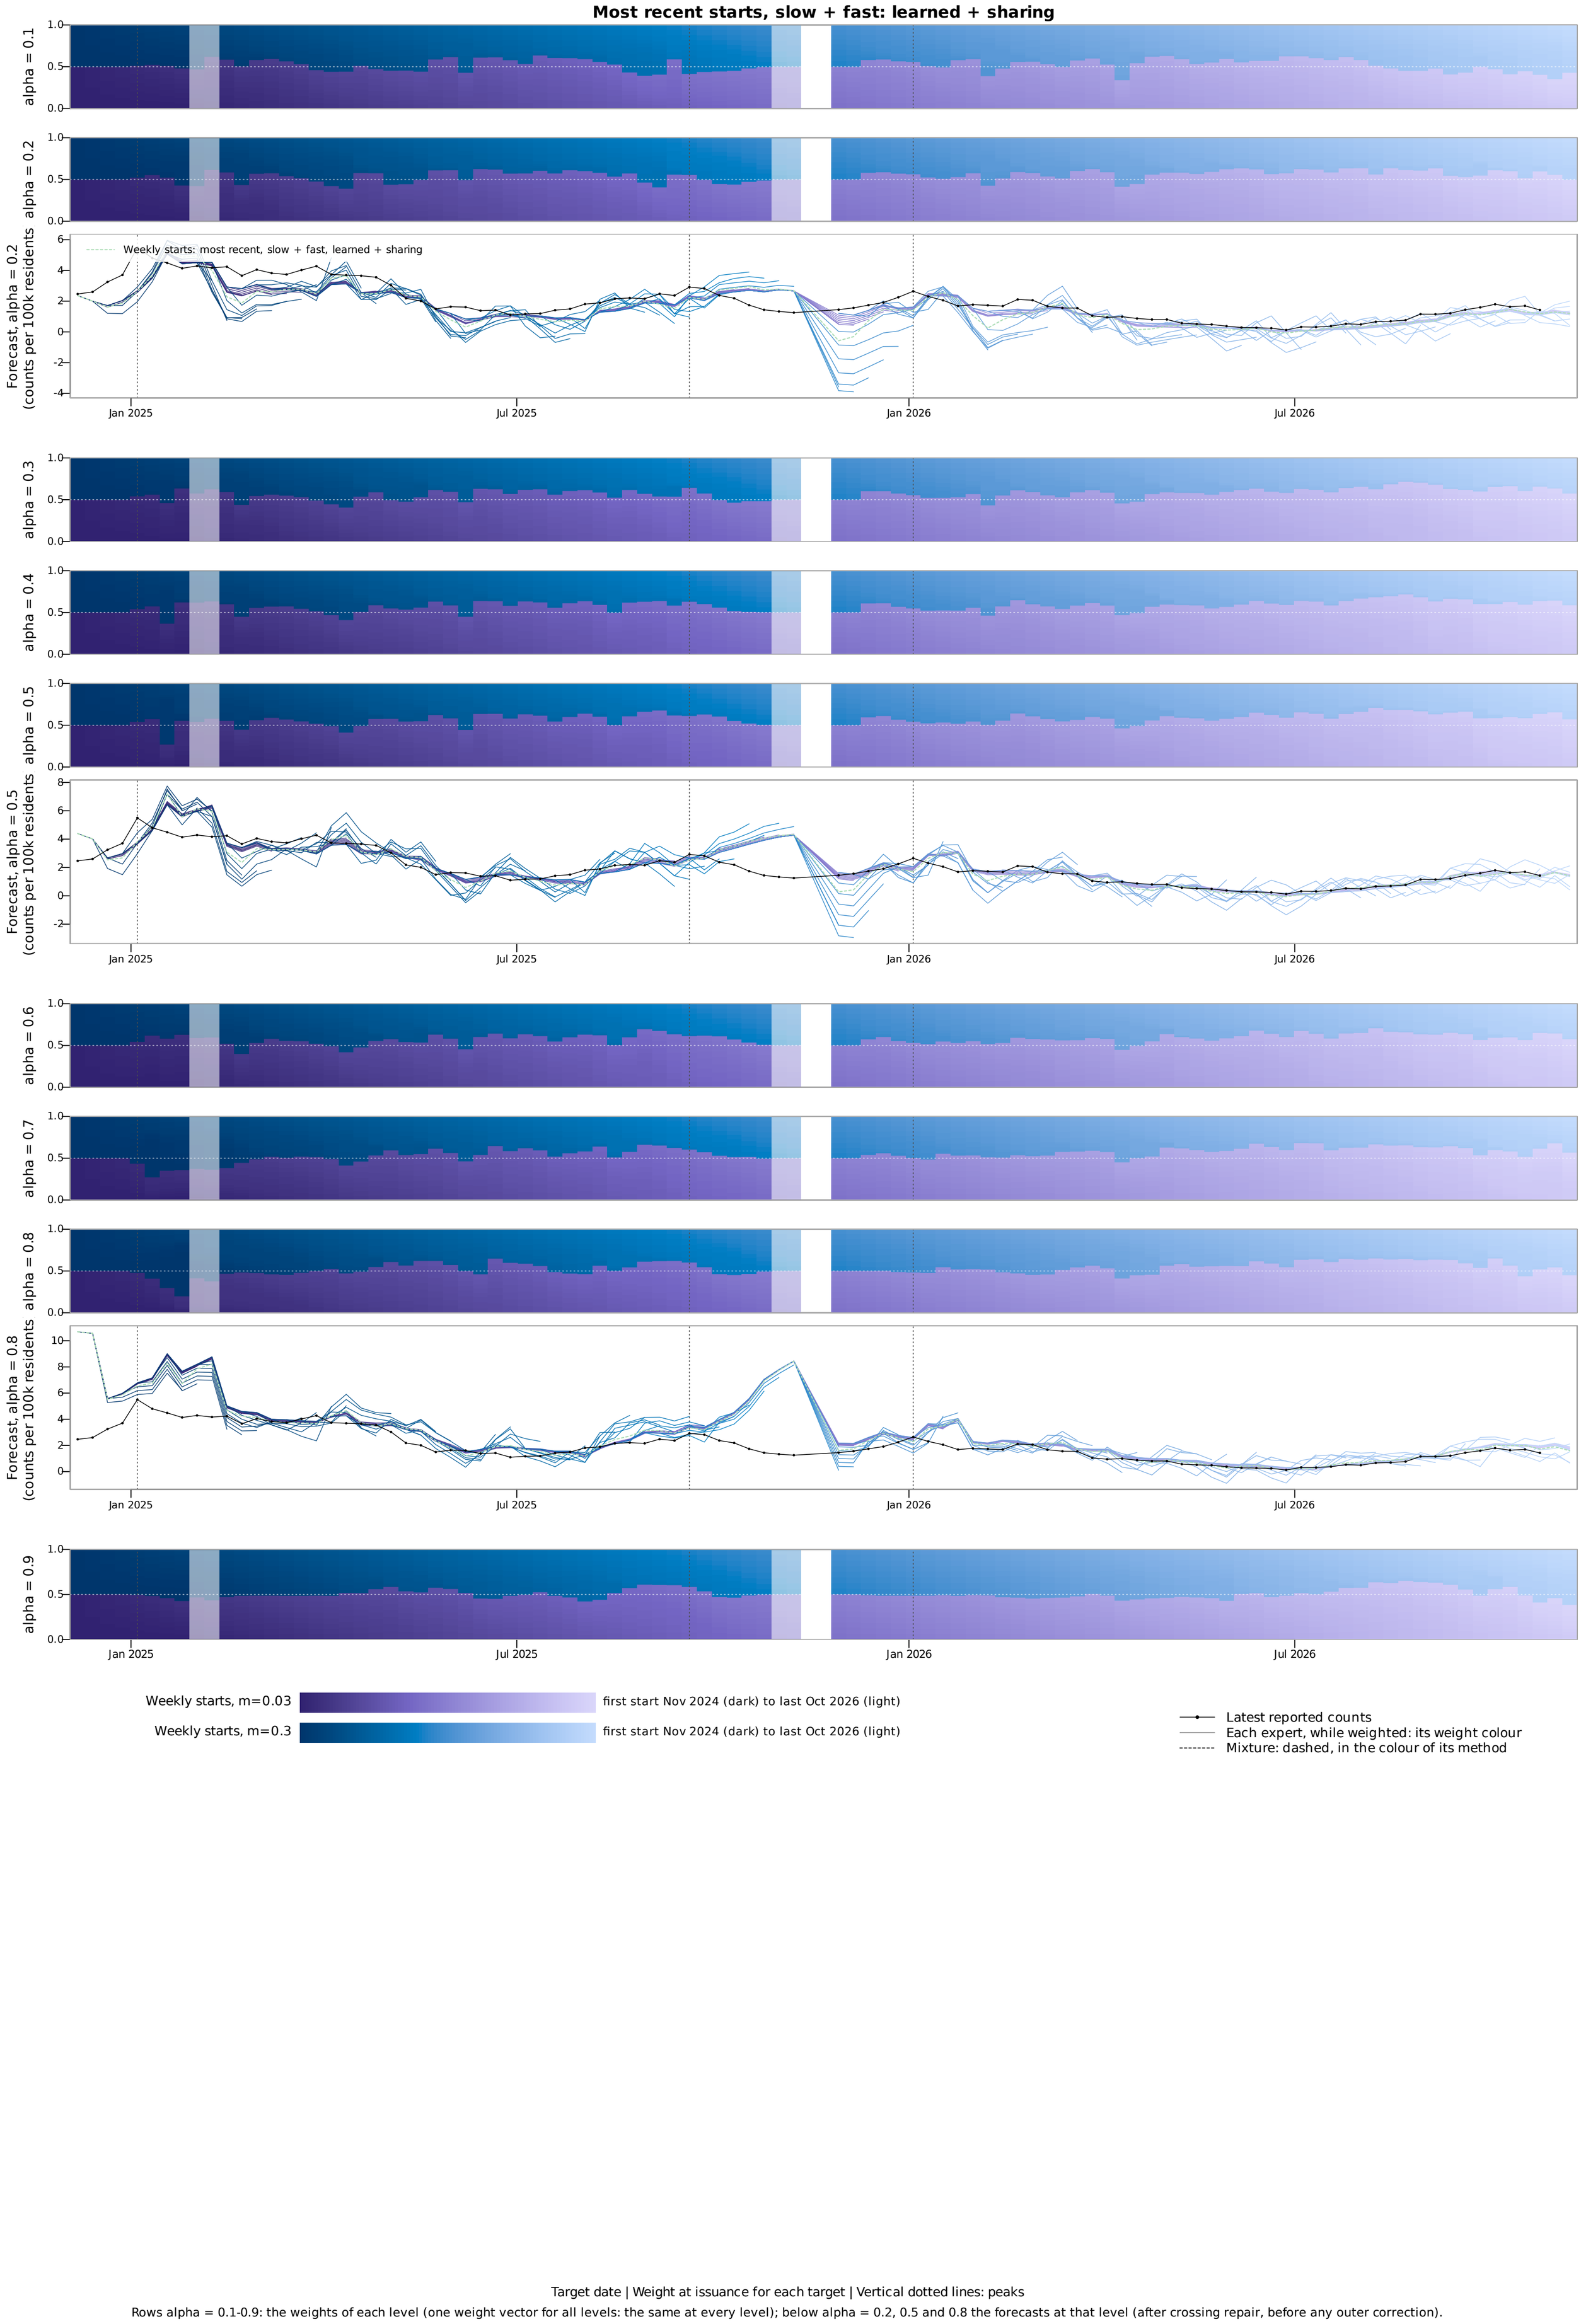
Rows alpha = 0.1-0.9: the weights of each level (one weight vector for all levels: the same at every level); below alpha = 0.2, 0.5 and 0.8 the forecasts at that level (after crossing repair, before any outer correction).

Lighter bars: rounds without a CMU forecast, filled with the latest earlier CMU forecast of the same target week.

Weekly-start weights: slow + fast copies

covid | horizon 1 | New York (36) | evaluation: Jun 2025 to Sep 2026 | raw-count updates | feedback re-run

Most recent starts, slow + fast: learned + sharing



Target date | Weight at issuance for each target | Vertical dotted lines: peaks

Rows alpha = 0.1-0.9: the weights of each level (one weight vector for all levels: the same at every level); below alpha = 0.2, 0.5 and 0.8 the forecasts at that level (after crossing repair, before any outer correction).

Lighter bars: rounds without a CMU forecast, filled with the latest earlier CMU forecast of the same target week.

Frozen rate-grid weights

covid | horizon 1 | Pennsylvania (42) | evaluation: Jun 2025 to Sep 2026 | raw-count updates | feedback re-run

Equal weights

Learned, no sharing

Learned + sharing

alpha = 0.1

alpha = 0.2

Forecast, alpha = 0.2 (counts per 100k residents)

Rate grid: equal

Rate grid: learned, no sharing

Rate grid: learned + sharing

alpha = 0.3

alpha = 0.4

alpha = 0.5

Forecast, alpha = 0.5 (counts per 100k residents)

alpha = 0.6

alpha = 0.7

alpha = 0.8

Forecast, alpha = 0.8 (counts per 100k residents)

alpha = 0.9

All levels: weights

All levels: forecast, alpha = 0.5 (counts per 100k residents)

Rate grid: equal

Rate grid: learned, common weights

Rate grid: learned + sharing, common weights

Jan 2025 | Jul 2025 | Jan 2026 | Jul 2026

Experts: weights and forecasts

m=0 | m=0.01 | m=0.015 | m=0.02 | m=0.03 | m=0.05 | m=0.075 | m=0.1 | m=0.15 | m=0.2 | m=0.3 | m=0.45

Latest reported counts

Each expert, while weighted: its weight colour

Mixture: dashed, in the colour of its method

Target date | Weight at issuance for each target | Vertical dotted lines: peaks

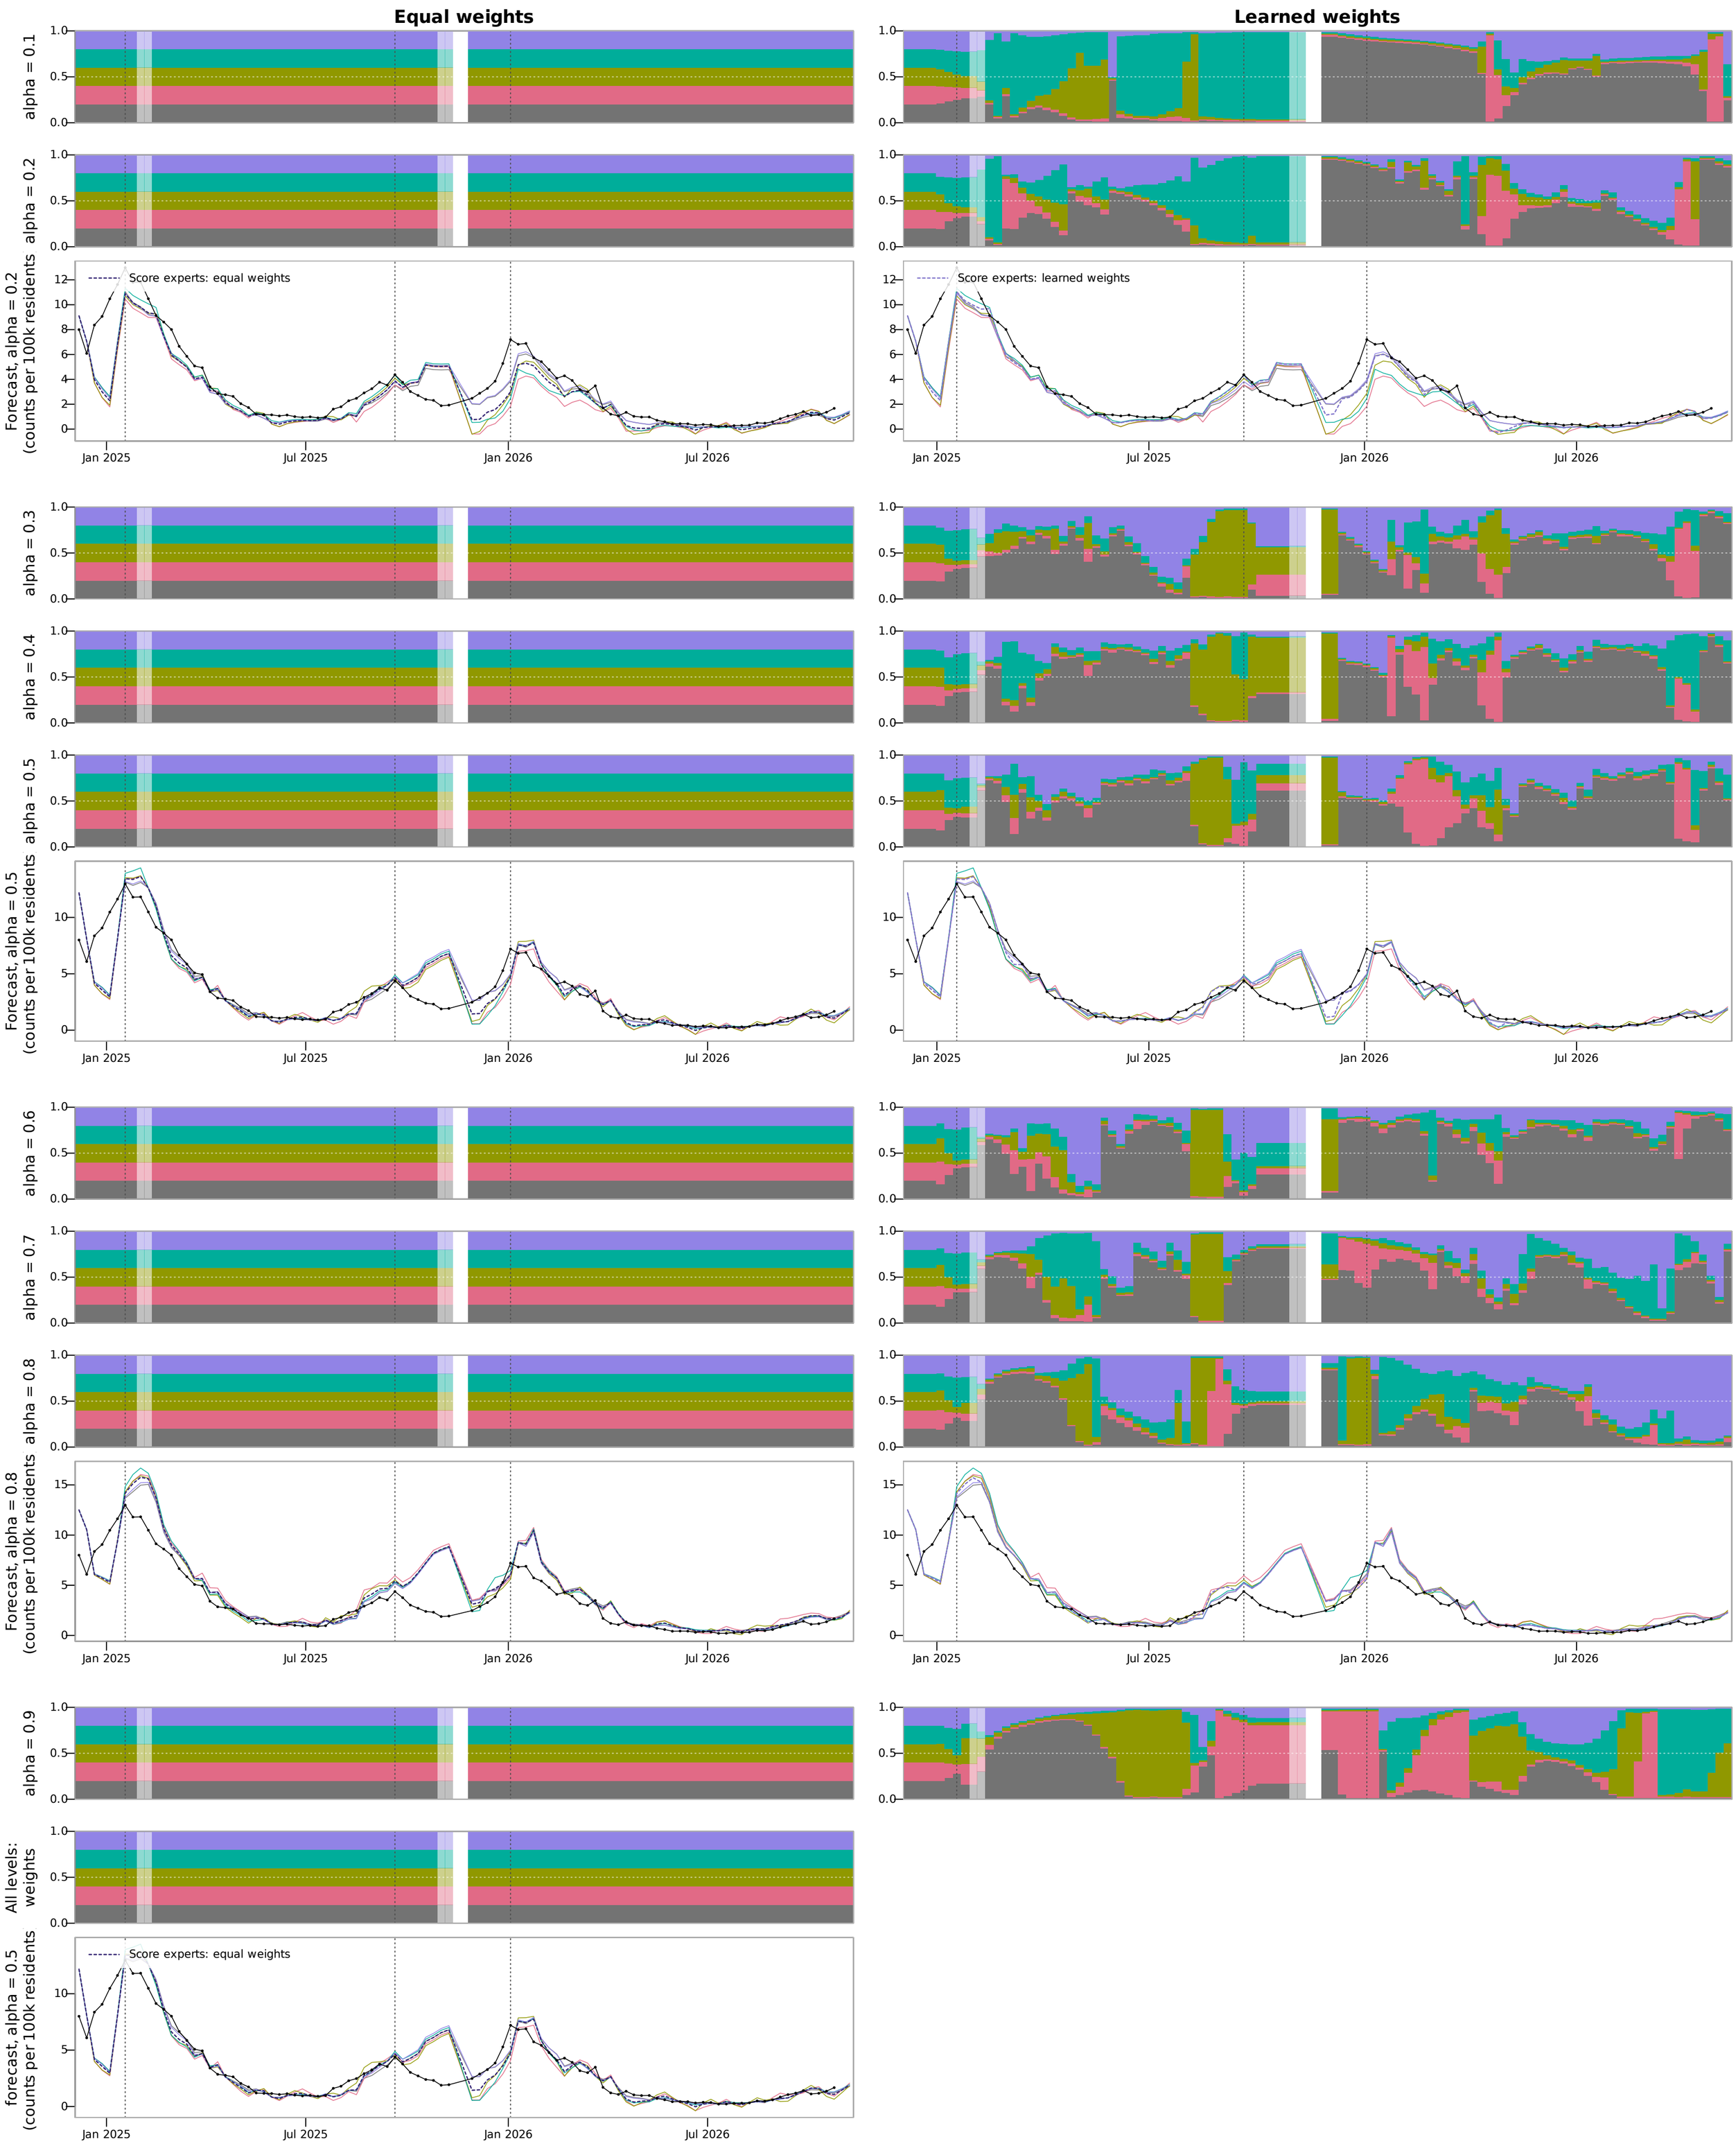
Rows alpha = 0.1-0.9: the weights of each level (one weight vector for all levels: the same at every level); below alpha = 0.2, 0.5 and 0.8 the forecasts at that level (after crossing repair, before any outer correction).

Rows 'All levels': the mixture with one weight vector for all levels (the column's mixture, or its common-weight variant), learned from the mean quantile loss over the levels (equal weights: not learned); its weights and its forecasts at alpha = 0.5.

Lighter bars: rounds without a CMU forecast, filled with the latest earlier CMU forecast of the same target week.

Score-based expert weights

covid | horizon 1 | Pennsylvania (42) | evaluation: Jun 2025 to Sep 2026 | raw-count updates | feedback re-run



Target date | Weight at issuance for each target | Vertical dotted lines: peaks

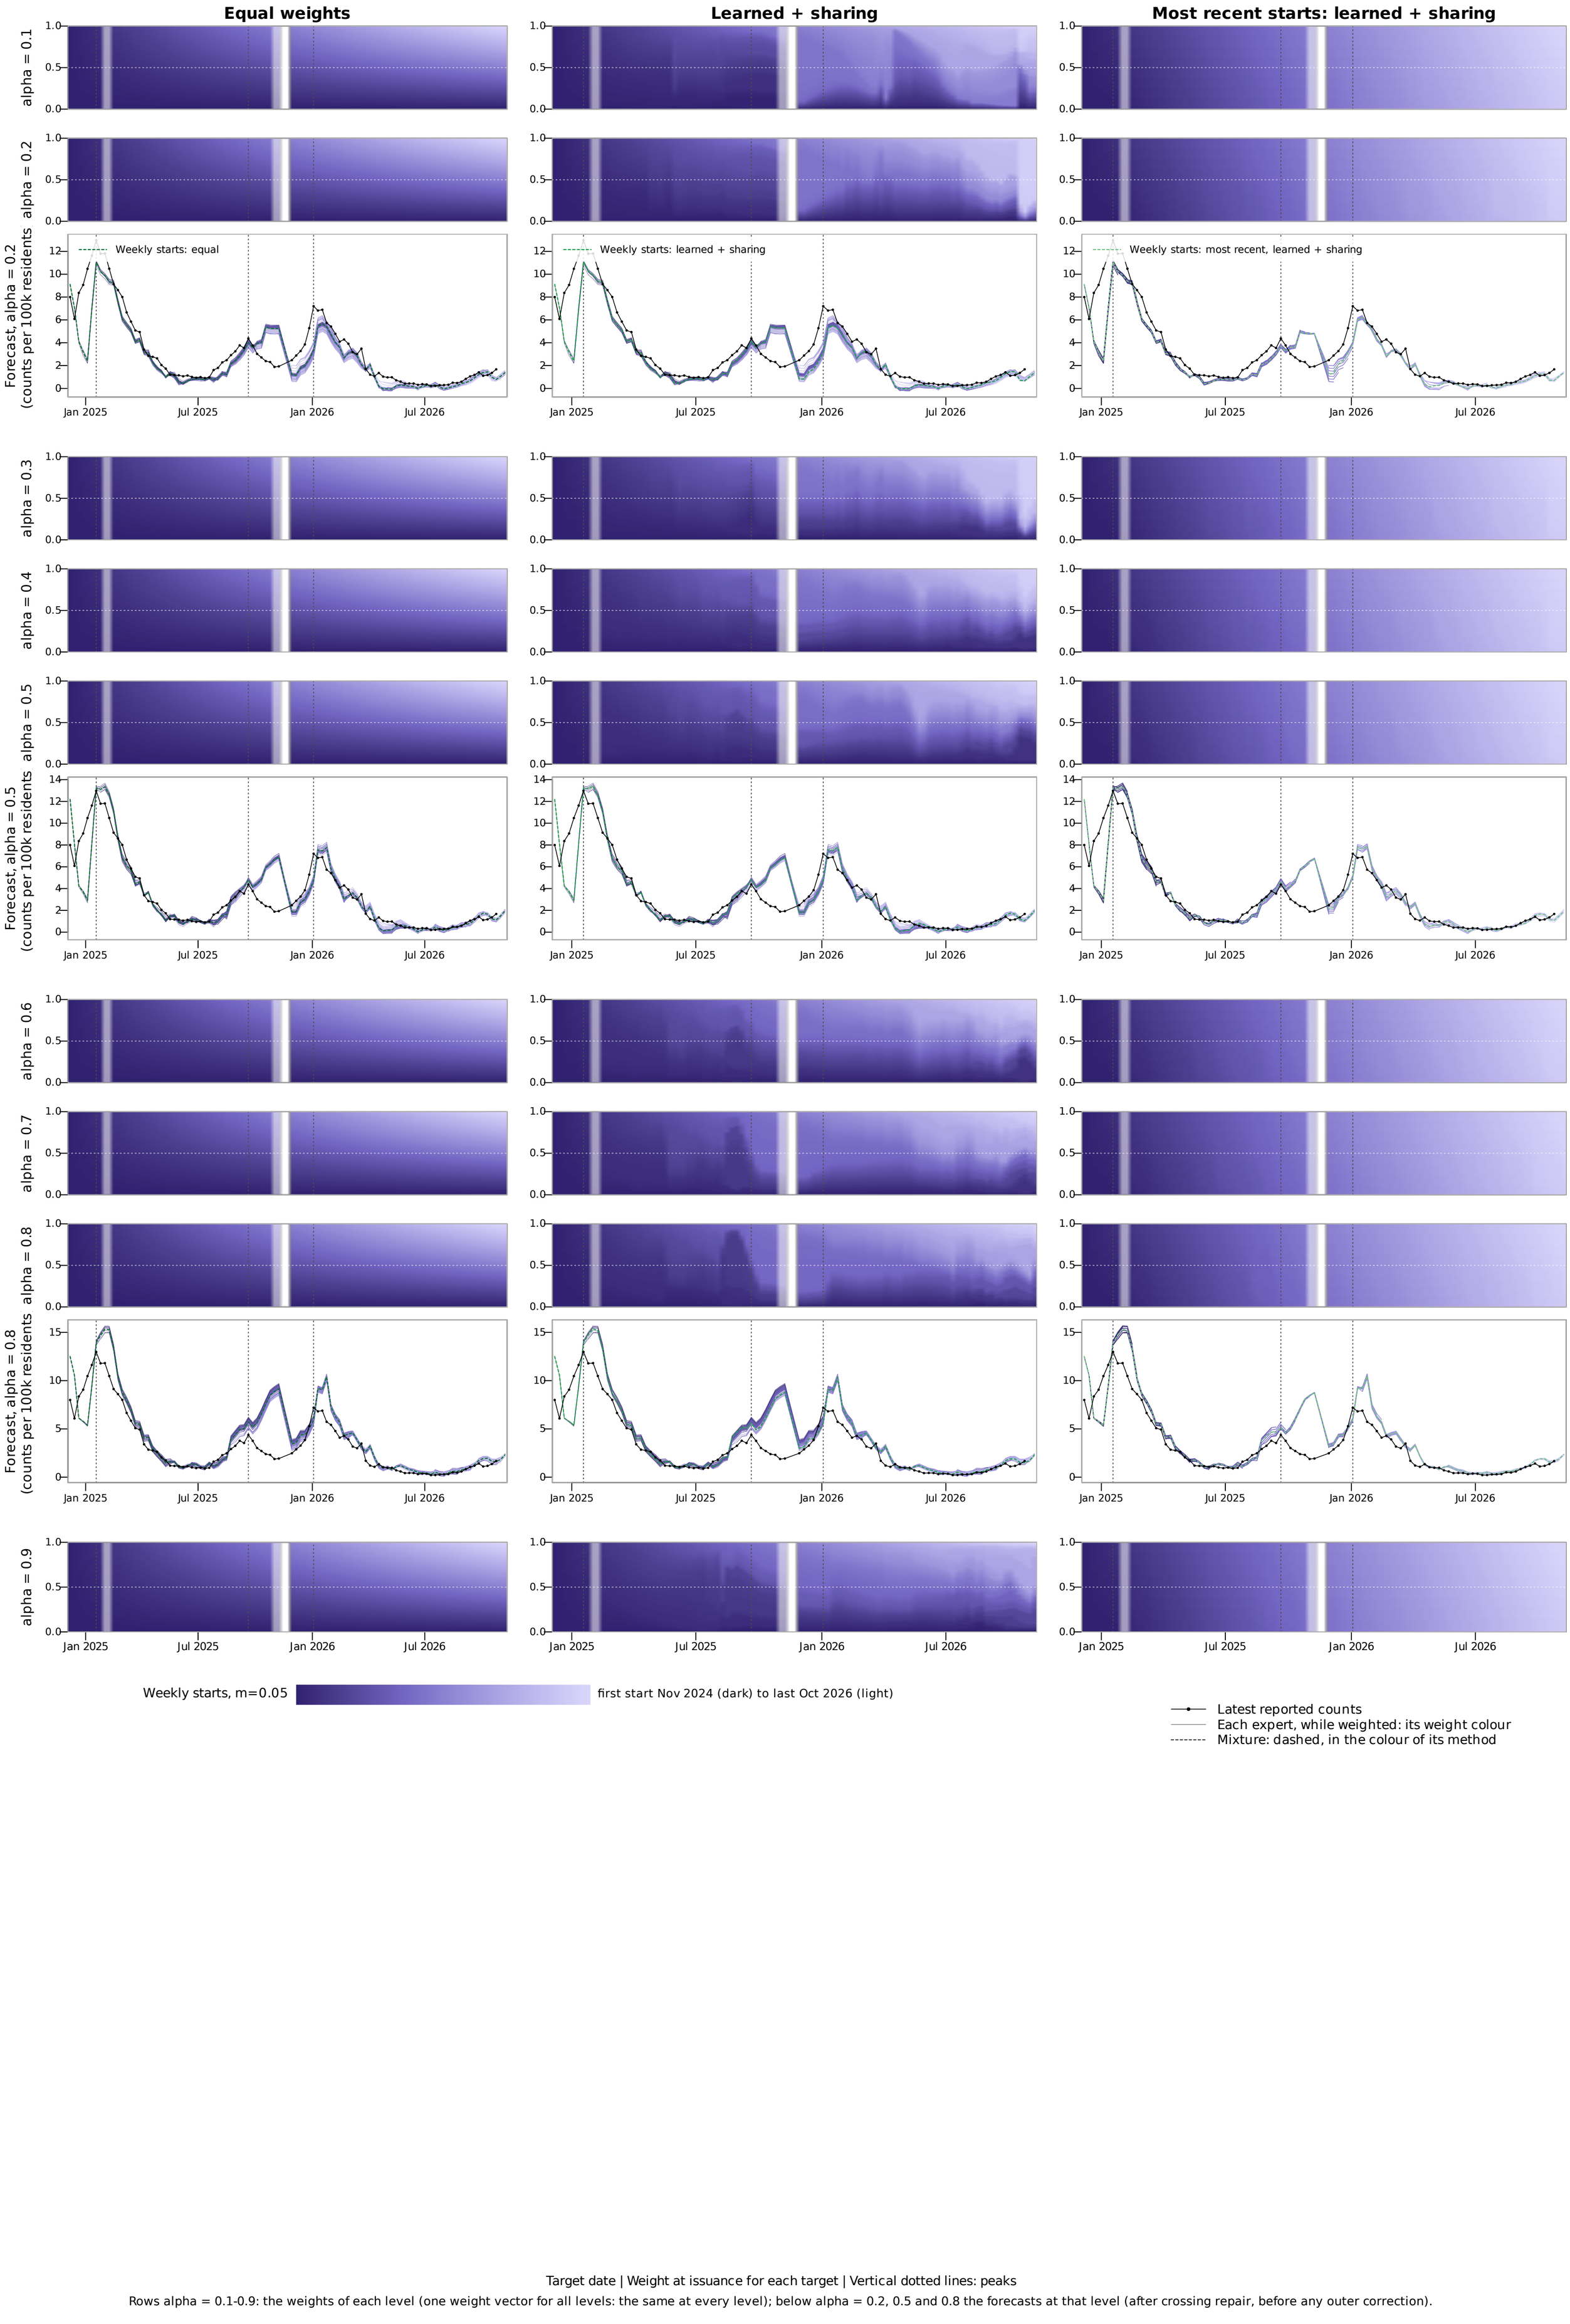
Rows alpha = 0.1-0.9: the weights of each level (one weight vector for all levels: the same at every level); below alpha = 0.2, 0.5 and 0.8 the forecasts at that level (after crossing repair, before any outer correction).

Rows 'All levels': the mixture with one weight vector for all levels (the column's mixture, or its common-weight variant), learned from the mean quantile loss over the levels (equal weights: not learned); its weights and its forecasts at alpha = 0.5.

Lighter bars: rounds without a CMU forecast, filled with the latest earlier CMU forecast of the same target week.

Weekly-start weights

covid | horizon 1 | Pennsylvania (42) | evaluation: Jun 2025 to Sep 2026 | raw-count updates | feedback re-run



Target date | Weight at issuance for each target | Vertical dotted lines: peaks

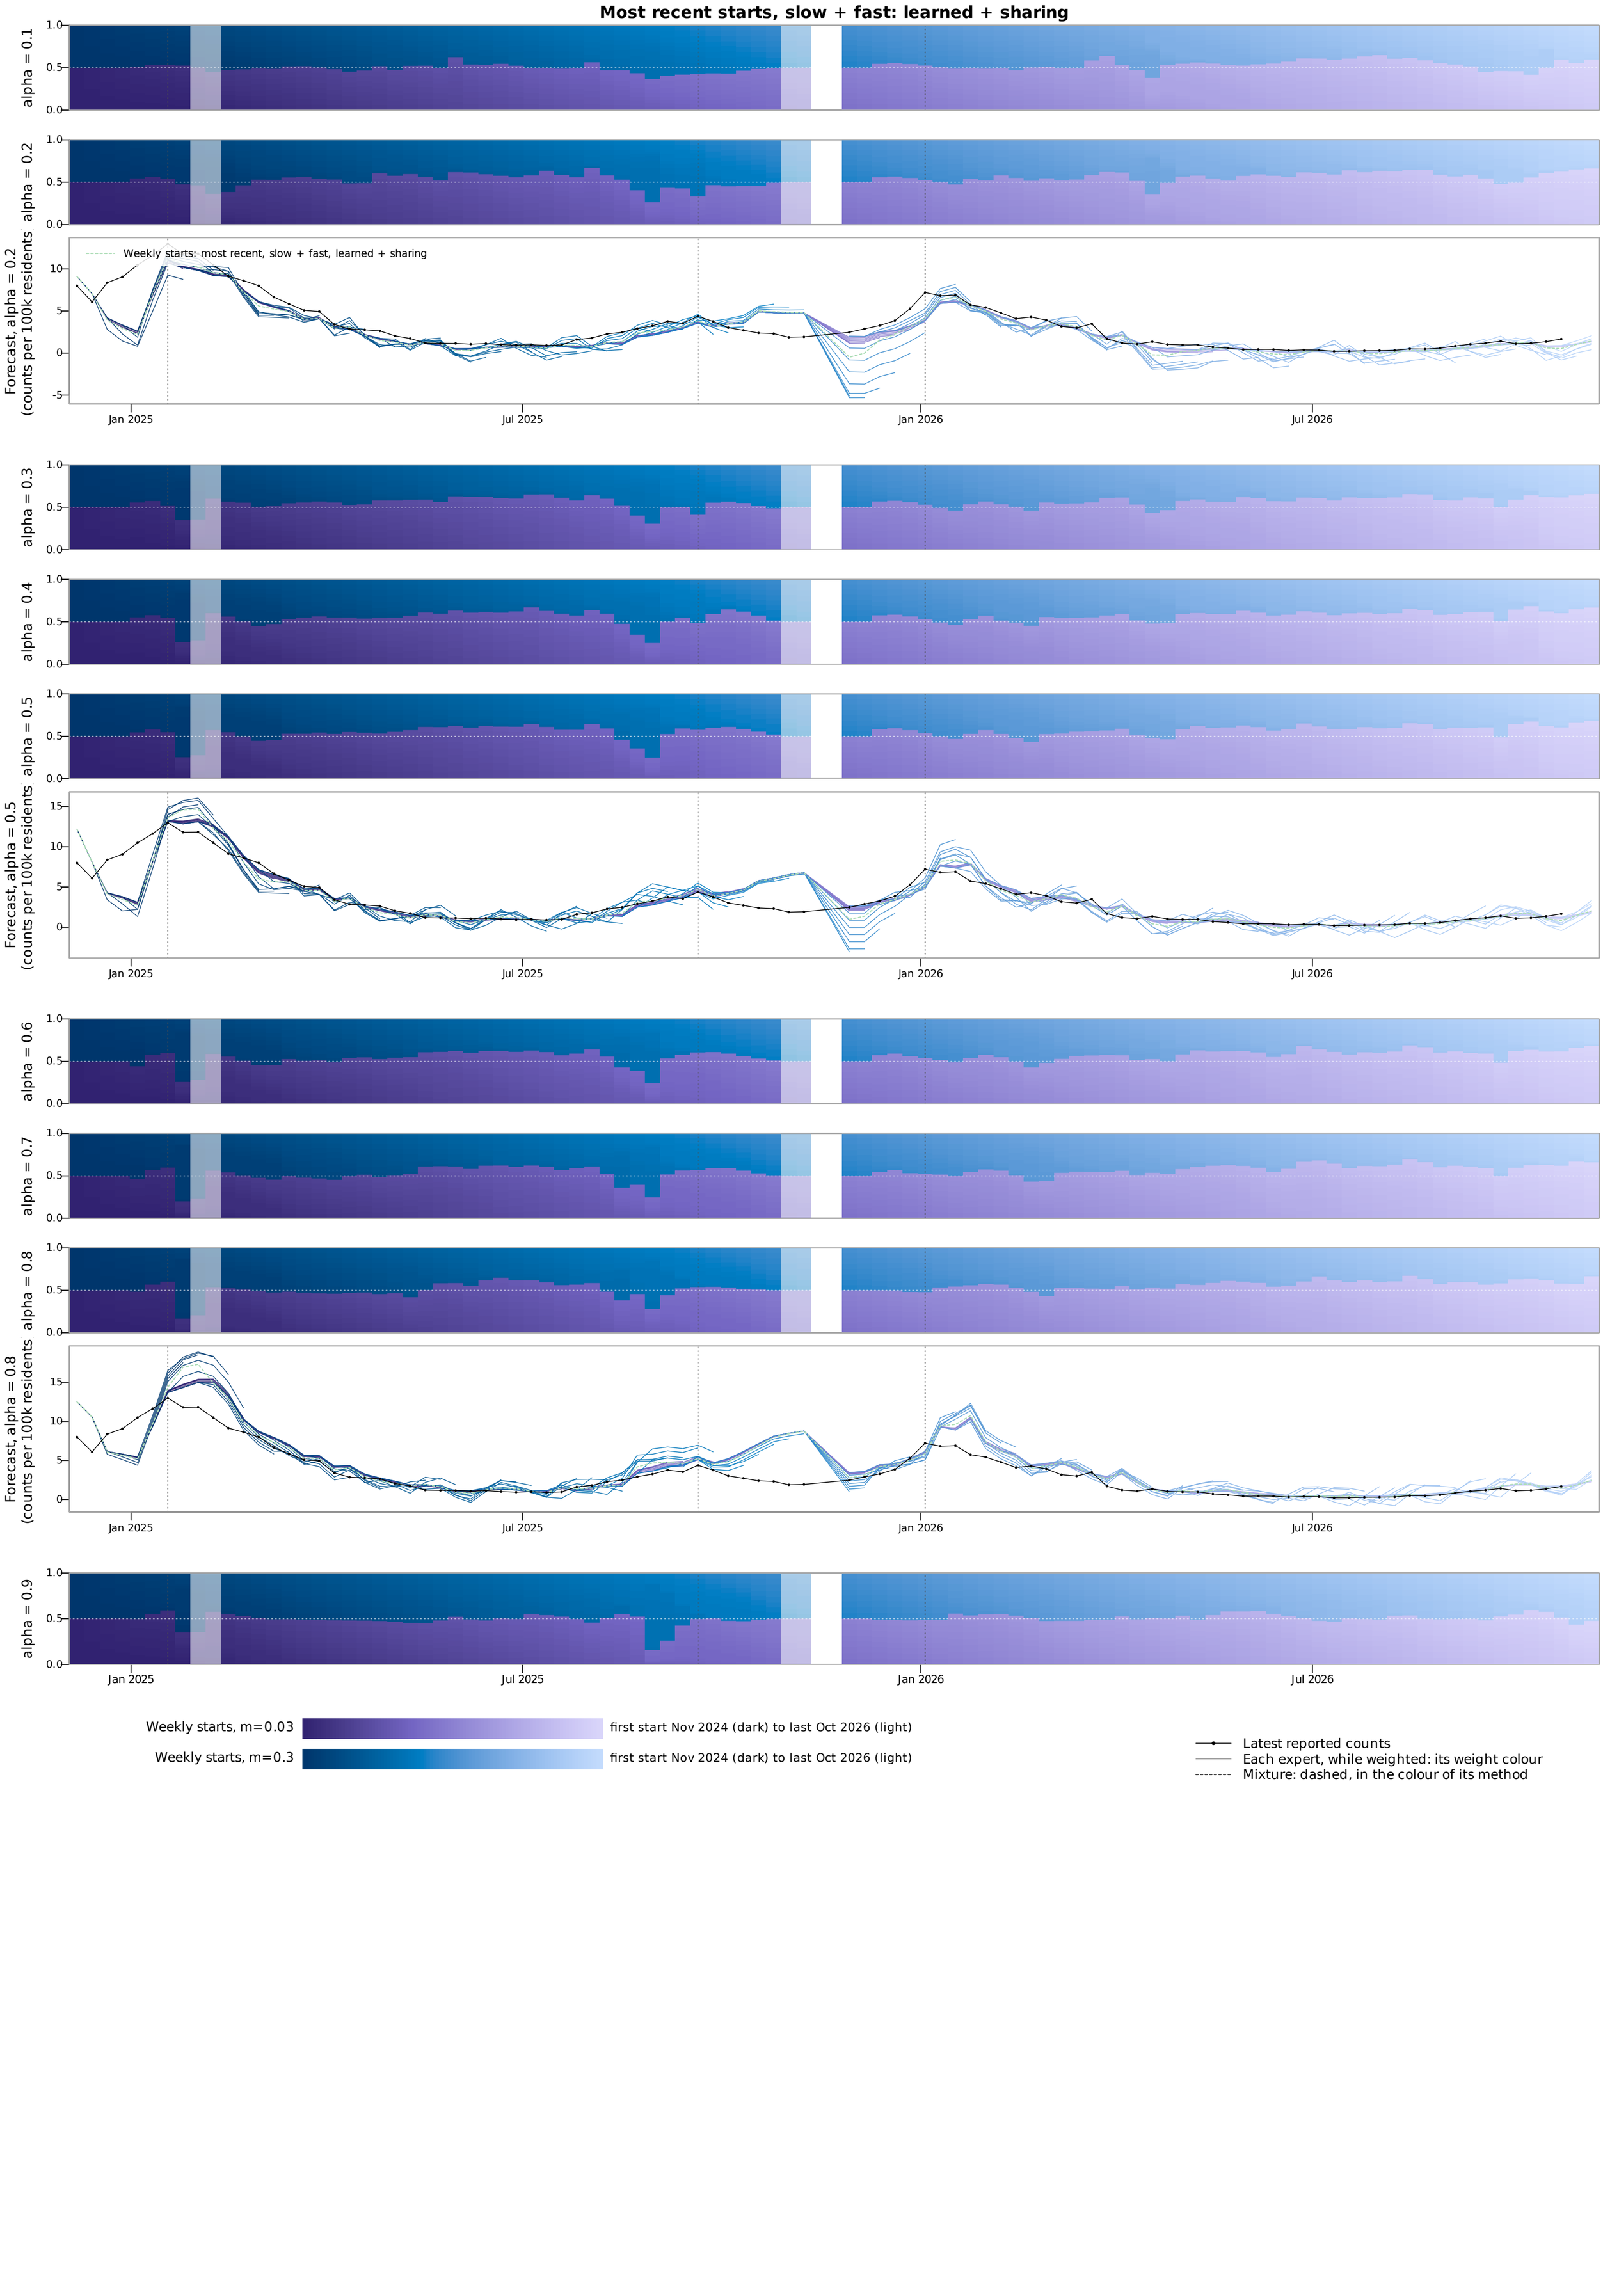
Rows alpha = 0.1-0.9: the weights of each level (one weight vector for all levels: the same at every level); below alpha = 0.2, 0.5 and 0.8 the forecasts at that level (after crossing repair, before any outer correction).

Lighter bars: rounds without a CMU forecast, filled with the latest earlier CMU forecast of the same target week.

Weekly-start weights: slow + fast copies

covid | horizon 1 | Pennsylvania (42) | evaluation: Jun 2025 to Sep 2026 | raw-count updates | feedback re-run

Most recent starts, slow + fast: learned + sharing



Target date | Weight at issuance for each target | Vertical dotted lines: peaks

Rows alpha = 0.1-0.9: the weights of each level (one weight vector for all levels: the same at every level); below alpha = 0.2, 0.5 and 0.8 the forecasts at that level (after crossing repair, before any outer correction).

Lighter bars: rounds without a CMU forecast, filled with the latest earlier CMU forecast of the same target week.